

# Triton-distributed: Programming Overlapping Kernels on Distributed AI Systems with the Triton Compiler

---

Size Zheng<sup>1,2,†</sup>, Wenlei Bao<sup>1,†</sup>, Qi Hou<sup>1</sup>, Xuegui Zheng<sup>1</sup>, Jin Fang<sup>1</sup>, Chenhui Huang<sup>1</sup>, Tianqi Li<sup>1,3</sup>, Haojie Duanmu<sup>1,4</sup>, Renze Chen<sup>1,3</sup>, Ruifan Xu<sup>1,3</sup>, Yifan Guo<sup>1,5</sup>, Ningxin Zheng<sup>1</sup>, Ziheng Jiang<sup>1</sup>, Xinyi Di<sup>1</sup>, Dongyang Wang<sup>1</sup>, Jianxi Ye<sup>1</sup>, Haibin Lin<sup>1</sup>, Li-Wen Chang<sup>1</sup>, Liqiang Lu<sup>5</sup>, Yun Liang<sup>3</sup>, Jidong Zhai<sup>2</sup>, Xin Liu<sup>1,†</sup>

<sup>1</sup>ByteDance Seed, <sup>2</sup>Tsinghua University, <sup>3</sup>Peking University,

<sup>4</sup>Shanghai Jiao Tong University, <sup>5</sup>Zhejiang University

†Corresponding authors

Shared by Zhou Ouxiang

2025.6.17



# TL;DR (Too Long; Didn't Read)

- ◆ The authors developed **Triton-distributed** to achieve performance competitive with low-level CUDA/C++ at a **fraction of the development cost**.
- ◆ The approach requires **minimal changes** to existing Triton compute kernels.
- ◆ It enables **rapid hardware support**, making it ideal for adapting AI workloads across a diverse ecosystem of chips.



# Outline

- ◆ **Background**
- ◆ **The Triton-distributed Architecture & Programming Model**
- ◆ **Overlapping Optimizations in Triton-distributed**
- ◆ **Experiments & Evaluations**
- ◆ **Conclusion**



# Outline

## ◆ Background

- ◆ The Triton-distributed Architecture & Programming Model
- ◆ Overlapping Optimizations in Triton-distributed
- ◆ Experiments & Evaluations
- ◆ Conclusion



# Background: Beyond a Single Chip

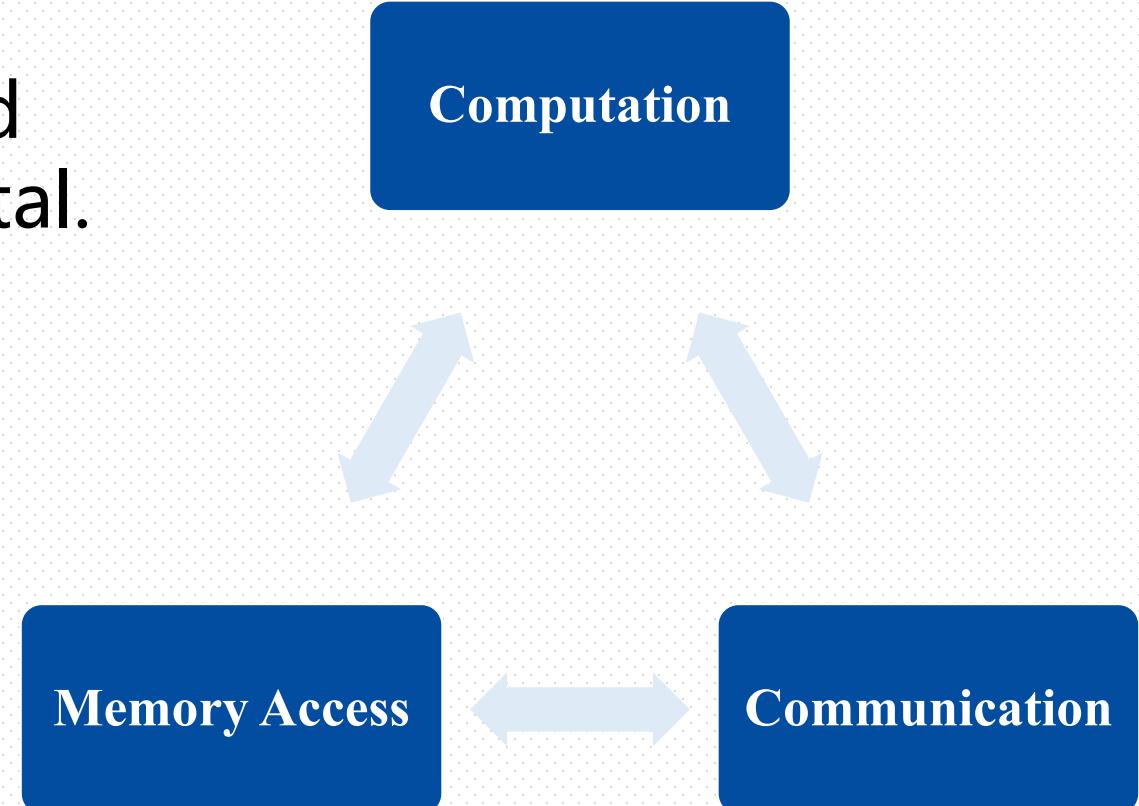
- ◆ Large Language Models (LLMs) have outgrown the memory and compute capacity of single accelerators.
- ◆ Distributed systems, composed of multiple accelerators, are now essential for both training and inference.
- ◆ This shift introduces significant new complexities.



# Background: Computation-Communication Overlap

- ◆ As cluster scale grows, hiding communication latency behind computation time becomes vital.

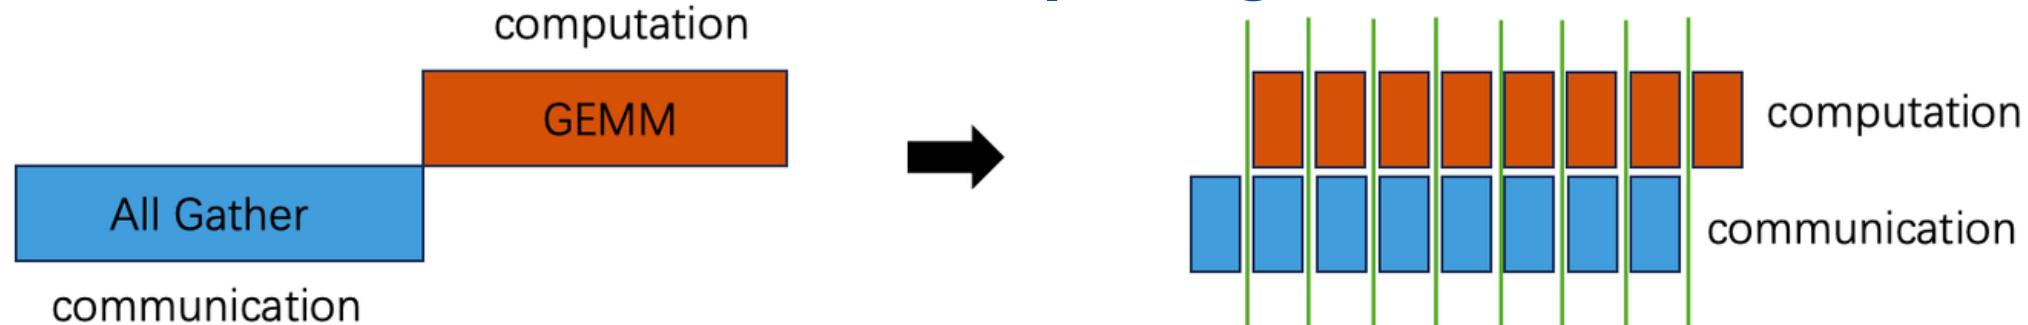
- ◆ Effective overlap can save millions of GPU hours and significant operational costs (e.g., ByteDance's COMET project).



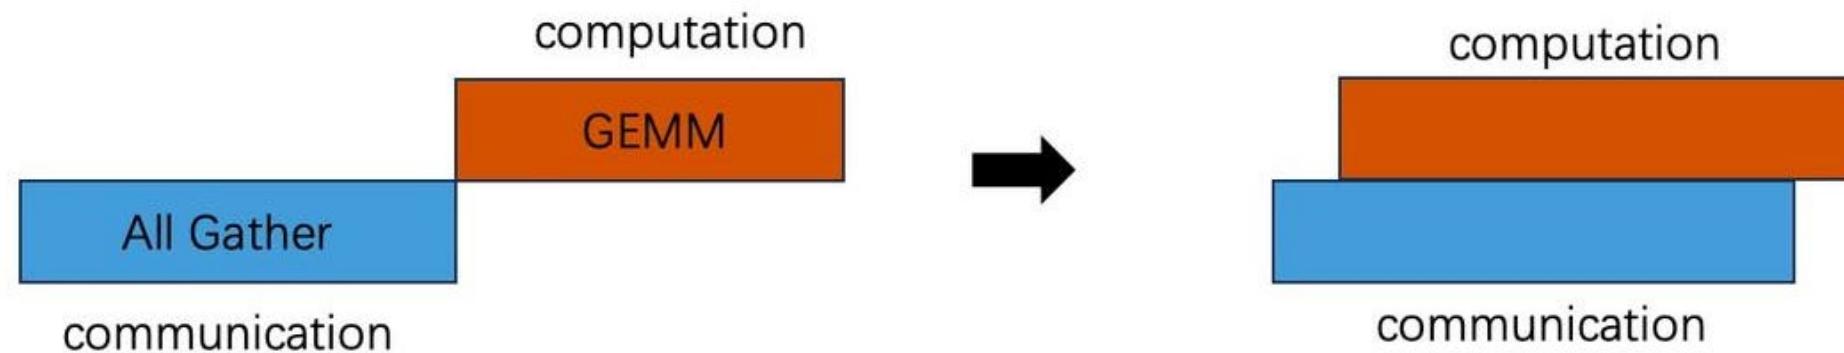


# Background: Computation-Communication Overlap

## Kernel Splitting

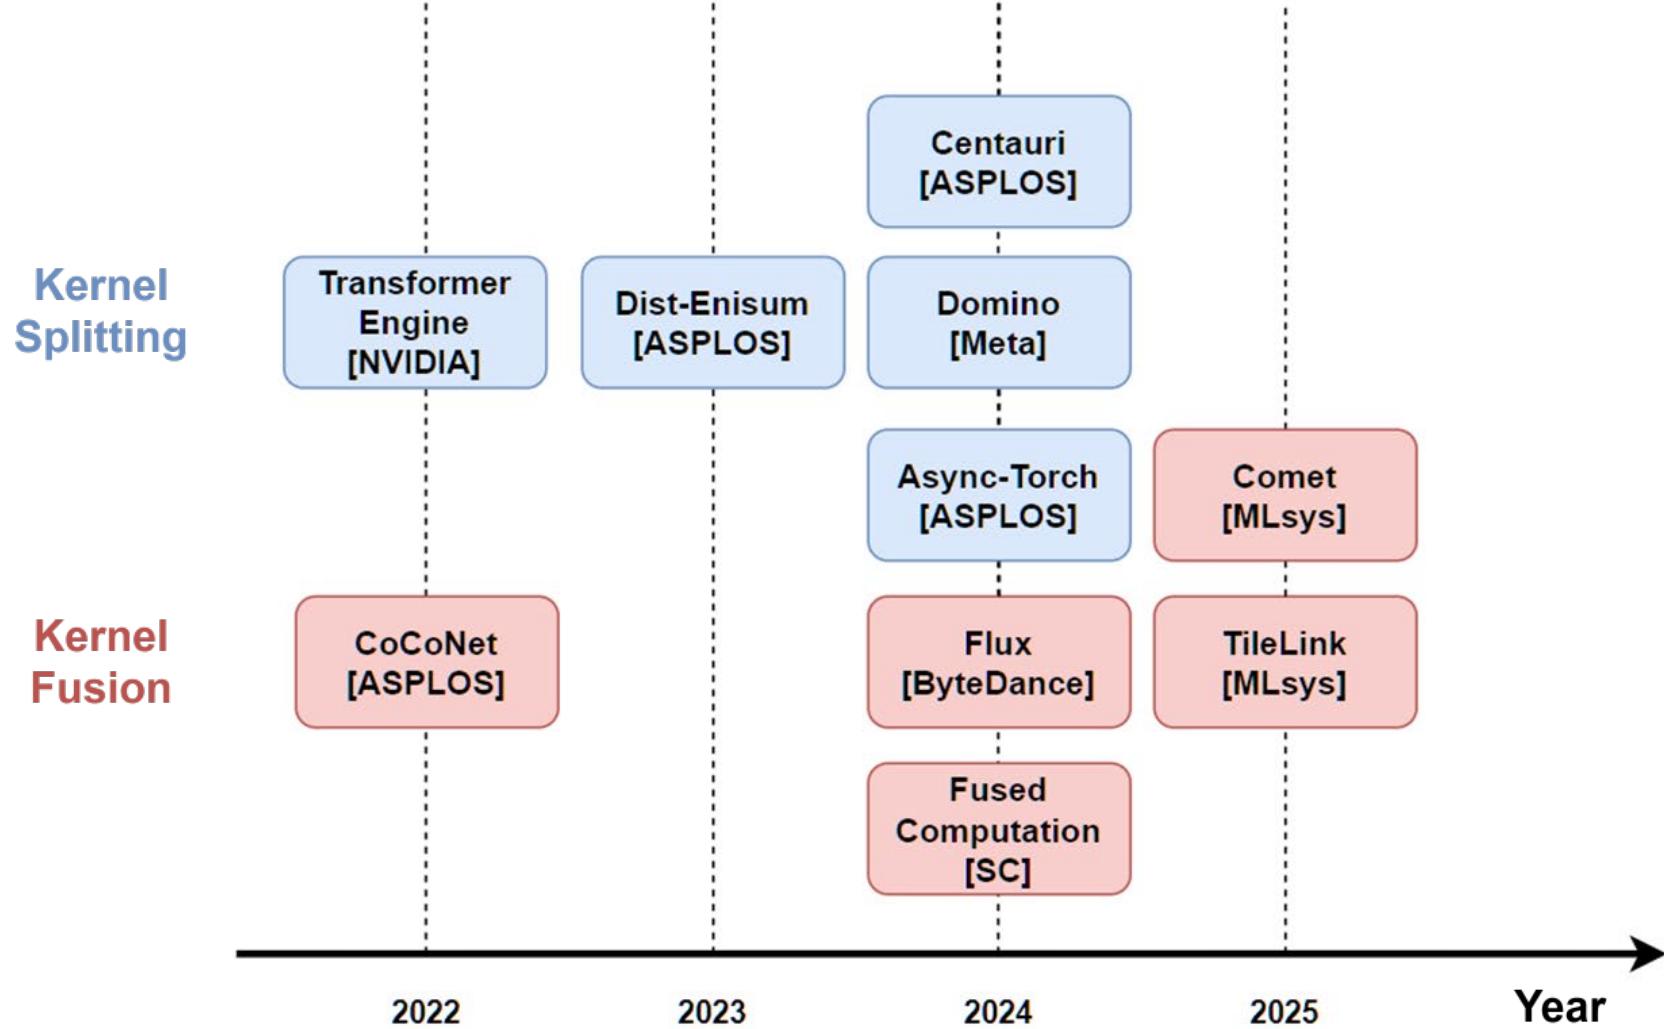


## Kernel Fusion





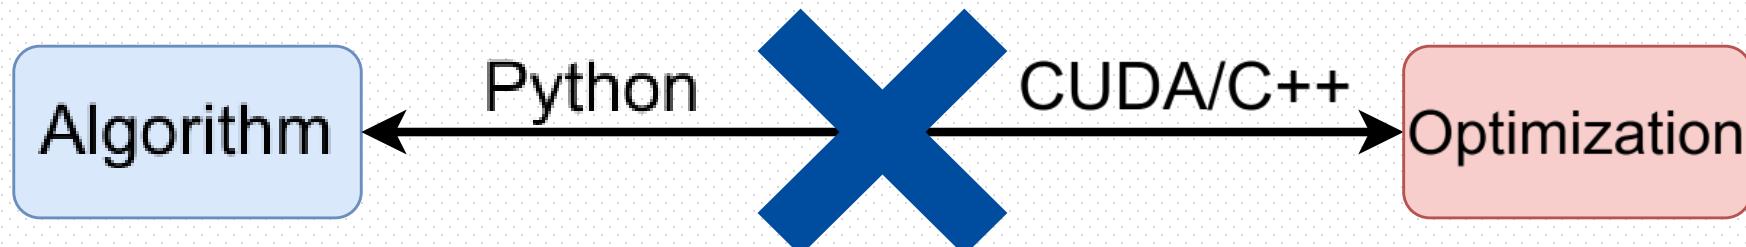
# Background: Computation-Communication Overlap





# The Gap Between Programming

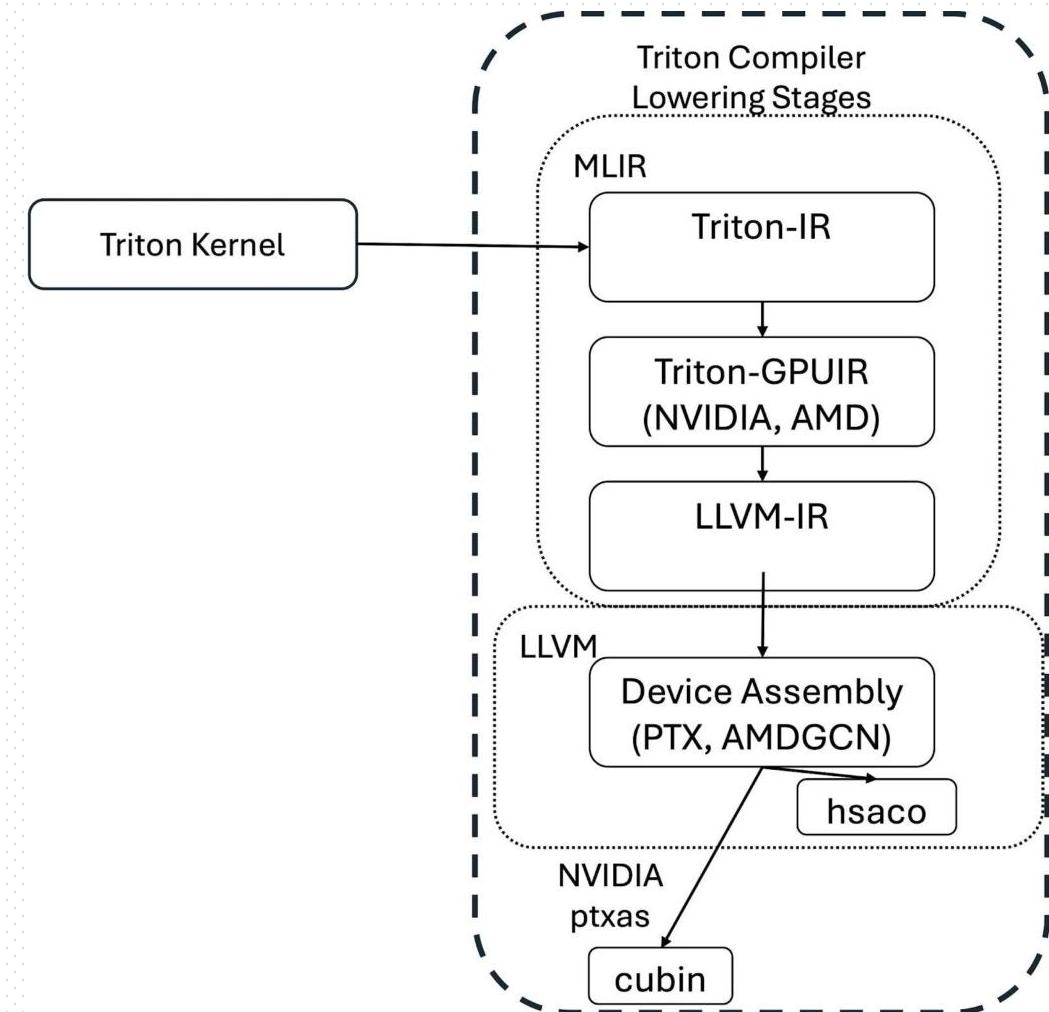
- ◆ AI Algorithms are developed in high-level **Python**.
- ◆ Performance-critical optimizations require low-level **CUDA/C++**.





# Background: What is Triton?

- ◆ A Python-based language and compiler for writing high-performance GPU kernels.
- ◆ It solves the problem for a **single GPU**; Triton-distributed extends this to **distributed systems**.



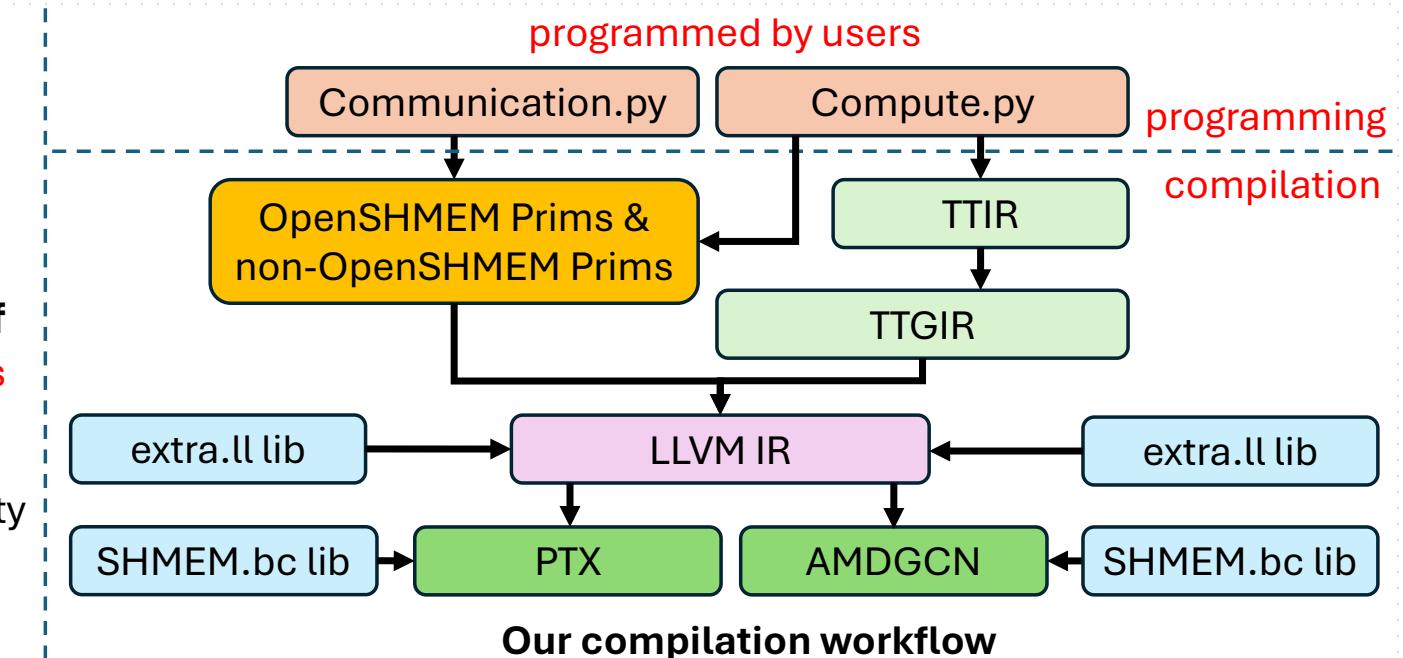
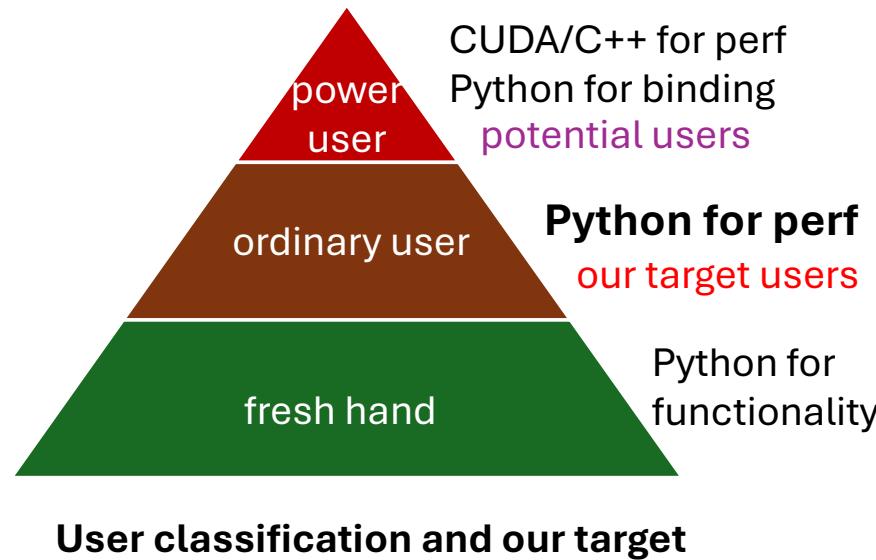


# Outline

- ◆ Background
- ◆ **The Triton-distributed Architecture & Programming Model**
- ◆ Overlapping Optimizations in Triton-distributed
- ◆ Experiments & Evaluations
- ◆ Conclusion



# Triton-distributed Architecture



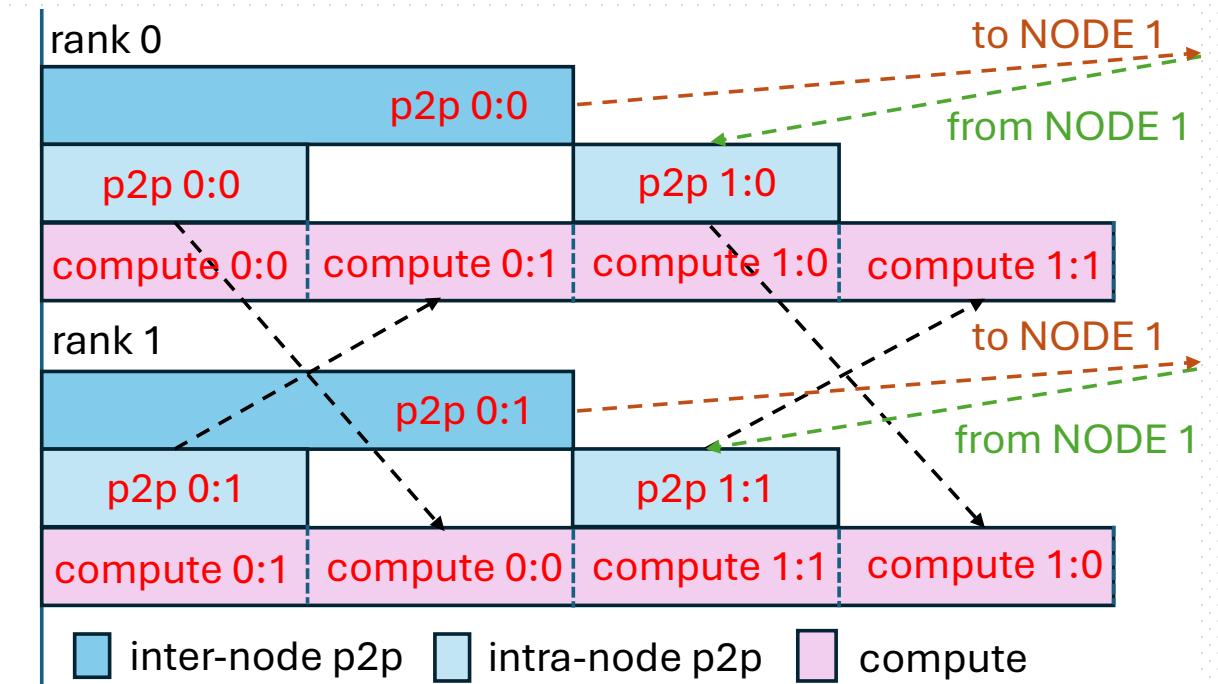
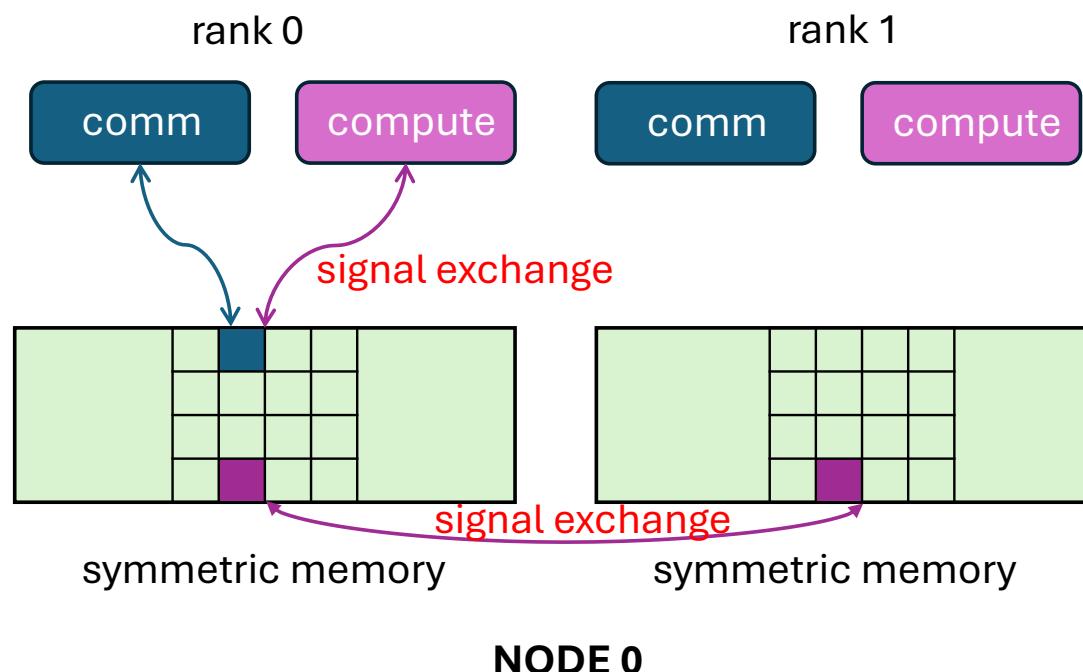


# Triton-distributed Programming Model

## ◆ Symmetric Memory

## ◆ Signal Exchange

## ◆ Async-Task





# Communication Primitives of Triton-distributed

Primitive	Explanation
<b>OpenSHMEM Primitives</b>	
<i>my_pe</i>	Get the current device id
<i>n_pes</i>	The number of devices in the world
<i>int_p</i>	Put an integer to remote device
<i>remote_ptr</i>	Convert local shared memory pointer to remote pointer
<i>barrier_all</i>	Barrier all the devices
<i>sync_all</i>	Synchronize all the devices
<i>quiet</i>	Ensure completion of shared memory operation of calling device
<i>fence</i>	Ensure order of shared memory operation of calling device
<i>getmem</i>	Blocking get data from remote device
<i>getmem_nbi</i>	Non-blocking get data from remote device
<i>putmem</i>	Blocking put data to remote device
<i>putmem_nbi</i>	Non-blocking put data to remote device
<i>putmem_signal</i>	Blocking put data and write signal to remote device
<i>putmem_signal_nbi</i>	Non-blocking put data and write signal to remote device
<i>signal_op</i>	Perform signal set/add operation to remote
<i>signal_wait_until</i>	Wait local signal until condition is meet
<i>broadcast</i>	Broadcast data into all the other ranks



# Communication Primitives of Triton-distributed

Primitive	Explanation
<b>non-OpenSHMEM Primitives</b>	
<i>wait</i>	Locally wait a signal until it equals to some given value
<i>consume_token</i>	used with <i>wait</i> primitive to create data dependency
<i>notify</i>	Notify a remote signal, similar to <i>signal_op</i> primitive
<i>atomic_cas</i>	Atomic compare and swap
<i>atomic_add</i>	Atomic add value
<i>ld_acquire</i>	Load with acquire semantic
<i>red_release</i>	Reduction add with release semantic
<i>multimem_ld_reduce</i>	Multimem load data and perform reduction
<i>multimem_st</i>	Multimem broadcast data



# Example: Inter-node Overlapping AllGather GEMM

Intra-node

```
@triton.jit
def producer_allgather(
    A, signal_num_elem_per_rank,
    rank, local_world_size, world_size):
    pid = tl.program_id(0)
    node = rank // local_world_size
    local_rank = rank % local_world_size
    n_nodes = world_size // local_world_size

    if pid < local_world_size - 1:
        peer = (local_rank + pid + 1) % local_world_size \
            + node * local_world_size
        for i in range(n_nodes):
            seg = (local_rank +
                   ((node + i) % n_nodes) * local_world_size)
            if tid(0) == 0:
                signal_wait_until(signal + seg, EQ, 1)
            __syncthreads()
            putmem_signal(
                A + seg * num_elem_per_rank,
                A + seg * num_elem_per_rank,
                signal + seg,
                1, SET, peer)
```

```
else:
    pid = pid - local_world_size + 1
    if tid(0) == 0:
        signal_wait_until(signal + rank, EQ, 1)
    __syncthreads()
    peer = (local_rank + (node + pid + 1) % n_nodes \
        * local_world_size)
    putmem_signal(
        A + rank * num_elem_per_rank,
        A + rank * num_elem_per_rank,
        signal + rank,
        1, SET, peer)
```

Inter-node

```
@triton.jit
def consumer_gemm(A, B, C, signal):
    pid = tl.program_id(0)
    pid_m, pid_n = ...
    offs_A, offs_B, offs_C = ...
    acc = tl.zeros([TILE_M, TILE_N])
    for k in range(K // TILE_K):
        token = wait(
            signal + pid_m, 1, "gpu", "acquire", waitValue=1)
        a_ptrs = consume_token(A + offs_A, token)
        a_data = tl.load(a_ptrs)
        b_data = tl.load(b_ptrs)
        tl.dot(a_data, b_data, acc)
        offs_A, offs_B = ...
        tl.store(C + offs_C, acc)
```

```
def ag_gemm(A, B, C, signal):
    with comm_stream():
        grid = (local_world_size + n_nodes - 2, 1, 1)
        producer_allgather[grid](
            A, signal, num_elem_per_rank,
            rank, local_world_size, world_size)
    with compute_stream():
        grid = ((M//TILE_M) * (N//TILE_N), 1, 1)
        consumer_gemm[grid](A, B, C, signal)

A = create_tensor([global_M, K])
signal = create_tensor([world_size])
```



# Outline

- ◆ Background
- ◆ The Triton-distributed Architecture & Programming Model
- ◆ Overlapping Optimizations in Triton-distributed
- ◆ Experiments & Evaluations
- ◆ Conclusion



# Optimization Approaches and Comparison with Other Frameworks

Name	NCCL	PyTorch	TE	Pallas	CoCoNet	FLUX	DeepEP	Ours
Intra-Node Swizzle	■	✓	✓	■	✓	✓	■	✓
Inter-Node Swizzle	■	✗	✗	■	✓	✓	■	✓
Inter-NUMA Swizzle	■	✗	✗	✗	✗	✗	■	✓
Copy Engine	✓	✓	✓	✓	✓	✓	✗	✓
High-BW Link	✓	✓	✓	✓	✓	✓	✓	✓
Network Comm.	✓	✓	✗	✓	✓	✓	✓	✓
PCIe Comm.	✓	✓	✗	✗	✗	✓	✗	✓
OpenSHMEM Support	✗	✗	✗	✗	✗	✓	✓	✓
Low-latency Protocol	✓	✗	✗	✗	✗	✓	✗	✓
Multimem Feature	■	✗	✗	✗	✗	✗	✗	✓
Fusion	✗	✗	✗	■	✓	✓	✓	✓
Code Generation	✗	✗	✗	✓	✓	✗	✗	✓
Nvidia/AMD	✓/✗	✓/✓	✓/✗	■/✗	✓/✗	✓/✗	✓/✗	✓/✓



# Communication Kernels (1/3): Intra-Node AllGather

- ◆ Primarily utilizes the dedicated Copy Engine to offload data transfer from compute cores.
- ◆ Offers two implementation modes:
  - ◆ **Push Mode (Algo 1):** Sender-initiated. Lower sync overhead, but uncontrolled arrival order.
  - ◆ **Pull Mode (Algo 2):** Receiver-initiated. Controlled order, but requires an extra barrier synchronization.

---

## Algorithm 1 One-sided Push-mode Intra-node AllGather

---

```

1: Input: Symmetric Buffer  $T$ , Signal  $S$ , Local Buffer  $L$ 
2: for  $r$  in range(WORLD_SIZE) do
3:   remote_buf = make_buffer(remote_ptr( $T, r$ ) + RANK  $\times L.size()$ )
4:   remote_buf.copy_( $L, L.size()$ ) // Memory Copy
5:   remote_sig = remote_ptr( $S, r$ ) + RANK
6:   set_signal(remote_sig) // Notify the consumer
7: end for

```

---



---

## Algorithm 2 One-sided Pull-mode Intra-node AllGather

---

```

1: Input: Symmetric Buffer  $T$ , Signal  $S$ , Local Buffer  $L$ 
2: local_t_buf = make_buffer( $T + RANK \times L.size()$ )
3: local_t_buf.copy_( $L, L.size()$ )
4: set_signal( $S + RANK$ )
5: barrier_all() // Make the local copy visible to all the other ranks
6: for  $r$  in range(WORLD_SIZE) do
7:   if  $r$  is not RANK then
8:     remote_buf = make_buffer(remote_ptr( $T, r$ ) +  $r \times L.size()$ )
9:     local_t_buf = make_buffer( $T + r \times L.size()$ )
10:    local_t_buf.copy_(remote_buf,  $L.size()$ )
11:    set_signal( $S + r$ )
12:   end if
13: end for

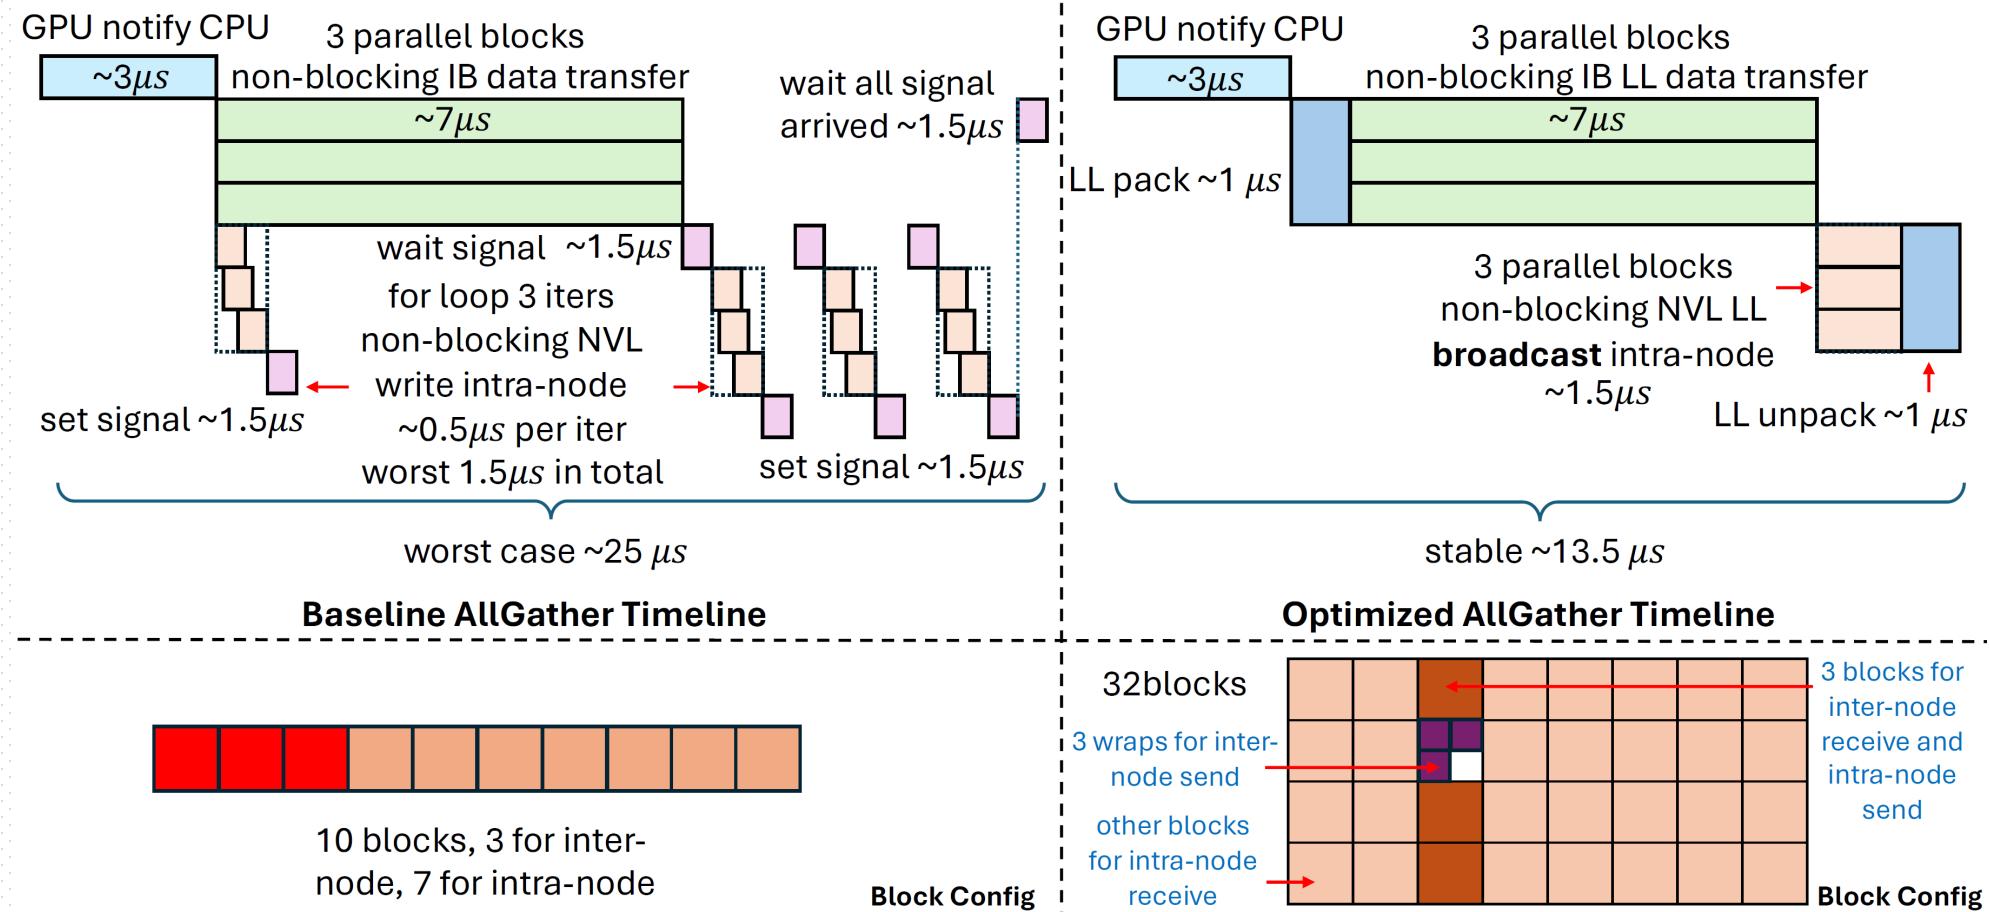
```

---



# Communication Kernels (2/3): Low-Latency Inter-Node AllGather

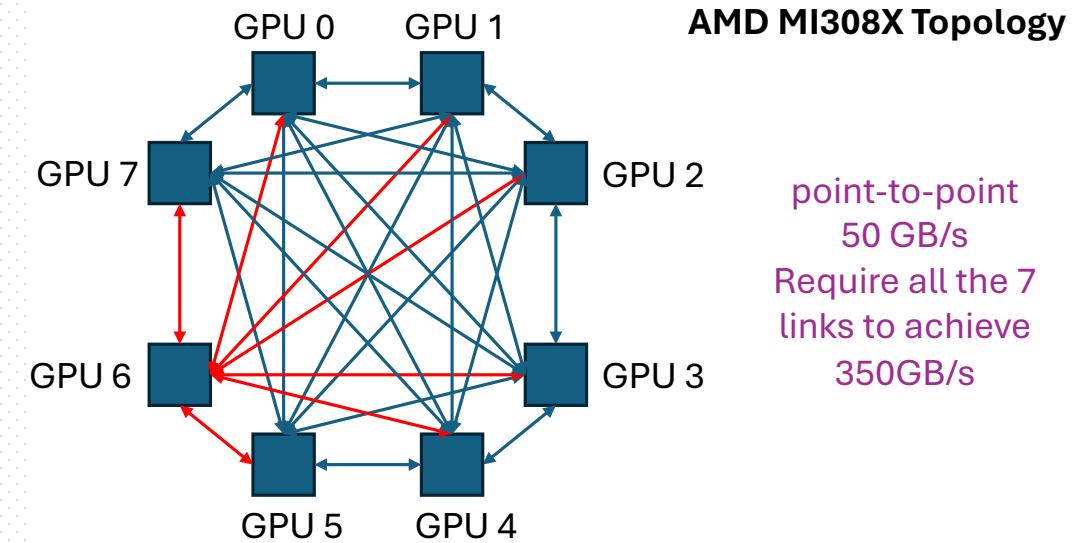
◆ **Problem:** Baseline implementations can suffer from "skew," turning parallel sends into sequential ones and increasing latency.





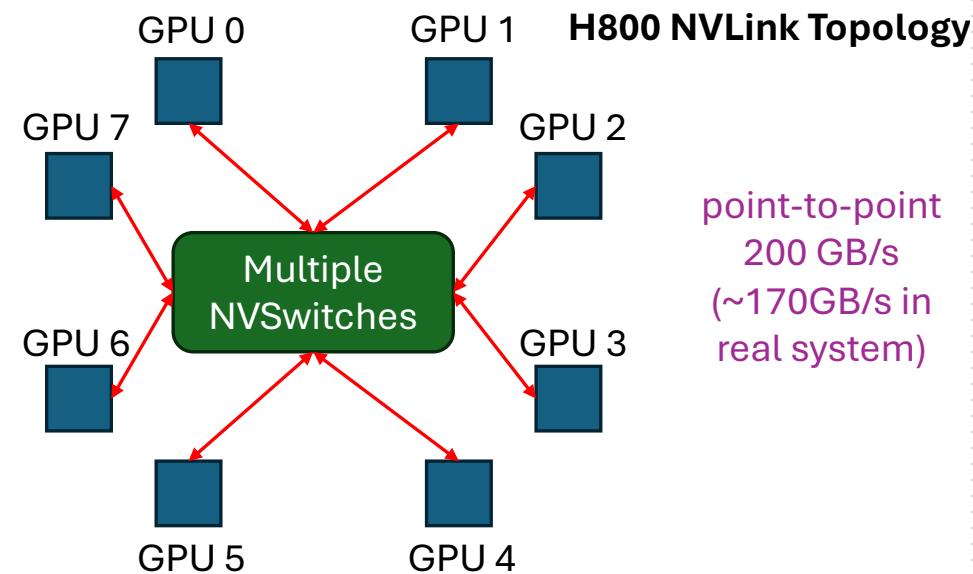
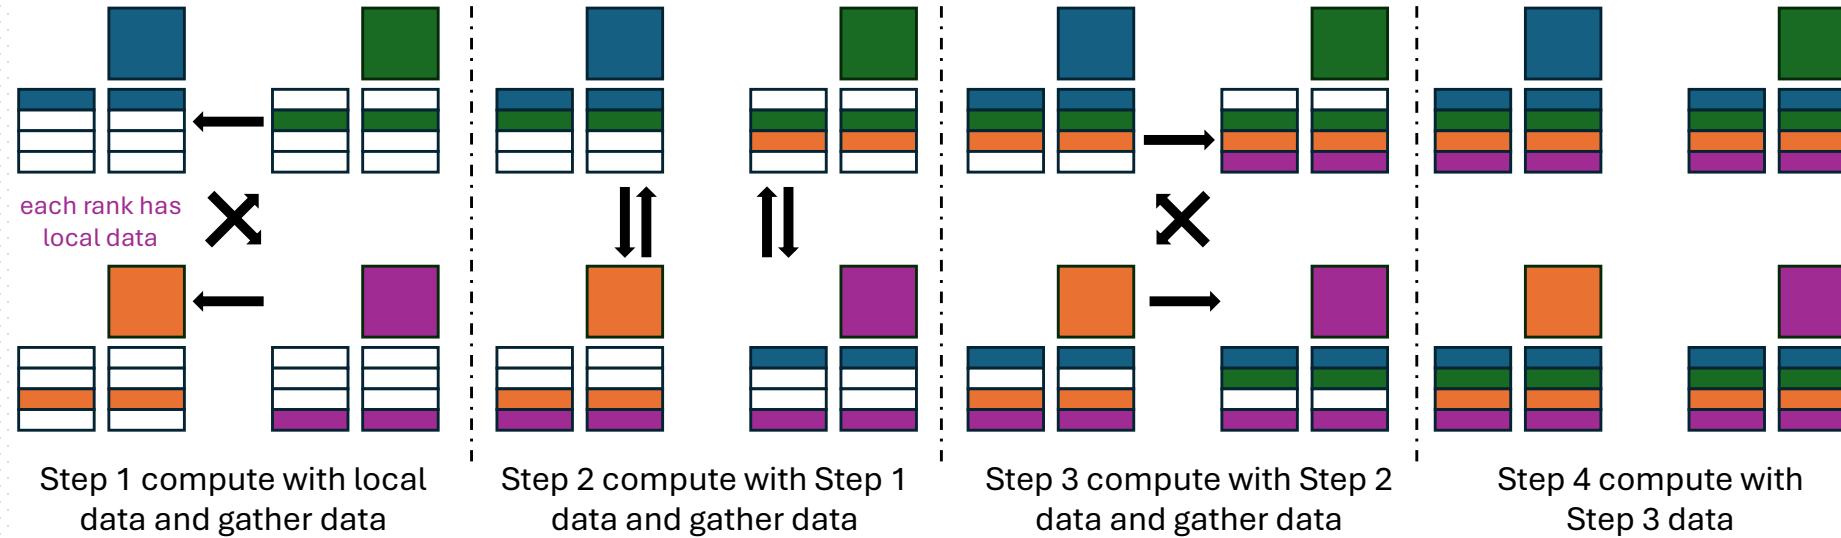
## Communication Kernels (3/3): Platform Adaptation (AMD)

- ◆ The framework adapts to different hardware topologies and behaviors.
- ◆ **On AMD MI308X:**
  - ◆ Requires launching transfers on multiple streams simultaneously to maximize bandwidth on its full-mesh topology.
  - ◆ Works around problematic driver APIs by fusing the scatter operation directly into the producer compute kernel, avoiding the API call.



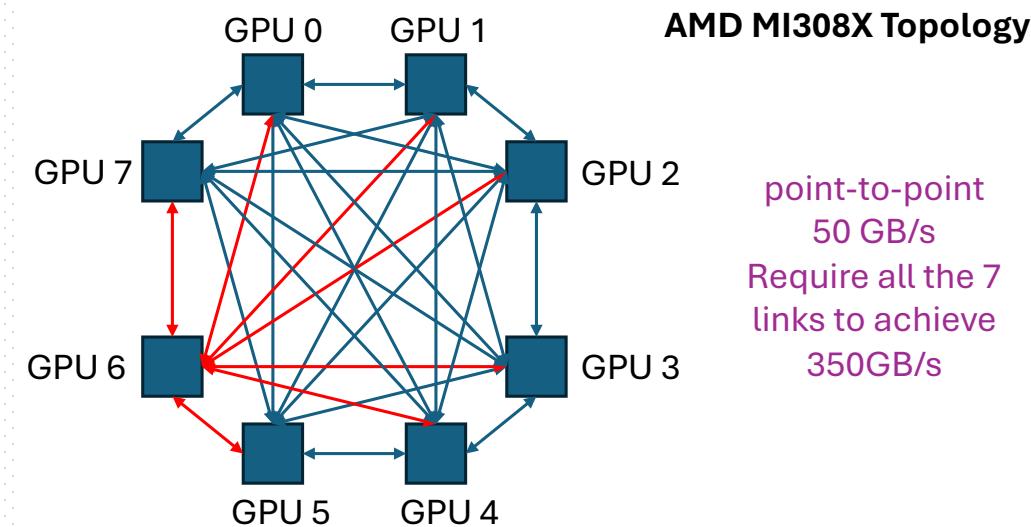
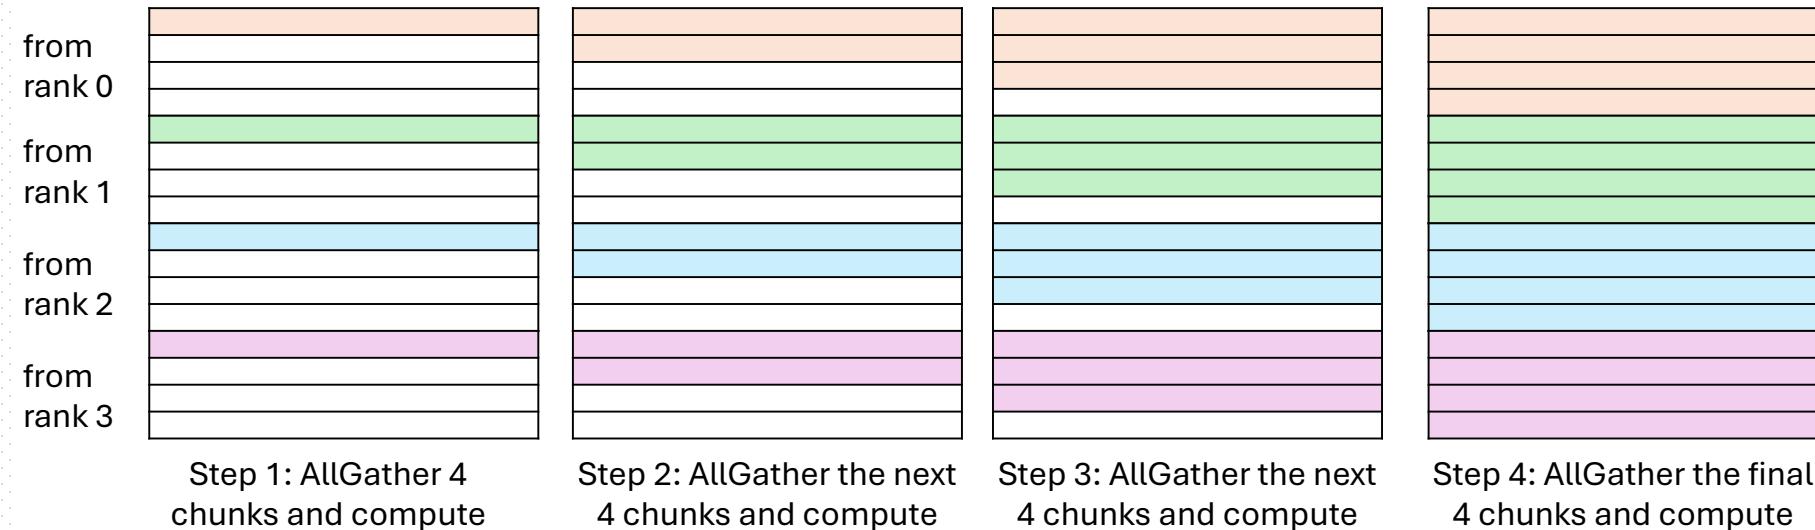


# Overlapping Computation with Swizzling Optimization





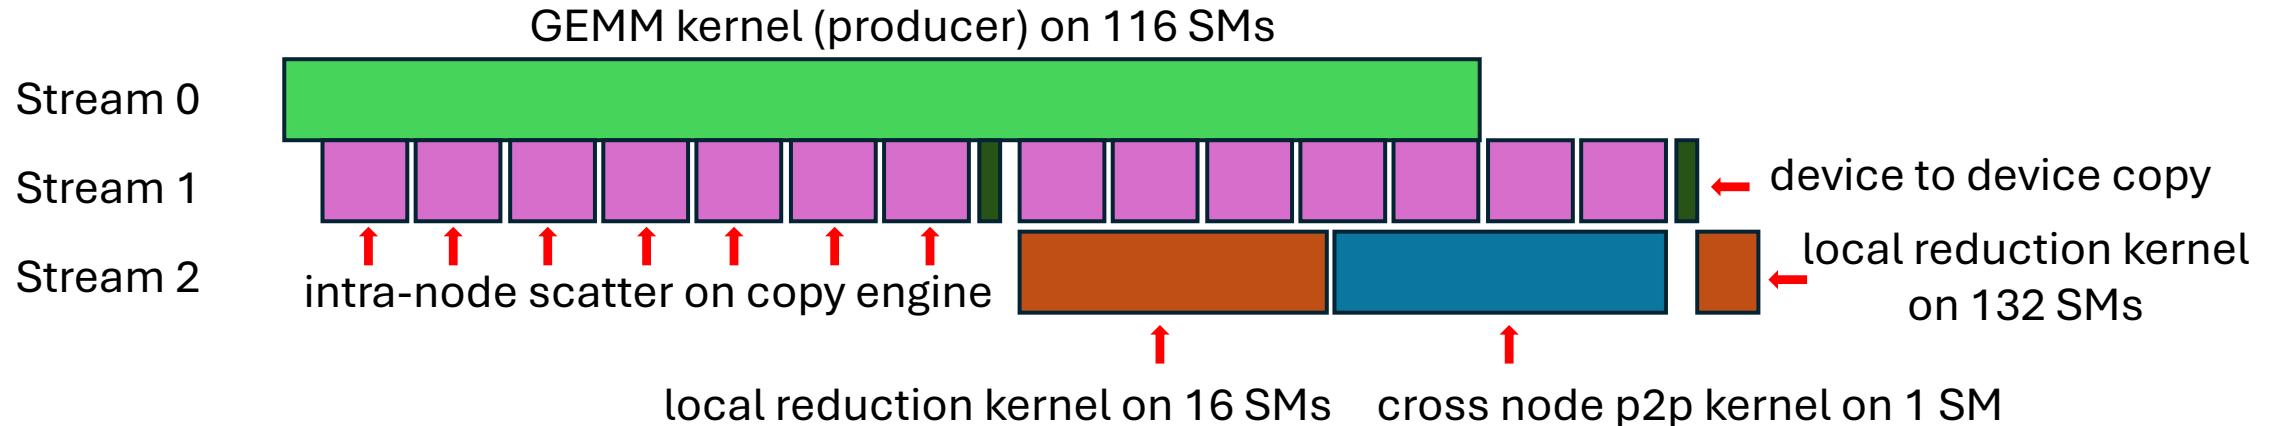
# Overlapping Computation with Swizzling Optimization





# Distributed Auto-Tuning and Resource Partitioning

- ◆ **Distributed Auto-Tuning:** A novel auto-tuner designed specifically for distributed, overlapping kernels.
- ◆ **Resource Partitioning:** A spatial optimization that maps tasks to different hardware units to balance load and prevent bottlenecks.





# Outline

- ◆ Background
- ◆ The Triton-distributed Architecture & Programming Model
- ◆ Overlapping Optimizations in Triton-distributed
- ◆ Experiments & Evaluations
- ◆ Conclusion

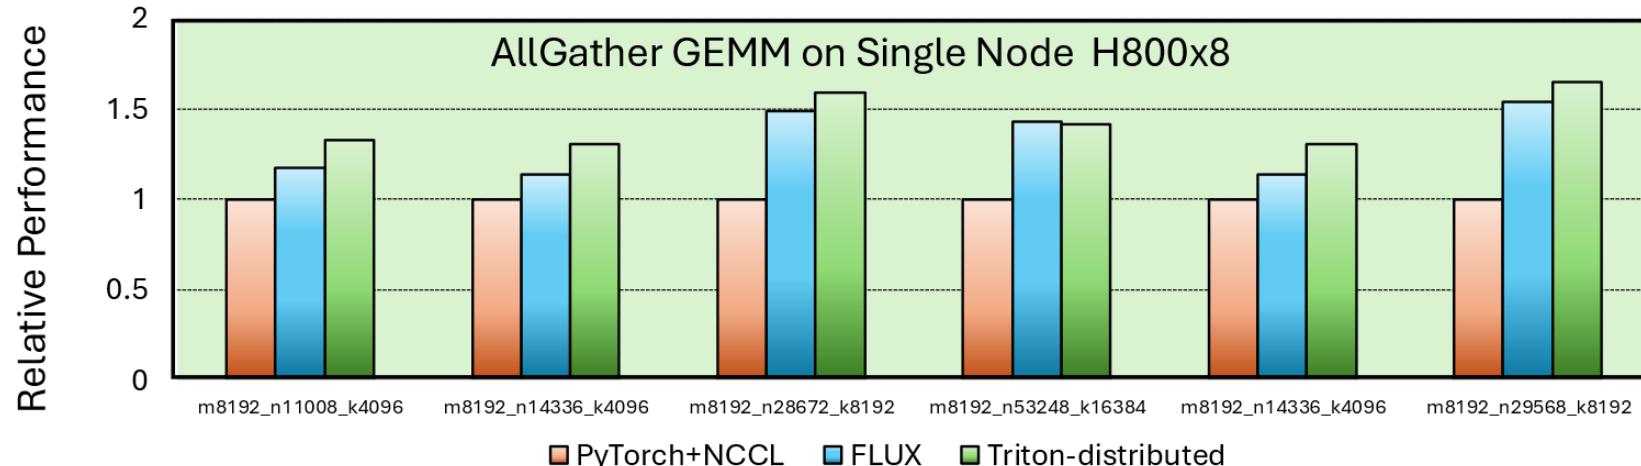


# List of Optimized Kernels

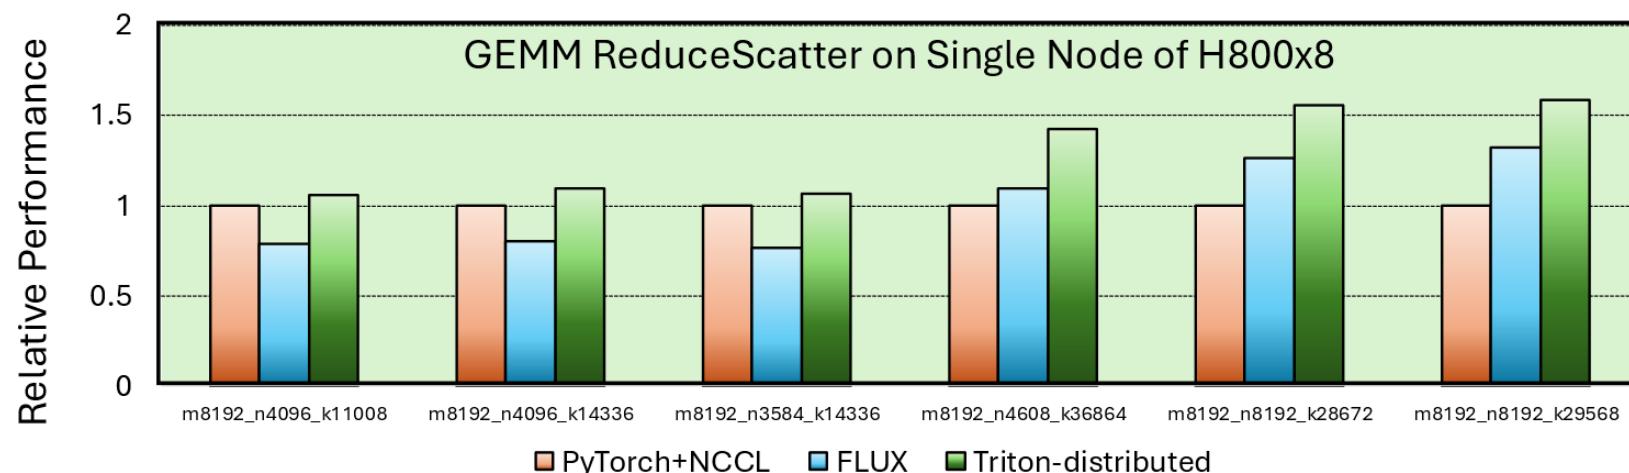
Name	Explanation	Tested Hardware Cluster
AG+GEMM-intra	Intra-node AllGather GEMM Overlapping	8 H800 and MI308X GPUs.
GEMM+RS-intra	Intra-node GEMM ReduceScatter Overlapping	8 H800 and MI308X GPUs.
AG+MoE-intra	Intra-node AllGather MoE GroupGEMM Overlapping	8 H800 GPUs
MoE+RS-intra	Intra-node MoE GroupGEMM ReduceScatter Overlapping	8 H800 GPUs
FlashDecode+AG-intra	Intra-node Flash Decode AllGather and Combine	8 H800 GPUs
AllToAll-intra	Intra-node Low-latency AllToAll	8 H800 GPUs
AG+GEMM-inter	Inter-node AllGather GEMM Overlapping	16 H800 GPUs
GEMM+RS-inter	Inter-node GEMM ReduceScatter Overlapping	16 H800 GPUs
AG+MoE-inter	Inter-node AllGather MoE GroupGEMM Overlapping	16 H800 GPUs
MoE+RS-inter	Inter-node MoE GroupGEMM ReduceScatter	16 H800 GPUs
FlashDecode+AG-inter	Inter-node Flash Decode AllGather and Combine	16 and 32 H800 GPUs
AllToAll-inter	Inter-node Low-latency AllToAll	16, 32, and 64 H800 GPUs



# Intra-node Kernel Performance on Nvidia GPUs



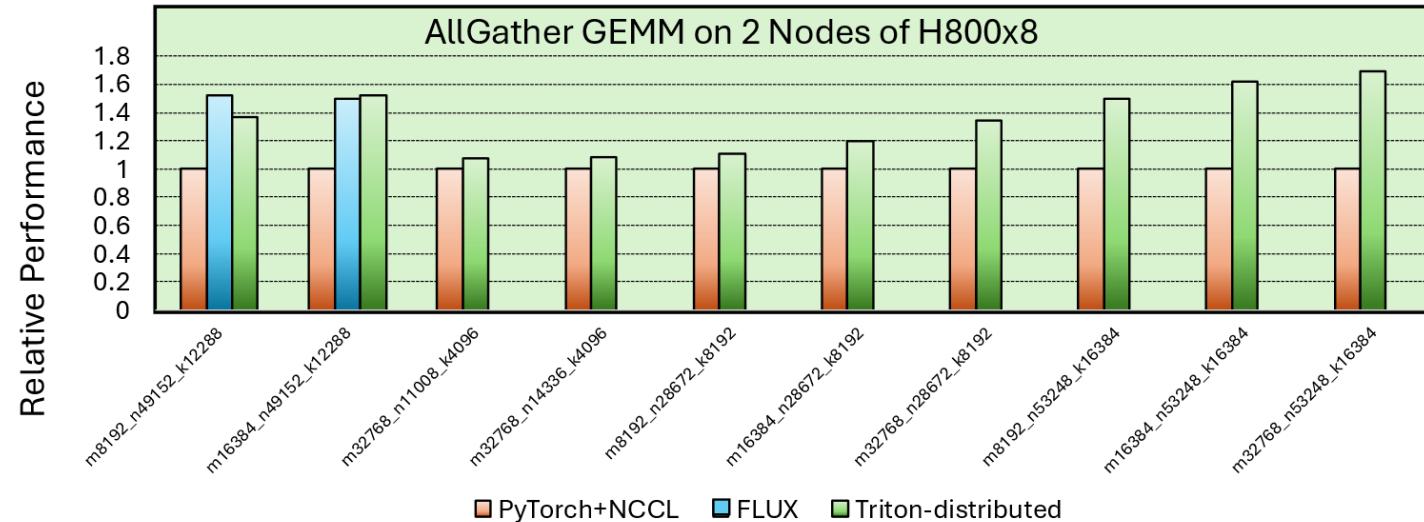
Performance of Intra-node AllGather GEMM on 8 H800 GPUs.



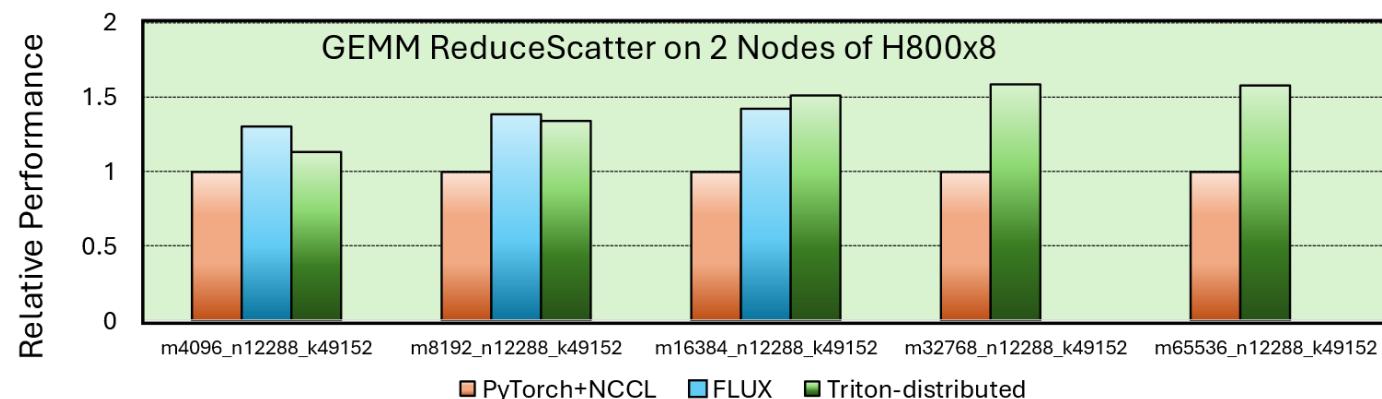
Performance of Intra-node GEMM ReduceScatter on 8 H800 GPUs.



# Inter-node Kernel Performance on Nvidia GPUs



Performance of Inter-node AllGather GEMM on 16 H800 GPUs.



Performance of Inter-node GEMM ReduceScatter on 16 H800 GPUs.



# MoE Performance on Nvidia GPUs

Name	tokens/rank	in hidden	out hidden	experts	topk	Ours		PyTorch	
						Intra	Inter	Intra	Inter
AG+MoE-1	256	2048	1408	60	4	0.33	0.45	23.95	28.84
AG+MoE-2	512	2048	1408	60	4	0.40	1.37	26.25	29.77
AG+MoE-3	1024	2048	1408	60	4	0.58	1.80	30.42	43.31
AG+MoE-4	2048	2048	1408	60	4	0.97	3.07	55.63	63.73
AG+MoE-5	256	14336	4096	8	2	0.54	1.01	7.05	19.92
AG+MoE-6	512	14336	4096	8	2	0.72	1.89	26.34	36.07
AG+MoE-7	1024	14336	4096	8	2	1.19	3.41	52.99	67.61
AG+MoE-8	2048	14336	4096	8	2	2.10	6.51	107.32	129.30
AG+MoE-9	256	16384	6144	8	2	0.81	1.39	11.02	27.29
AG+MoE-10	512	16384	6144	8	2	1.06	2.21	39.65	52.32
AG+MoE-11	1024	16384	6144	8	2	1.66	4.32	80.46	101.61
AG+MoE-12	2048	16384	6144	8	2	2.92	8.28	159.69	192.67
AG+MoE-13	512	1408	2048	64	6	0.45	0.84	29.25	38.17
AG+MoE-14	1024	1408	2048	64	6	0.67	1.26	48.86	56.77
AG+MoE-15	2048	1408	2048	64	6	1.18	2.18	74.26	90.44

Test Shapes for AllGather MoE and Performance (ms).



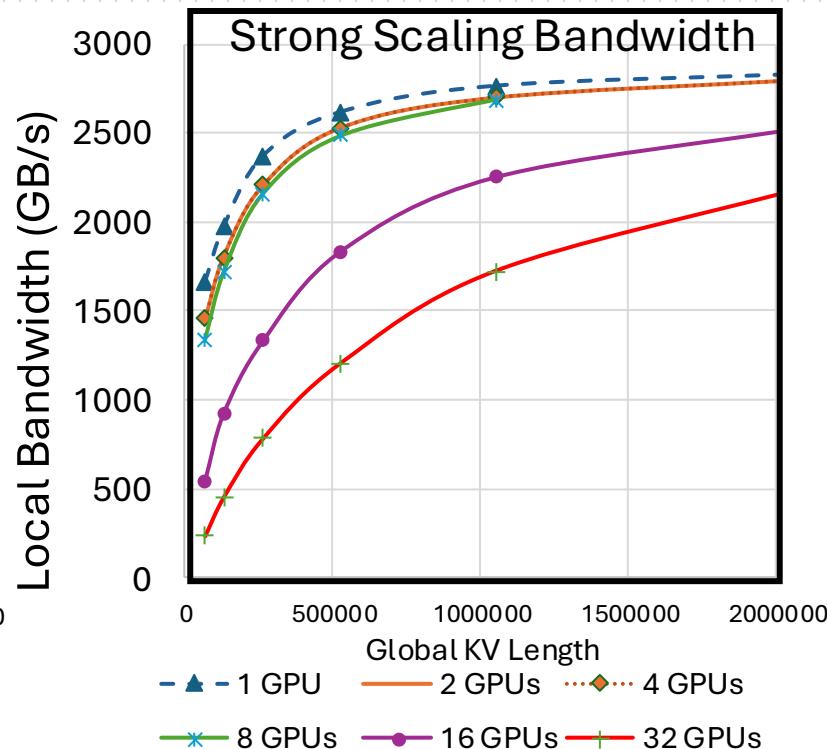
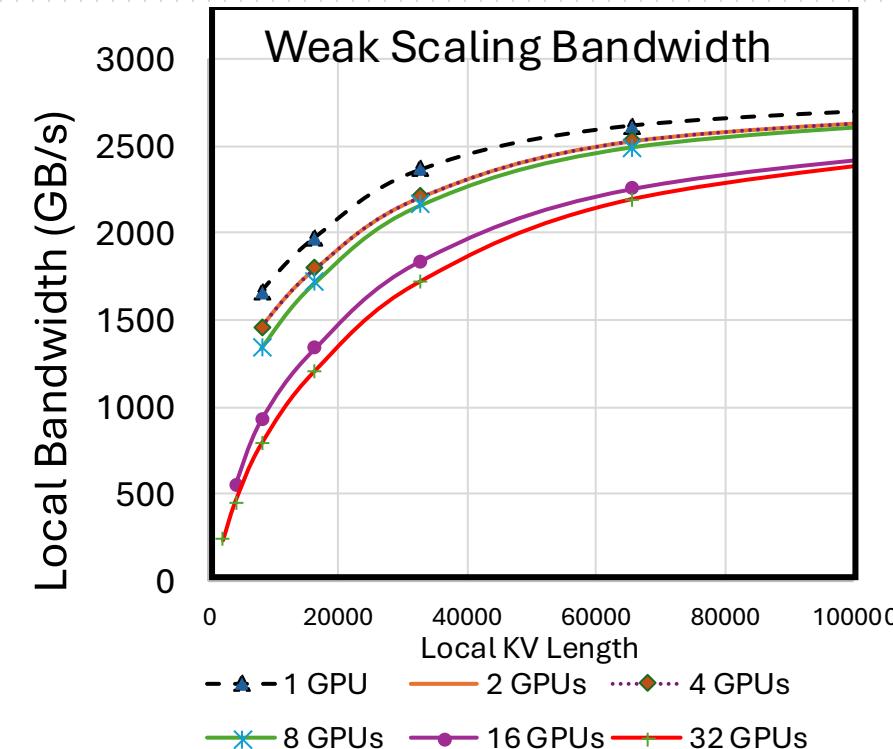
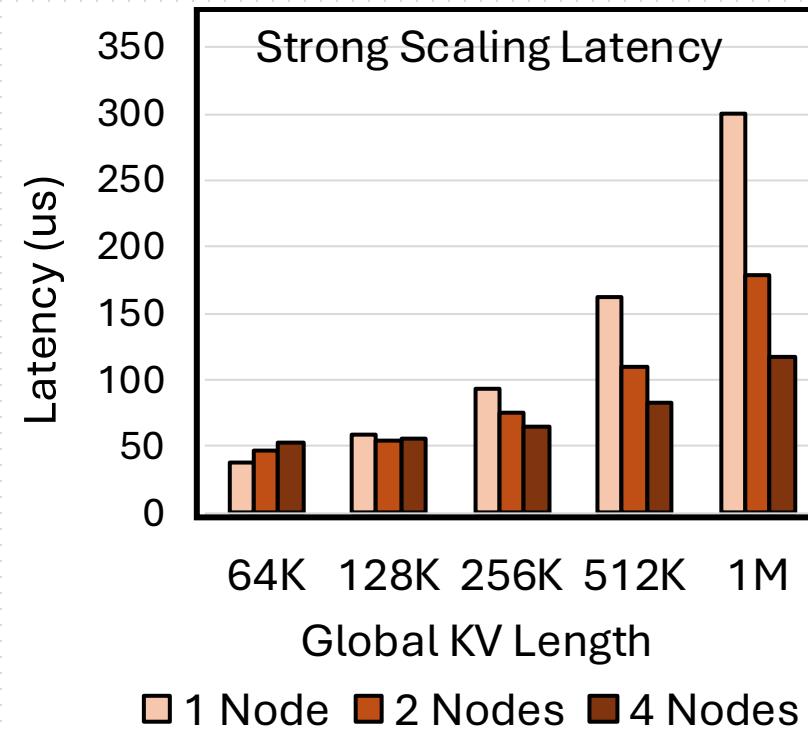
# MoE Performance on Nvidia GPUs

Name	tokens/rank	in hidden	out hidden	experts	topk	Ours		PyTorch	
						Intra	Inter	Intra	Inter
MoE-RS-1	1024	1536	2048	8	2	0.51	3.62	4.35	12.41
MoE-RS-2	1024	1536	2048	32	2	0.55	3.90	13.89	33.05
MoE-RS-3	1024	1536	2048	64	2	0.67	4.82	27.91	61.70
MoE-RS-4	1024	1536	2048	32	5	0.92	7.78	14.48	35.35
MoE-RS-5	1024	1536	2048	64	5	0.93	8.25	29.96	64.88
MoE-RS-6	1024	2048	4096	8	2	0.98	7.00	5.02	17.93
MoE-RS-7	1024	2048	4096	32	2	1.08	7.86	14.12	38.24
MoE-RS-8	1024	2048	4096	64	2	1.34	9.87	28.61	66.48
MoE-RS-9	1024	2048	4096	32	5	1.84	15.51	16.70	44.37
MoE-RS-10	1024	2048	4096	64	5	1.86	16.60	27.71	71.82

Test Shapes for MoE ReduceScatter and Performance (ms).



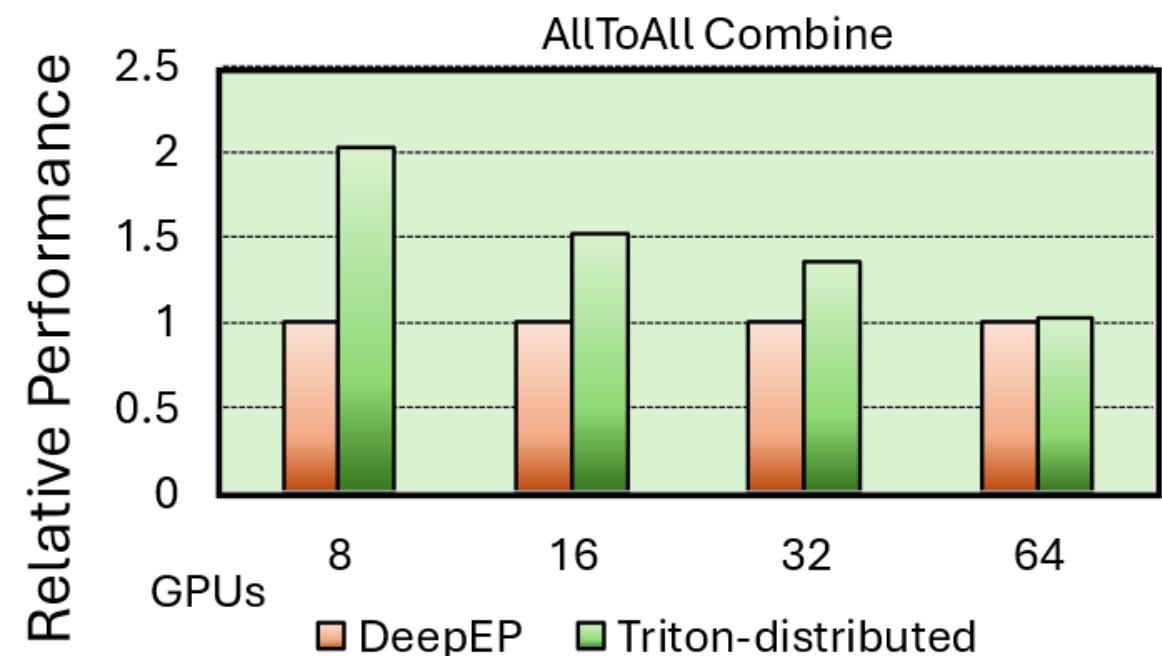
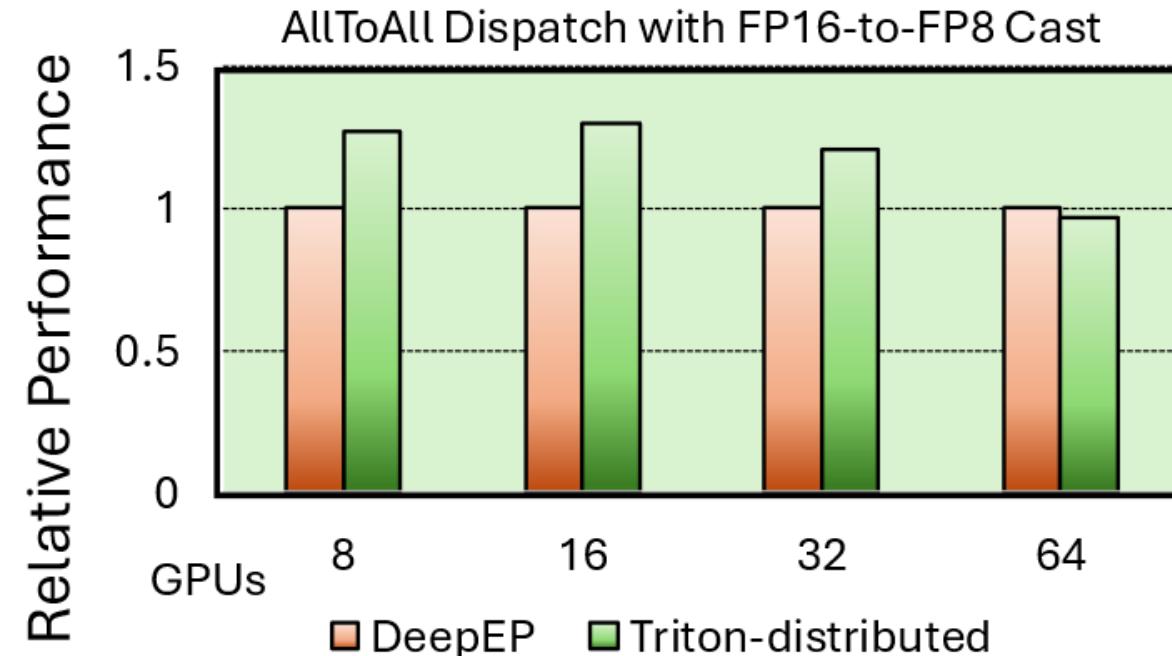
# Distributed Flash Decoding Performance



# Performance of Distributed Flash Decoding.



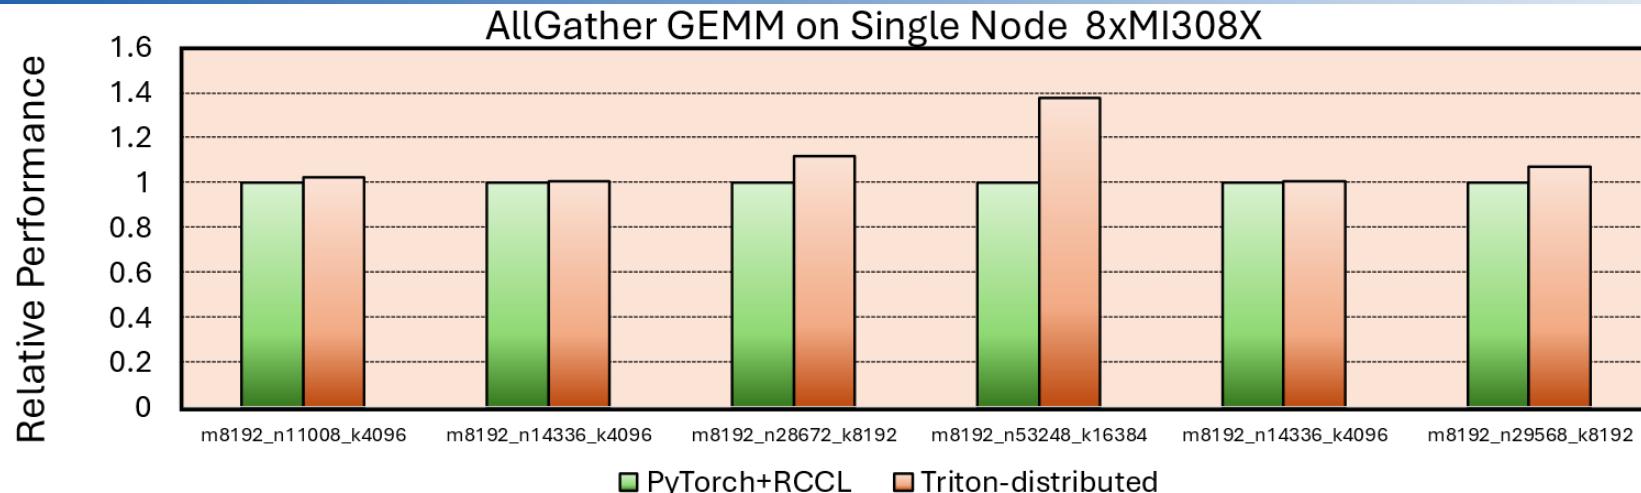
# Low Latency AllToAll Performance



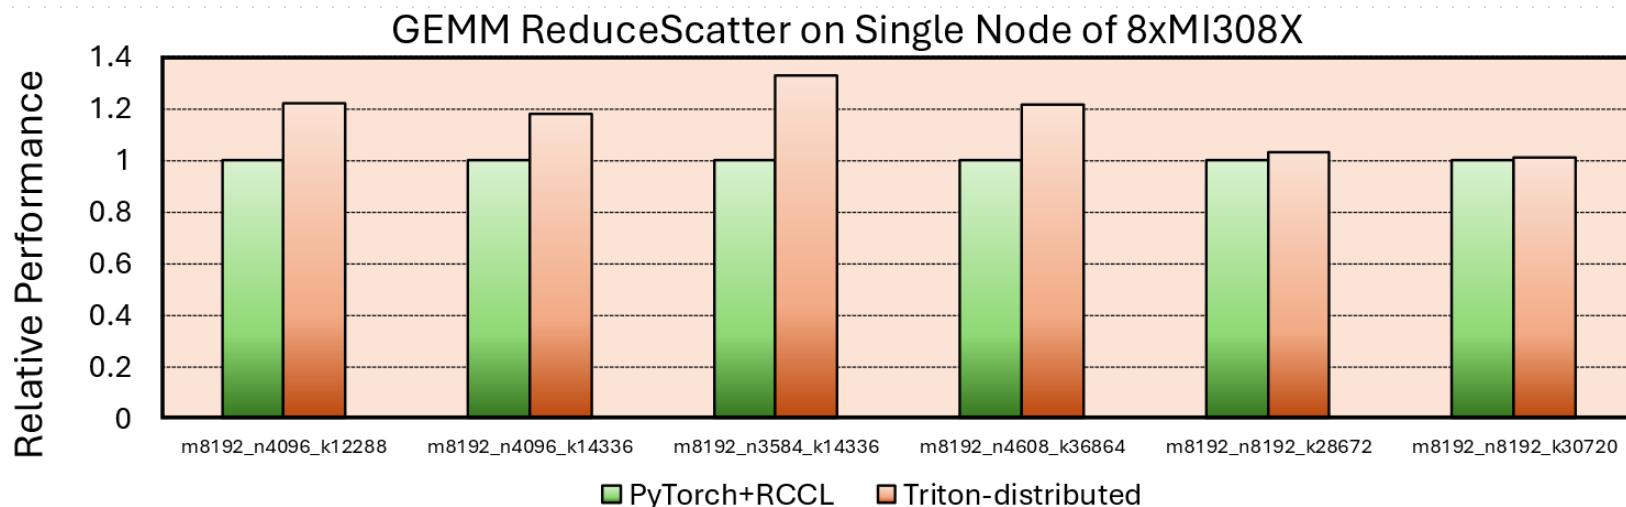
Performance of AllToAll



# Intra-node Kernel Performance on AMD GPUs



Performance of Intra-node AllGather GEMM on AMD GPUs.



Performance of Intra-node GEMM ReduceScatter on AMD GPUs.



# Outline

- ◆ **Background**
- ◆ **The Triton-distributed Architecture & Programming Model**
- ◆ **Overlapping Optimizations in Triton-distributed**
- ◆ **Experiments & Evaluations**
- ◆ **Conclusion**



# Conclusion

- ◆ Triton-distributed successfully **unifies distributed programming into Python**, drastically lowering the development barrier.
- ◆ The generated code achieves performance that is **competitive with, or superior to, hand-optimized low-level code**.
- ◆ The methodology is **portable across different hardware platforms**, demonstrating its general applicability.

# Thanks

---