

FAST-DLLM V2: Efficient Block-Diffusion LLM

Chengyue Wu^{1,2} Hao Zhang² Shuchen Xue² Shizhe Diao²
Yonggan Fu² Zhijian Liu² Pavlo Molchanov² Ping Luo¹
Song Han^{2,3} Enze Xie²

¹ The University of Hong Kong ²NVIDIA ³MIT

arXiv:2509.26328

Presented by Xiliang Xian

Outline

□ Background

□ Design

□ Evaluations

□ Discussion

Auto-Regressive(AR) LLM

Autoregression:



High quality



Arbitrary-length



KV caching



Not Parallelizable

Generation steps

There are three categories of the average
There are three categories of the average rate
There are three categories of the average rate of...

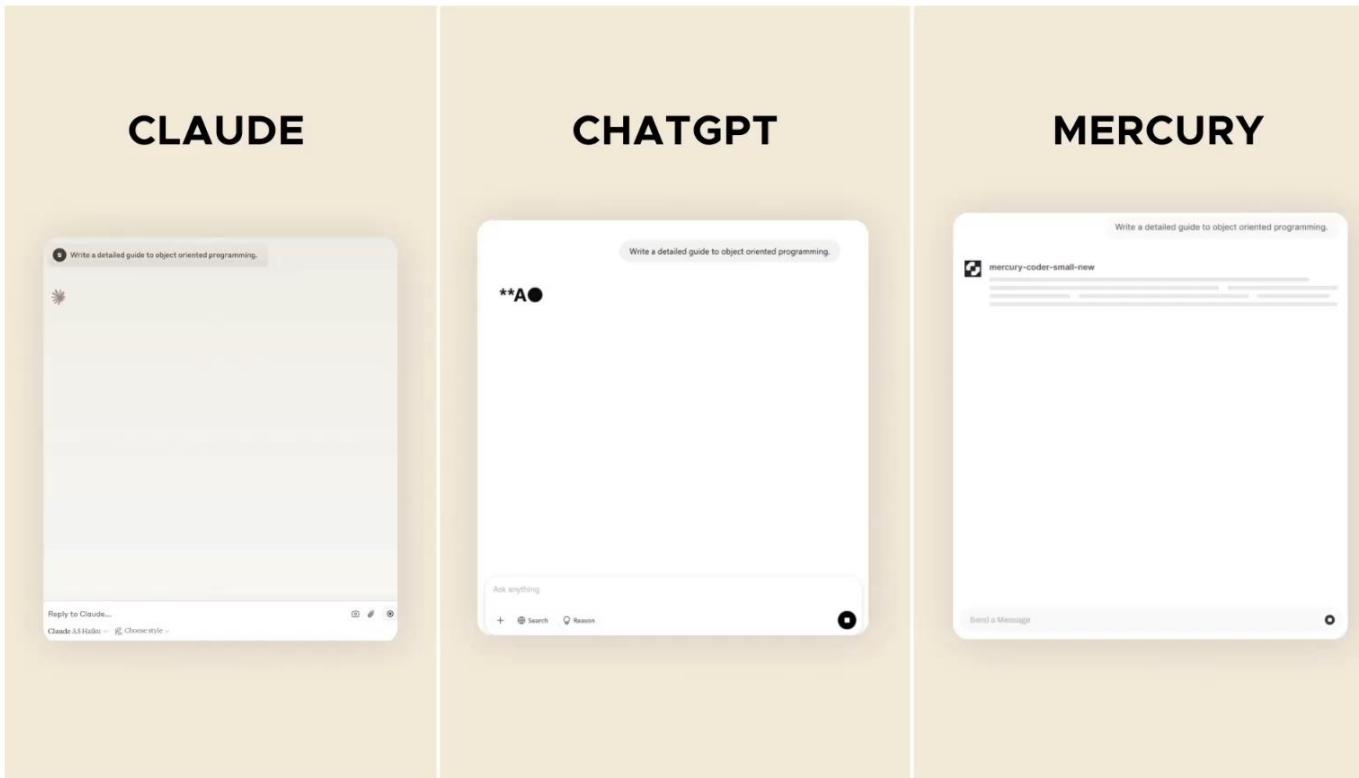
- produces one token at a time in strict left-to-right order

What is DLLM?

- Diffusion large language model

Why DLLM?

- DLLMs show **superior speed** vs Auto-Regressive(AR) LLM with **comparable performance**



A Glimpse into Recent Development

Closed-source model

- ❖ Gemini Diffusion
- ❖ Mercury (claimed to achieve 1109 tokens/s on H100s)
- ❖ Seed Diffusion (claimed to achieve 2146 tokens/s on H20s)

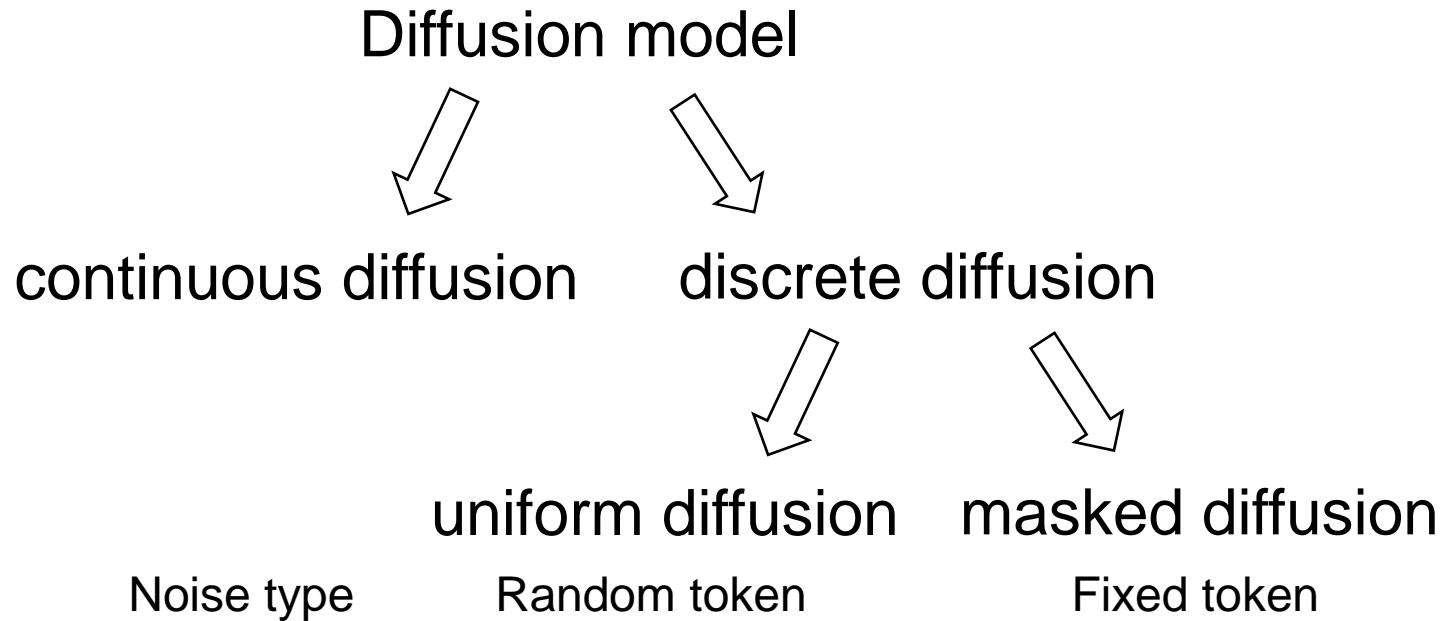
Open-source model

- ❖ LLaDA (train from scratch)
- ❖ Dream (fine-tune from Qwen-2.5 7B)

Introduction

- Model Apply **Block Diffusion**
- **Data Efficient** Fine-tune Training Model
- Introduce Approximate **KV Cache** for **Full Attention**
- Introduce **Tokens** Parallel Decoding

Classification



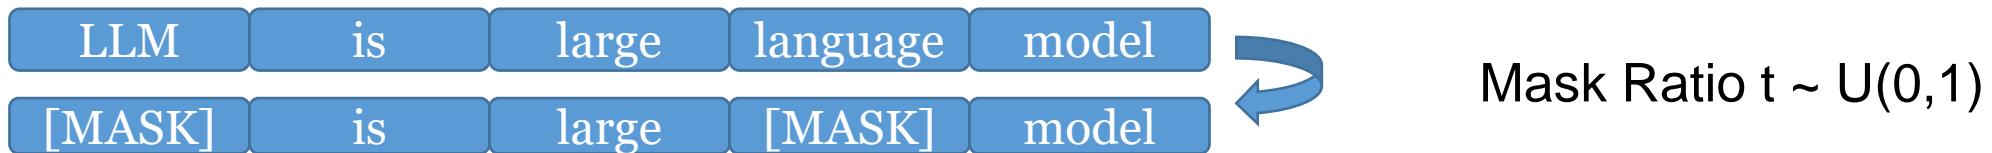
Masked Diffusion Model (MDM)

Mainstream Formulation in DLLM

❑ Training details:

most Encoder-only (full attention)

❖ Forward:

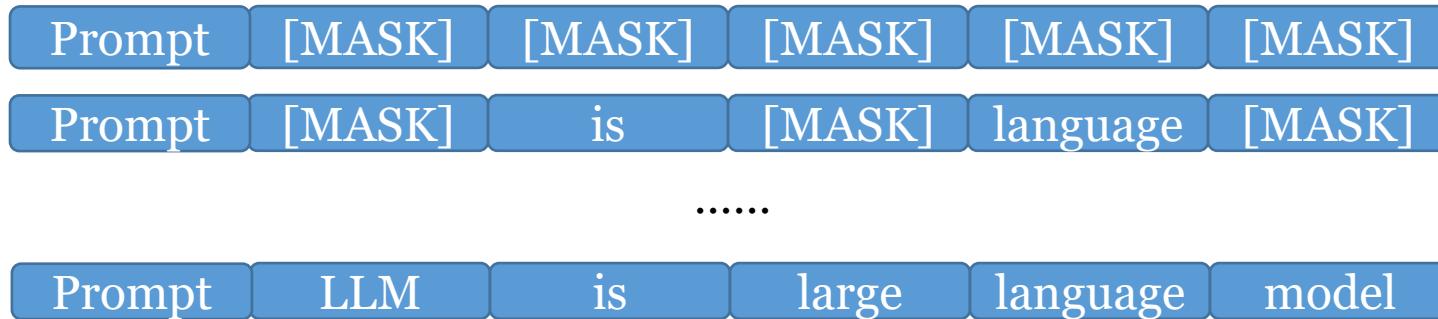


❖ Reverse denoising:



Masked Diffusion Model (MDM)

❑ Inference details:



Multi-steps: predict and remask

Predict all
masked tokens
simutaneously

Random remask,
low confidence
remask, etc

- ❖ Way for variable-length sequence:
 - Discard tokens after first <EOS>(end of sentence)

- ❑ Apply **Block Diffusion**
- ❑ **Data Efficient** Fine-tune Training
- ❑ Introduce Approximate **KV Cache** for Full Attention
- ❑ Introduce **Tokens** Parallel Decoding

Block Diffusion (ICLR 25)

Apply Block Diffusion

Data Efficient Fine-tune Training

Introduce Approximate KV Cache for Full Attention

Introduce Tokens Parallel Decoding

□ Semi-Autoregressive

Autoregression:

✓ High quality ✓ Arbitrary-length ✓ KV caching ✗ Not Parallelizable

Generation steps

There are three categories of the average
There are three categories of the average rate
There are three categories of the average rate of...

Diffusion:

✗ Lower quality ✗ Fixed-length ✗ No KV caching ✓ Parallelizable

the reusability will continue to the
Repeal the reusability cuts and the law will continue to reduce the
Repeal the reusability cuts and prove the law will continue to reduce the deficit.

Block Diffusion (Ours):

✓ High quality ✓ Arbitrary-length ✓ KV caching ✓ Parallelizable

On September 17, we be
On September 17, 2016, we will be giving the release of
On September 17, 2016, we will be giving the beta-release of the to our server testing ...

Limitations in MDMs

❑ Inductive bias conflict in training

- ❖ natural language is overwhelmingly processed in a sequential order

❑ Inference inefficiency

- ❖ take LLaDA as an example:
 - Full attention is slower than causal attention
 - **Cannot leverage KV cache**
 - Reduce the number of inference steps ➡ Decode more tokens simultaneously
➡ **Bad accuracy**

Apply **Block Diffusion**

Data Efficient Fine-tune Training

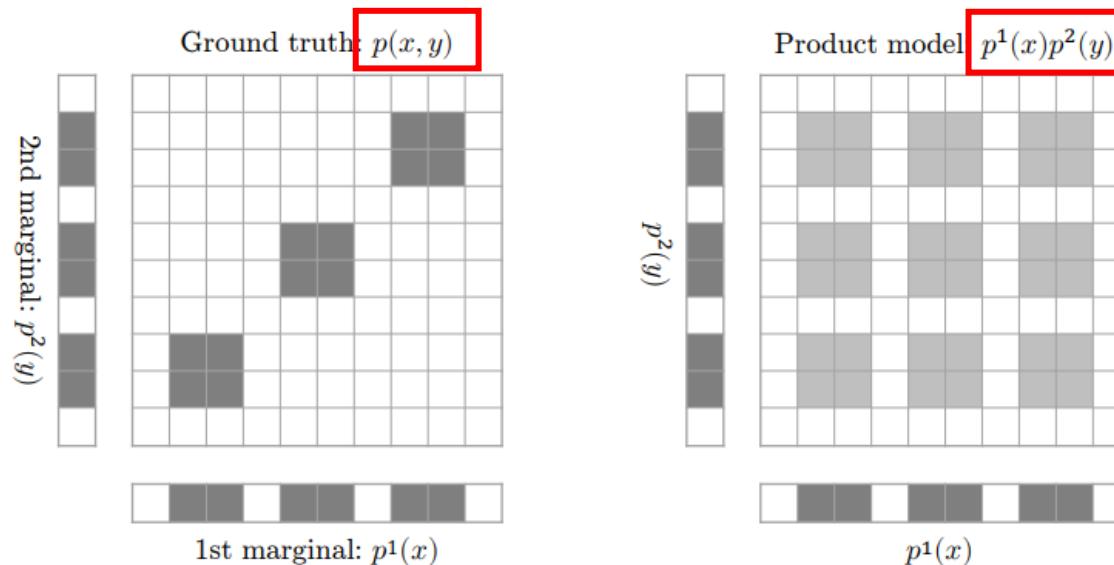
Introduce Approximate **KV Cache** for Full Attention

Introduce **Tokens** Parallel Decoding

Limitations in MDMs

☐ Inference inefficiency

-
- Reduce the number of inference steps → Decode more tokens simultaneously
 - **Bad accuracy**



I X am
he is

marginal distribution \neq joint distribution

- ☐ Apply Block Diffusion
- ☐ Data Efficient Fine-tune Training
- ☐ Introduce Approximate KV Cache for Full Attention
- ☐ **Introduce Tokens Parallel Decoding**

Outline

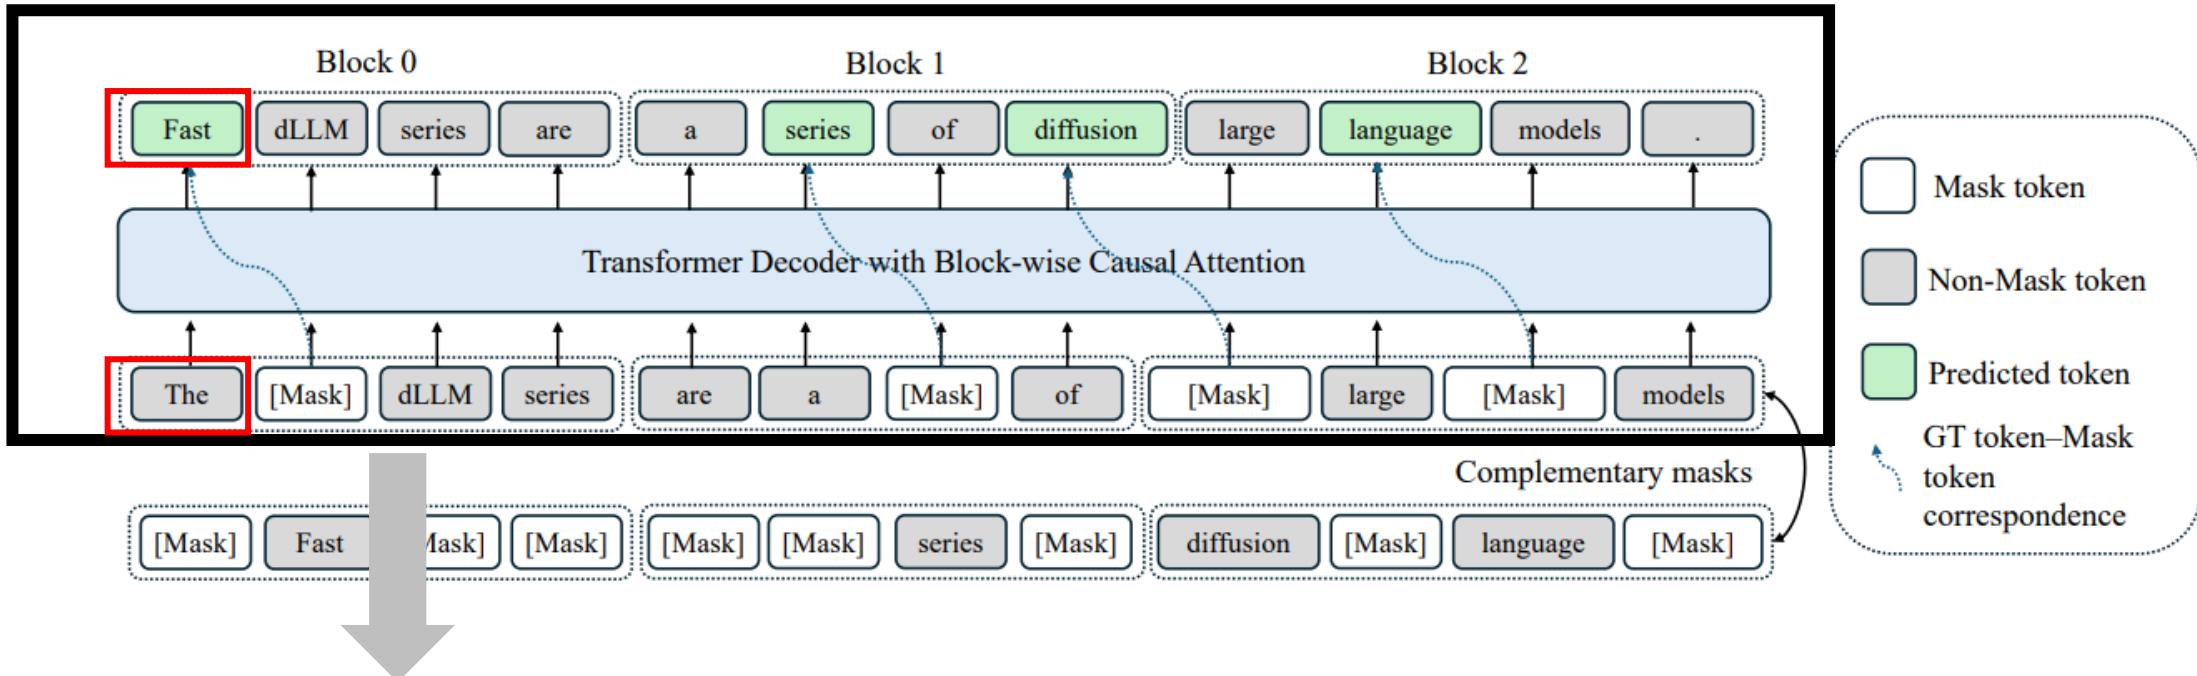
□ Background

□ Design

□ Evaluations

□ Discussion

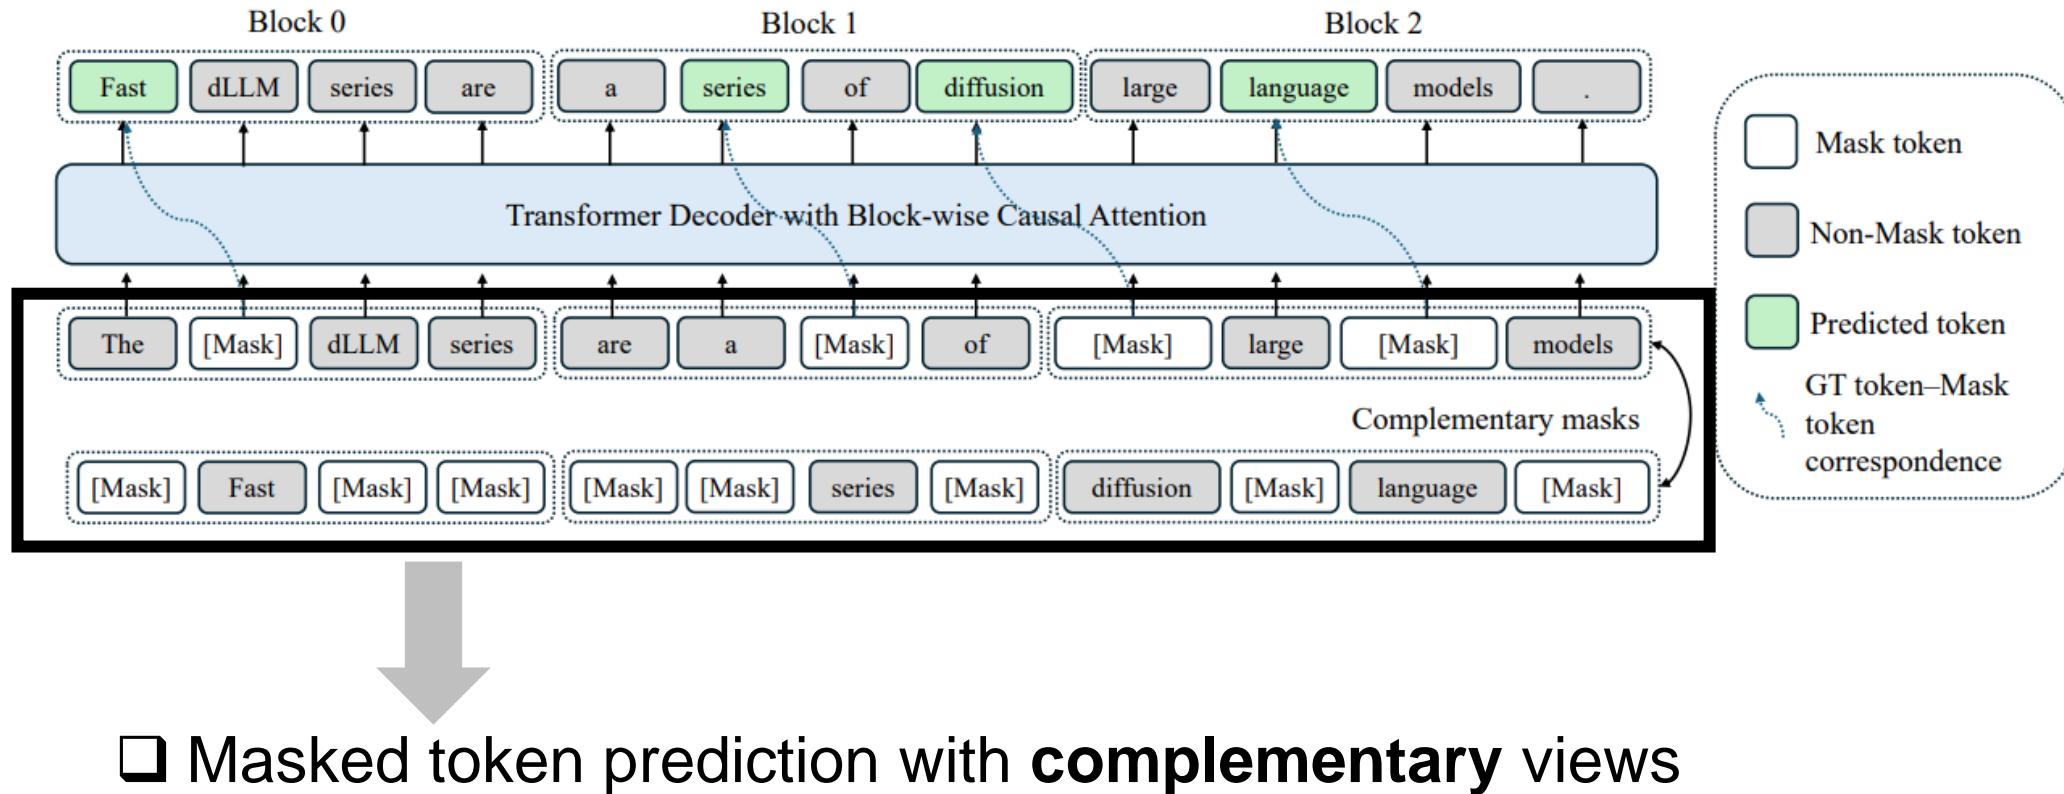
Training Process



❑ Token shift for prediction

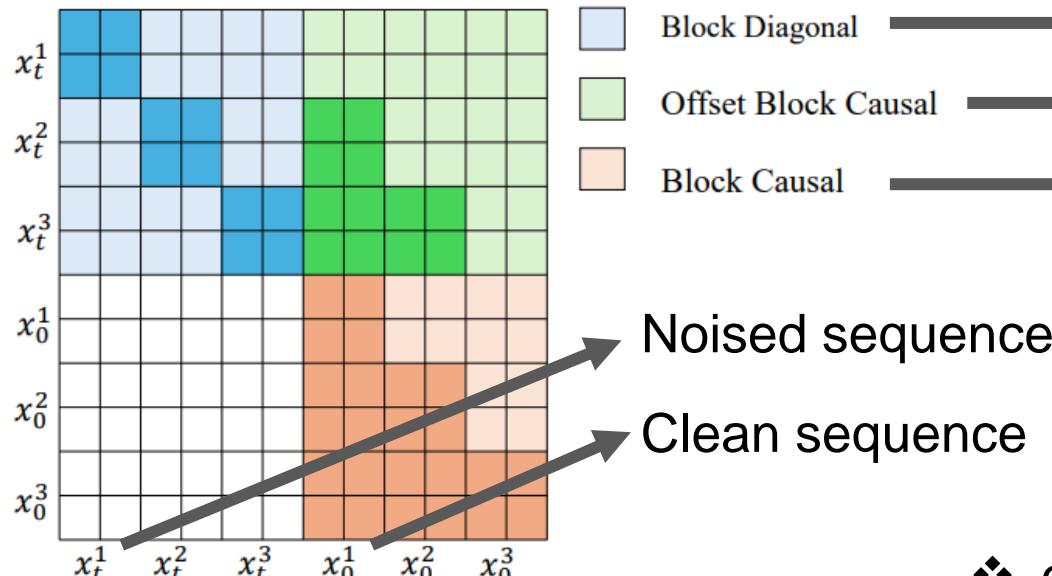
- ❖ Use hidden state at $i-1$ to predict x_i

Training Process

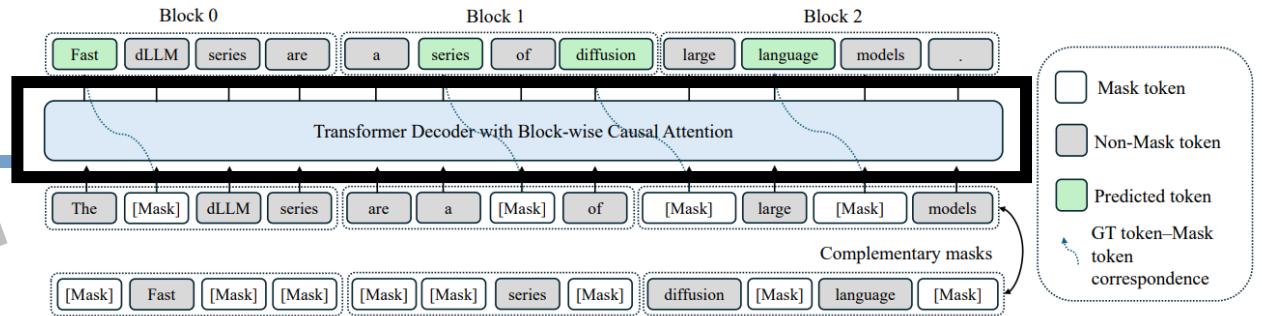


Training Process

Block-wise attention



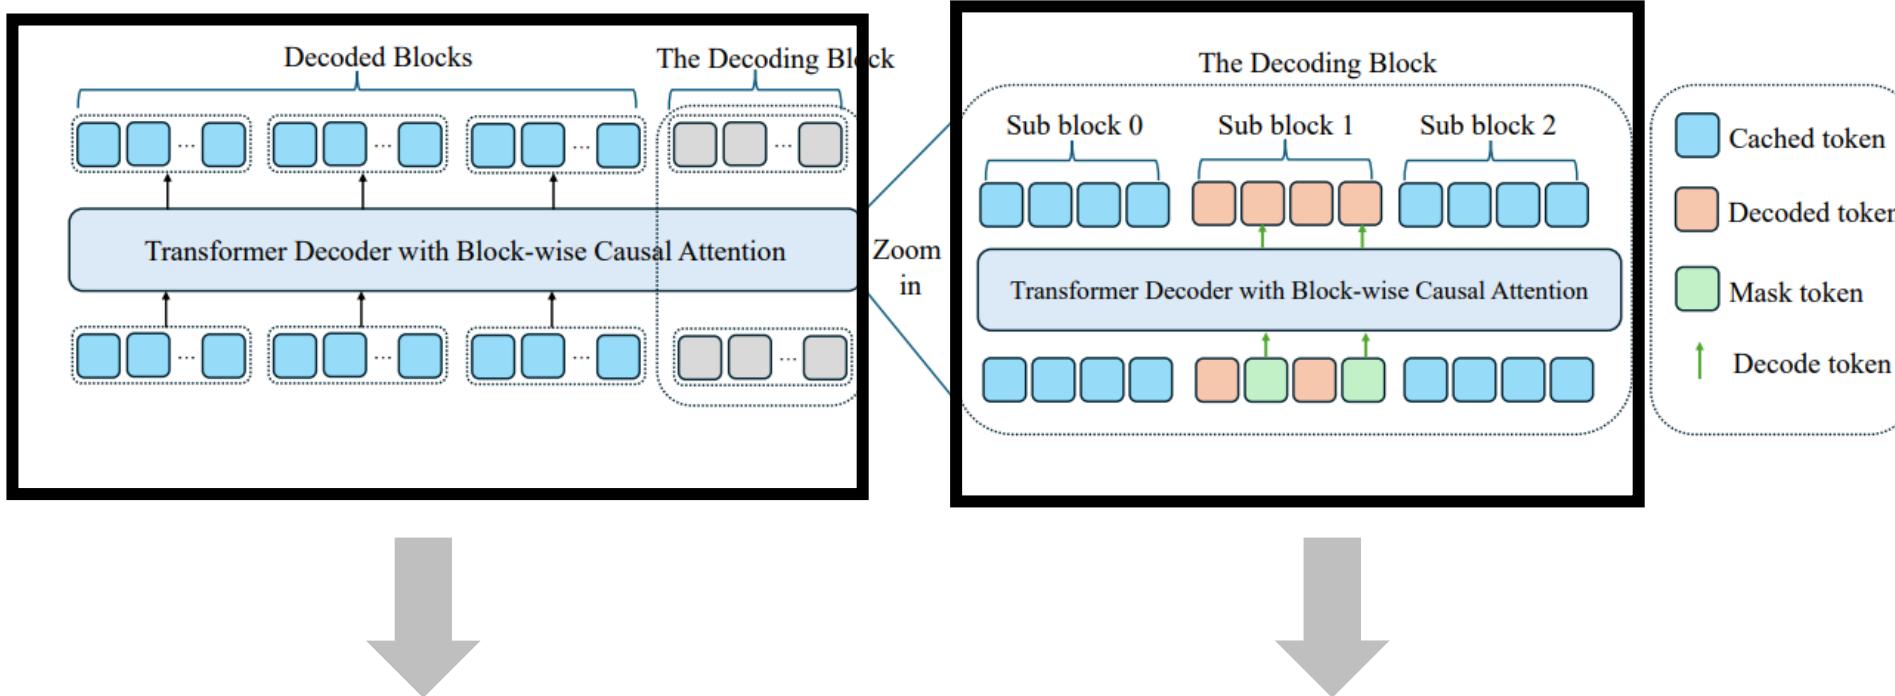
x^b denote set of tokens in the b-th
block (rather than the b-th **token**)



- Block Diagonal → bidirectional self-attention
 - Offset Block Causal → inter-block causal
 - Block Causal → Traditional AR
- Noised sequence
- Clean sequence
- ❖ employ the **flex-attention** implementation

Inference Pipeline

❑ Block-wise AR decoding with caching (Semi-AR)



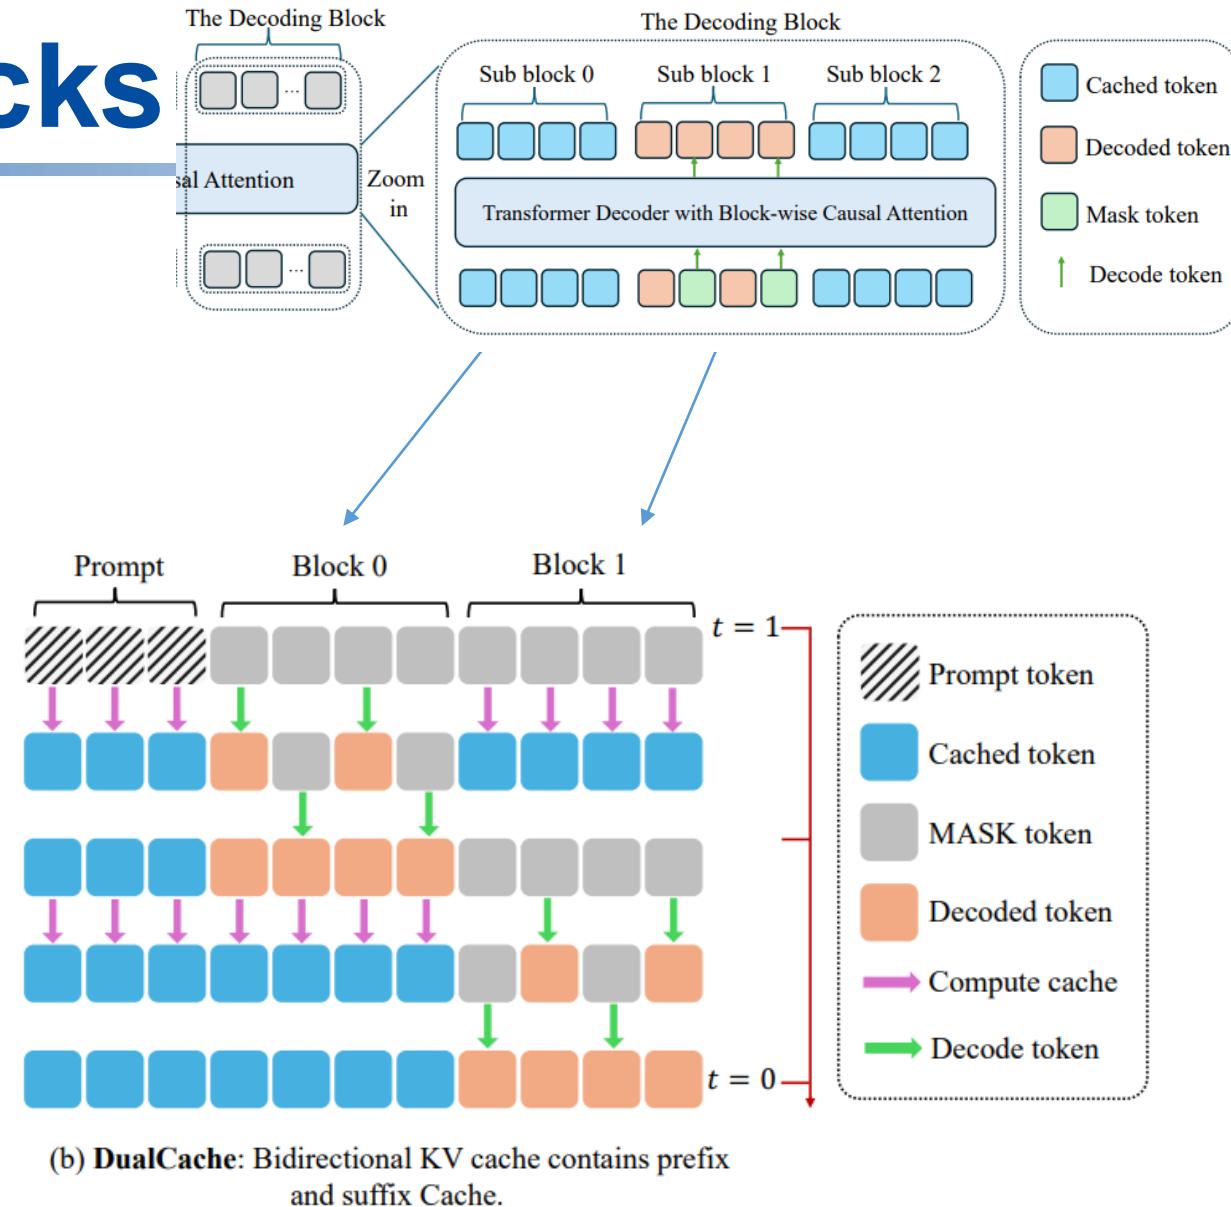
- ❑ AR across blocks
❖ Enable block-level KV cache

- ❑ Diffusion within blocks

Diffusion within blocks

Cannot leverage KV cache?

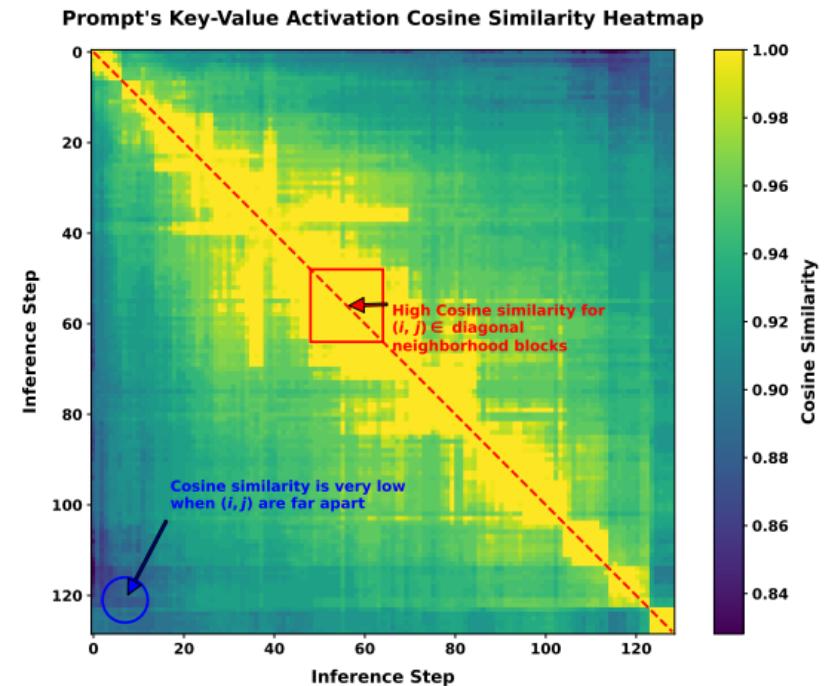
❑ **DualCache**: approximate KV cache



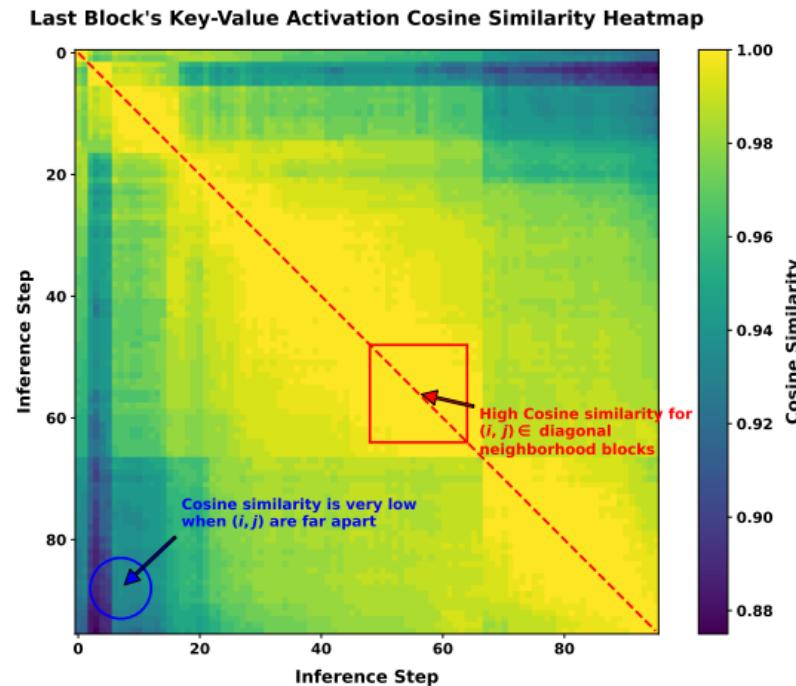
DualCache Intuition

Why works?

- Observation: KV activations exhibit **high similarity** across **adjacent inference steps within a block**



(a) Prompt block



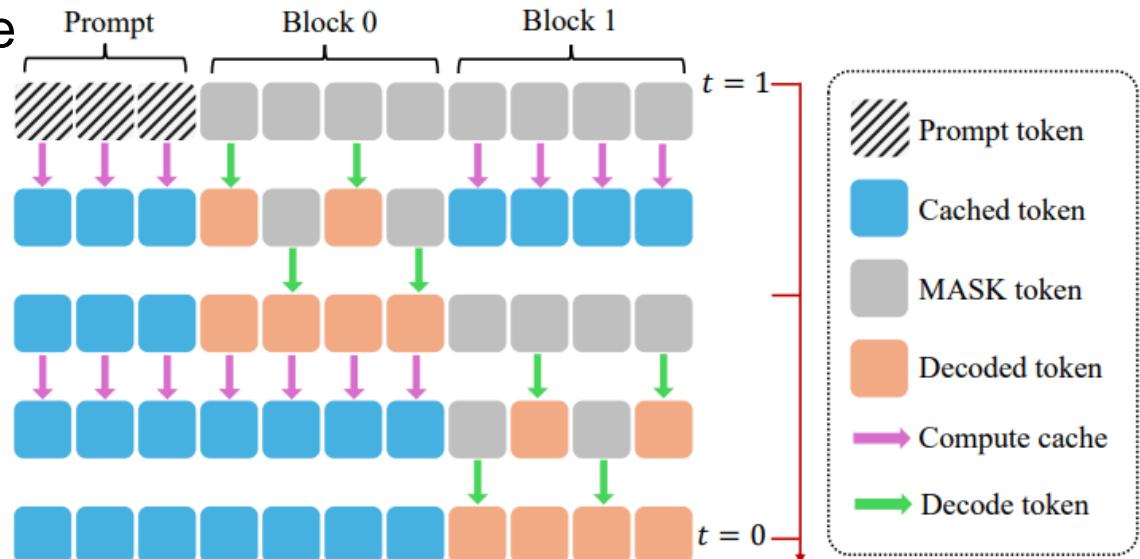
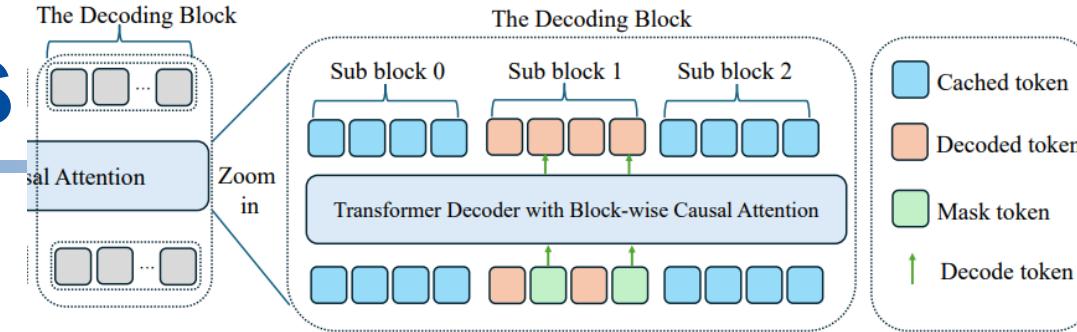
(b) Last block

Diffusion within blocks

Cannot leverage KV cache?

❑ DualCache: approximate KV cache

- ❖ Cache previous kv in one sub block inference
- ❖ **Update** when current sub block is finished



(b) **DualCache:** Bidirectional KV cache contains prefix and suffix Cache.

Performance under different cache block size

- Smaller block size incur **overhead**
- larger block size **diminish** accuracy
- Trade-off

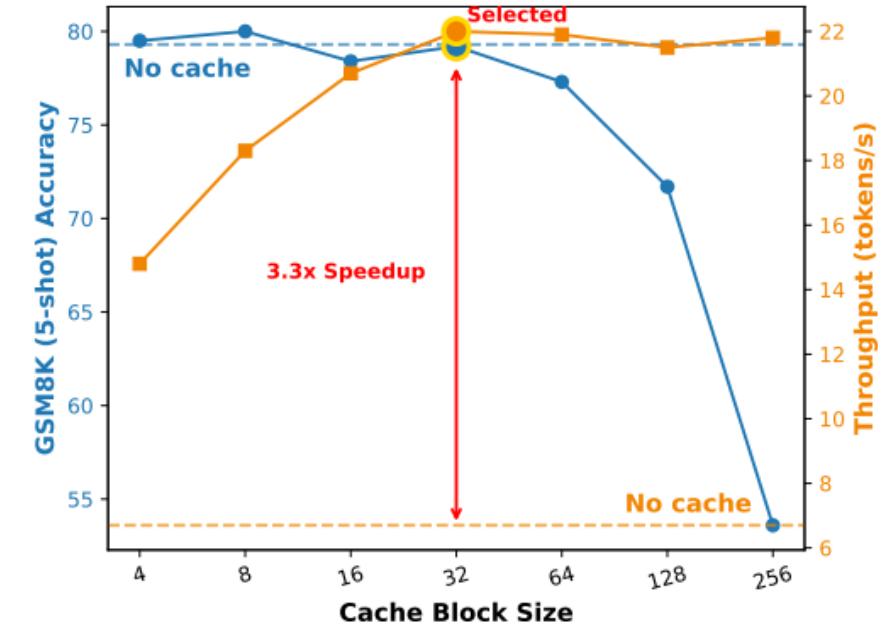
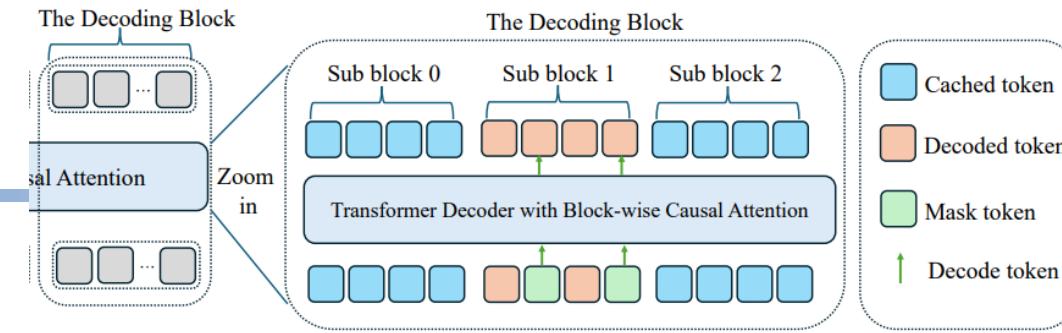


Figure 4 | **Impact of Cache Block Size on Accuracy and Throughput.** The orange line illustrates the effect of varying cache block size on throughput, while the blue line depicts accuracy.

Diffusion within blocks



□ Confidence-Aware Parallel Decoding

(TL;DR) If the model is confident on many positions, parallel decoding will not introduce errors.

Theorem 1 (Parallel Decoding under High Confidence). *Suppose there exists a specific sequence of tokens $\mathbf{x}^* = (x_{i_1}, \dots, x_{i_n})$ such that for each $j \in \{1, \dots, n\}$, the model has high confidence in x_{i_j} : $p_j(X_{i_j} = x_{i_j} | E) > 1 - \epsilon$ for some small $\epsilon > 0$. Then, the following results hold:*

1. *Equivalence for Greedy Decoding:* If $(n + 1)\epsilon \leq 1$ (i.e., $\epsilon \leq \frac{1}{n+1}$), then

$$\operatorname{argmax}_z p(z|E) = \operatorname{argmax}_z q(z|E) = \mathbf{x}^*. \quad (4)$$

This means that greedy parallel decoding (selecting $\operatorname{argmax} q$) yields the same result as greedy sequential decoding (selecting $\operatorname{argmax} p$).

This bound is **tight**: if $\epsilon > \frac{1}{n+1}$, there exist distributions $p(\mathbf{X}|E)$ satisfying the high-confidence marginal assumption for which $\operatorname{argmax}_z p(z|E) \neq \operatorname{argmax}_z q(z|E)$.

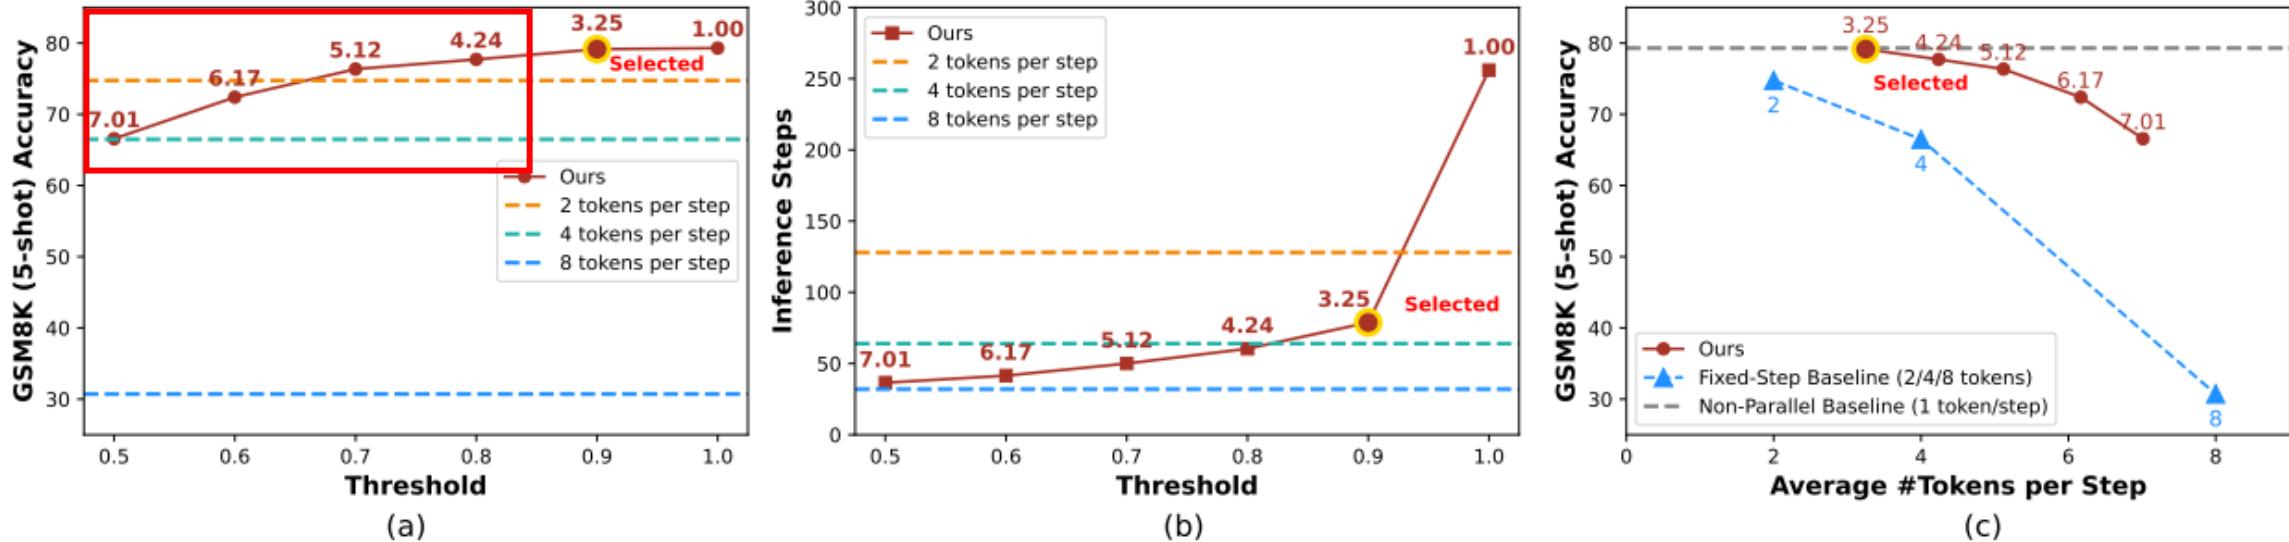
Diffusion within blocks

Algorithm 1 Block-wise Confidence-aware Parallel Decoding with (Dual) KV Cache

Require: p_θ , prompt p_0 , answer length L , blocks K , block size B , steps per block T , threshold τ , use_DualCache, strategy $\in \{\text{threshold, factor}\}$, factor f

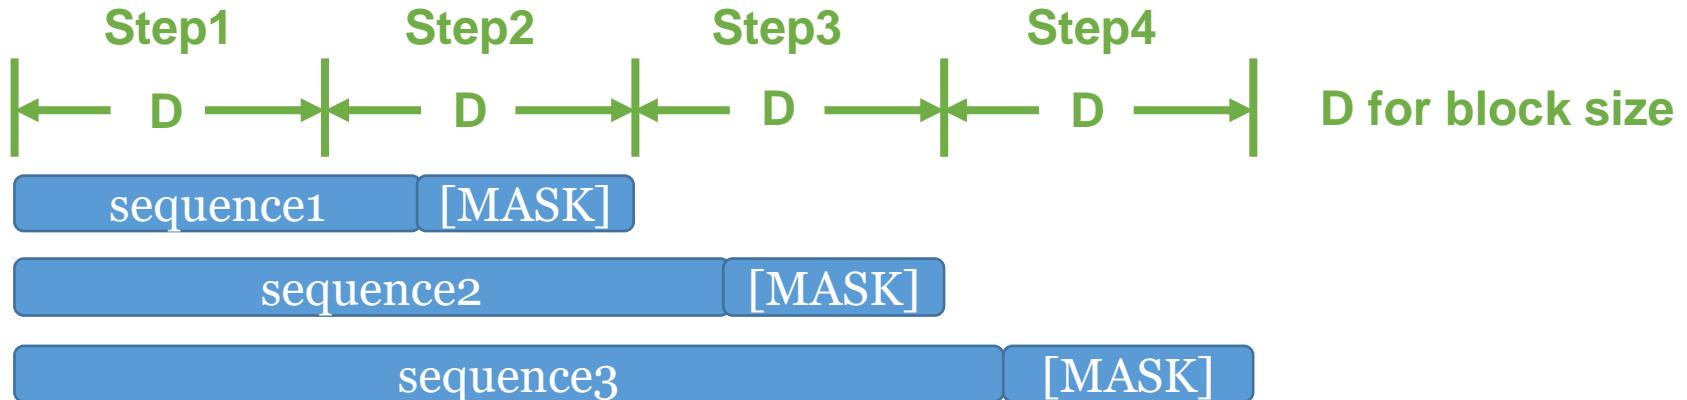
```
1:  $x \leftarrow [p_0; [\text{MASK}], \dots, [\text{MASK}]]$ 
2: Initialize KV Cache (single or dual) for  $x$  (fuse with decoding). // KV Cache Init
3: for  $k = 1$  to  $K$  do
4:    $s \leftarrow |p_0| + (k - 1)B$ ,  $e \leftarrow |p_0| + kB$ 
5:   for  $t = 1$  to  $T$  do
6:     Use cache, run  $p_\theta$  on  $x^{[s,e]}$  if use_DualCache else  $x^{[s,:]}$  // Cache Reuse
7:     For masked  $x^i$ , compute confidence  $c^i = \max_x p_\theta(x^i | \cdot)$  // Confidence scoring
8:     if strategy == threshold then
9:       Unmask all  $i$  in  $[s, e]$  with  $c^i \geq \tau$ , always unmask max  $c^i$ 
10:    else if strategy == factor then
11:      Sort  $c^i$  in descending order as  $(c^{(1)}, c^{(2)}, \dots, c^{(m)})$ 
12:      Find largest  $n$  such that  $(n + 1)(1 - c^{(n)}) < f$  → Unmask dynamic number of tokens
13:      Unmask top- $n$  tokens, always unmask the max  $c^i$ 
14:    end if
15:    if all  $x^{[s,e]}$  unmasked then
16:      break
17:    end if
18:  end for
19:  Update KV cache: if use_DualCache: prefix & suffix; else: prefix. // Cache Update
20: end for
21: return  $x$ 
```

Threshold Curve



- ❖ Baseline: Ilada previous fixed-N tokens
- ❖ Setting: GSM8K, 256 length

Batch decoding with padding



- Adaptation to block diffusion

Outline

□ Background

□ Design

□ Evaluations

□ Discussion

Setups

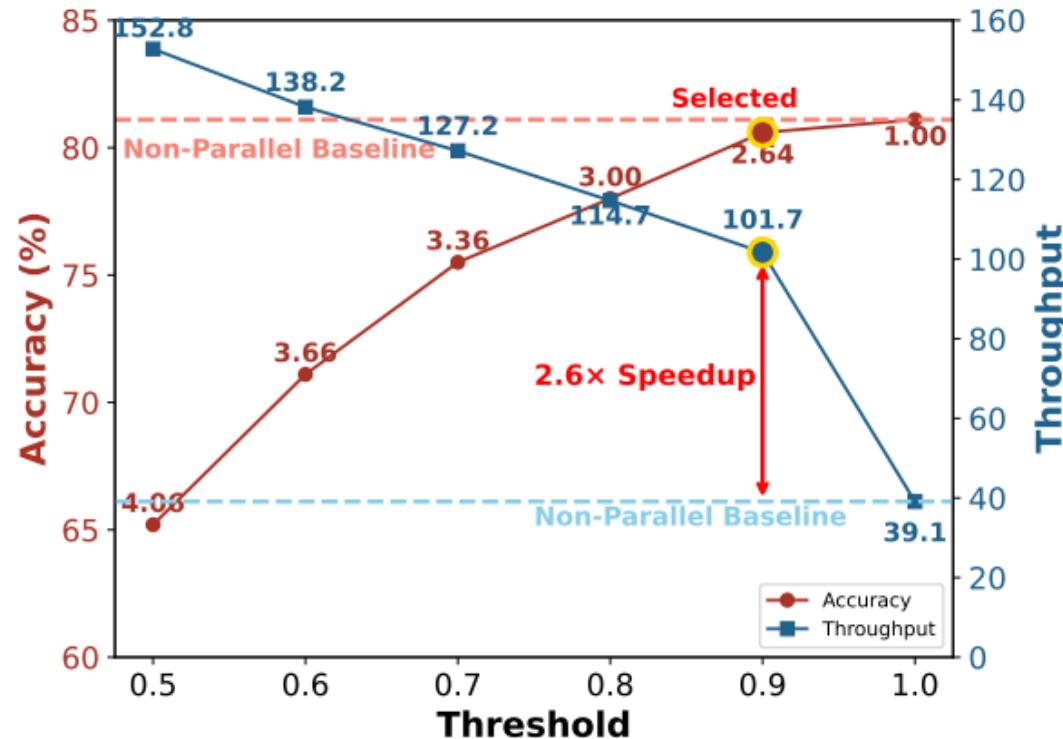
- ❑ Tuning model: Qwen-2.5 1.5B and 7B instruct
- ❑ Training dataset: LLaMA-Nemotron post-training dataset (batch 256)
- ❑ Training environment: 64 NVIDIA A100 GPUs
- ❑ Training configuration: 1.5B: learning rate 2×10^{-5} for 6,000 steps (costs 8h)
7B: learning rate 1×10^{-5} for 2500 steps (costs 12h)
- ❑ Block size: 32
- ❑ Sub-Block size: 8

Benchmark Results

Model	#Params	HumanEval		MBPP		GSM8K	Math	IFEval	MMLU	GPQA	Avg.
		Base	Plus	Base	Plus						
1B Models											
LlaMA-3.2	1.2B	34.1	31.1	34.1	29.4	43.0	23.8	58.9	44.4	24.1	35.9
SmolLM 2	1.7B	34.1	28.7	50.6	46.0	47.7	21.1	55.1	49.1	29.2	40.7
Qwen2.5-1.5B	1.5B	42.1	37.2	48.1	41.3	57.0	46.8	41.2	54.6	30.6	<u>44.3</u>
Qwen2.5-1.5B-Nemo-FT	1.5B	37.2	33.5	53.4	44.4	58.5	43.5	39.4	58.1	31.0	<u>44.3</u>
Fast-dLLM v2	1.5B	43.9	40.2	50.0	41.3	62.0	38.1	47.0	55.1	27.7	45.0
7B+ Models											
LLaDA	8B	35.4	31.7	31.5	28.6	78.6	26.6	59.9	65.5	31.8	43.3
LLaDA-1.5	8B	52.4	-	42.8	-	83.3	42.6	58.2	66.0	36.9	-
LLaDA-MoE	7B	61.6	-	70.0	-	82.4	58.7	59.3	67.2	-	-
Dream	7B	57.9	53.7	68.3	56.1	81.0	39.2	62.5	67.0	33.0	57.6
Qwen2.5-7B	7B	51.2	47.6	57.7	49.5	71.4	73.3	70.8	68.7	33.5	58.2
Qwen2.5-7B-Nemo-FT	7B	52.4	48.2	57.1	50.0	84.1	72.0	69.5	68.6	34.2	<u>59.6</u>
Fast-dLLM v2	7B	63.4	58.5	63.0	52.3	83.7	61.6	61.4	66.6	31.9	60.3

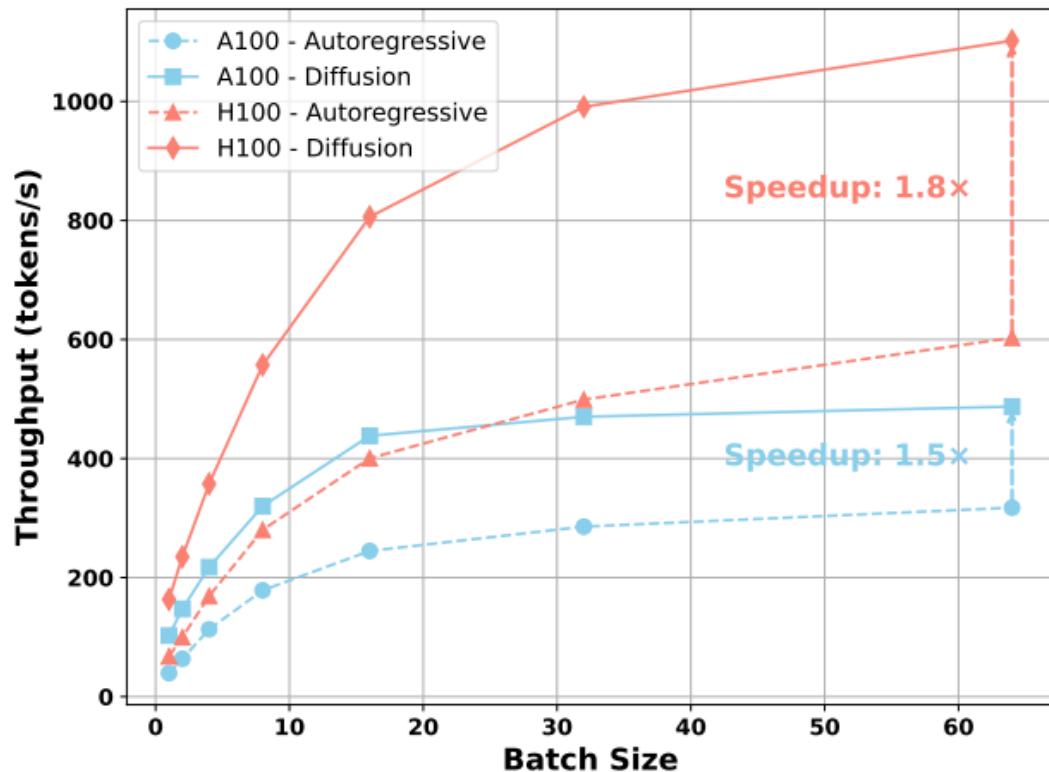
Performance

□ Accuracy and throughput under different thresholds on GSM8K



Performance

- Throughput comparison between AR(Qwen2.5-7B-Instruct) and diffusion(Fast-dLLM v2 7B) generation for GSM8K



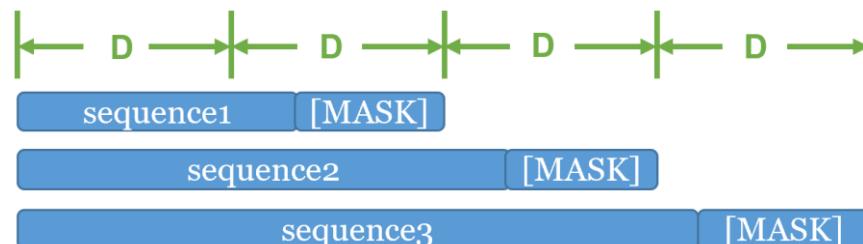
Ablation Study

- ❖ Based on Fast-dLLM v2 1.5B

Method	HumanEval		MBPP		GSM8K	Math	IFEval	MMLU	GPQA	Avg.
	Base	Plus	Base	Plus						
Naive token shift	<u>38.4</u>	32.9	44.4	<u>38.6</u>	59.0	<u>37.3</u>	39.9	52.9	27.9	41.3
+ pad	<u>38.4</u>	<u>34.1</u>	<u>45.2</u>	38.4	<u>60.1</u>	37.0	<u>45.8</u>	<u>53.5</u>	<u>27.7</u>	<u>42.2</u>
+ pad + CM	43.9	40.2	50.0	41.3	62.0	38.1	47.0	55.1	<u>27.7</u>	45.0

- ❖ Naive token shift: use hidden state at $i-1$ to predict x_i

- ❖ Pad:



- ❖ CM: complementary mask

Ablation Study

- Sub-Block size and Block size affect performance

Table 3 | Sub-Block size decoding improves performance, with size 8 being optimal.

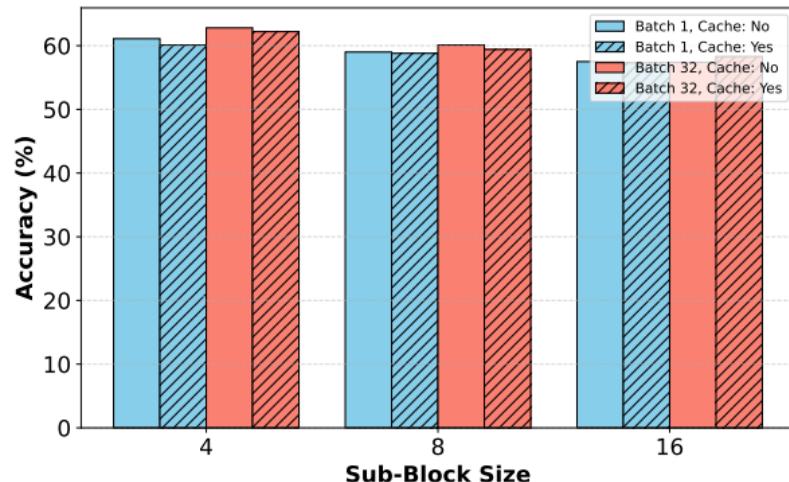
Sub-Block Size	2	4	8	16	32
GSM8K	62.8	61.8	<u>62.0</u>	61.3	60.2
HumanEval	42.7	<u>43.3</u>	43.9	39.6	38.4
HumanEval+	<u>39.6</u>	40.2	40.2	36.0	34.8

Table 4 | Inference with mismatched sizes reduces performance.

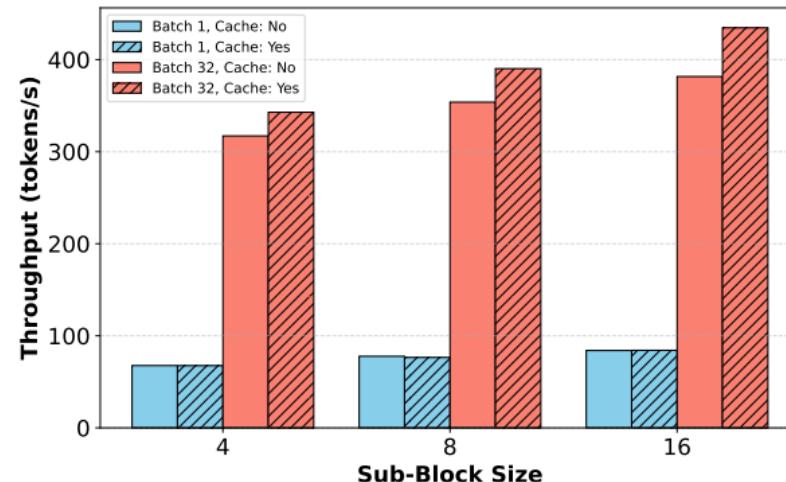
Block Size	2	4	8	16	32
GSM8K	53.2	56.8	58.5	<u>59.7</u>	60.2
HumanEval	37.8	43.3	43.3	<u>38.4</u>	<u>38.4</u>
HumanEval+	34.1	<u>39.0</u>	39.6	34.1	34.8

Ablation Study

- Cache is a purely efficiency-enhancing feature



(a)



(b)

Figure 6 | Effect of small block size and sub-block cache on model performance. **(a)** Accuracy remains largely unaffected by the use of sub-block cache across different block sizes and batch sizes. **(b)** Throughput increases as small block size grows due to higher decoding parallelism. While sub-block cache has negligible effect when batch size is small, it significantly improves throughput under compute-bound settings (e.g., batch size = 32).

Outline

□ Background

□ Design

□ Evaluations

□ Discussion

Discussion

□ Trade-off between accuracy and throughput in MDMs

- ❖ Parallel generation inevitably introduces conditional independence

□ Arbitrary-Order Autoregressive?

- ❖ Semi-AR Semi-diffusion

□ Data-efficient fine-tuning

- ❖ Fast-dLLM v2 achieves lossless adaptation with just ~1B tokens, compared to ~500B tokens required by Dream
- ❖ Diffusion Beats Autoregressive in Data-Constrained Settings (NeurIPS 25)

Thanks!