

Enabling Efficient GPU Communication over Multiple NICs with FuseLink

Group: Haiquan Wang, Tonghuan Xiao and Jiahui Tan

2025-11-04




Agenda

- Background**
- Motivation**
- Opportunities and Challenges**
- Design**
- Evaluations**
- Conclusion**



Background


- Inter-server communication is everywhere in large scale machine learning tasks
- PD disaggregation: KV cache transmission





Background


- Inter-server communication is everywhere in large scale machine learning tasks
- MoE model training: token dispatch and combine in expert parallel





Background

- Inter-server communication is everywhere in large scale machine learning tasks
- Recommend System Embedding Table Traffic






Background

- ❑ **Distributed ML tasks**
 - ❑ PD disaggregation: KV cache transmission
 - ❑ MoE model training: token dispatch and combine in expert parallel
 - ❑ Recommendation model training: embedding table traffic
- ❑ **Efficient communication among GPUs is crucial**



Background


- Key Hardware Technology: Accelerating Data Exchange
 - NVLink/NVSwitch: intra-server
 - RDMA: inter-server





Background

- The bandwidth of single NIC is much lower than NVLink
- How about stack multiple network interface cards ?





Agenda

□ Background

□ Motivation

□ Opportunities and Challenges

□ Design

□ Evaluations

□ Conclusion





Motivation

□ Data Distribution lead to resource waste

□ Blanced Traffic & Full Utilzation


□ UnBlanced Trafficc & Underutilized





Motivation


- PD disaggregation: KV cache transmission
- Stochastic Task Arrival & multiple model instance





Motivation


- ❑ PD disaggregation: KV cache transmission
- ❑ Stochastic Task Arrival & multiple model instance





Motivation


- ❑ PD disaggregation: KV cache transmission
- ❑ Varying Requests Size





Motivation


- ❑ PD disaggregation: KV cache transmission
- ❑ Varying Requests Size





Motivation



- ❑ PD disaggregation: KV cache transmission
- ❑ Varying Requests Size





Motivation

- MoE model training: token dispatch and combine in expert parallel
- Different Token Batches







Motivation

- MoE model training: token dispatch and combine in expert parallel

- Asynchronous Transmission



Multiple, concurrent data overwhelm and congest a single NIC.





Motivation

- Recommend System Embedding Table Traffic
- Different Embeddings





Motivation

- Inter-server communication Dynamic traffic in large scale machine learning tasks
- PD disaggregation: KV cache transmission
- MoE model training: token dispatch and combine in expert parallel
- Recommend System Embedding Table Traffic

ML Tasks	Communication Pattern Cause			NIC util.	Comm. ratio
	via direct NICs	Concurrent transmission	Same traffic volume		
Distributed LLM Serving	✓	✗ stochastic task arrival	✗ varying requests sizes	13%-53%	11%-82%
Expert-parallel MoE	✓	✗ async. transmission	✗ different token batches	29%-65%	15%-42%
DLRM	✓	✓	✗ different embeddings	59%-82%	28%-55%

Table 1: Dynamic-traffic ML tasks with communication patterns, ✓(✗) marks satisfied (unsatisfied) patterns





Agenda

- ❑ Background
- ❑ Motivation
- ❑ Opportunities and Challenges
- ❑ Design
- ❑ Evaluations
- ❑ Conclusion



Existing Method: PXN(PCI x Nvlink)


- When the load is unbalanced, PXN still fails to make full use of the bandwidth of each NIC.
- PXN enables a GPU to communicate with an NIC on the node through NVLink and then PCI.
- NCCL with PXN is limited by statically bind traffic to paths and cannot support transmitting data through NVLink on receivers.





Opportunities

- ❑ Multi-NICs acceleration under imbalanced traffic
 - ❑ Leveraging idle NICs
 - ❑ NIC load status aware
 - ❑ Dynamic NVLink routing to bypass PCIe





Challenges

The incompatibility of intra-server and inter-server networks presents following technical challenges

□ Relaying overhead

□ GPU-based relaying incurs significant synchronization overhead, resulting in minimal NIC throughput improvement.

□ Interruption and contention risks

□ Routing traffic to idle NICs risks bandwidth contention and performance disruption when new communications arise from peer workers on relay GPUs.

□ Scheduling overhead

□ Dynamic NIC scheduling incurs high control overhead.



Agenda


- ❑ Background
- ❑ Motivation
- ❑ Opportunities and Challenges
- ❑ Design
- ❑ Evaluations
- ❑ Conclusion



Design - Overview

Workflow of FuseLink


- ❑ Explores the intra-server topology and establishes RDMA connections
- ❑ Identifies available NICs for inter-server communication
- ❑ Selects the optimal data path for each NIC and GPU





Design - Relaying

- ❑ Inter-server communication (w/o relaying)
 - ❑ The GPU fills network buffer (pinned for RDMA communication)
 - ❑ The GPU notifies CPU to initiate RDMA transmission





Design - Relaying


- ❑ Inter-server communication (w/o relaying)
 - ❑ The GPU fills network buffer (pinned for RDMA communication)
 - ❑ The GPU notifies CPU to initiate RDMA transmission

- ❑ Inter-server communication (w/ relaying)
 - ❑ The GPU0 fills network buffer of GPU1
 - ❑ The GPU0 notifies CPU to initiate RDMA transmission



Design - Relaying

- ❑ 4 relaying path as candidates
 - ❑ D1: GPU writes data to remapped buffers on relay GPUs
 - ❑ D2: CPU copies data to relay GPUs then RDMA






Design - Relaying

❑ 4 relaying path as candidates

- ❑ D1: GPU writes data to remapped buffers on relay GPUs
- ❑ D2: CPU copies data to relay GPUs then RDMA






Design - Relaying

- ❑ 4 relaying path as candidates
- ❑ D3: GPU threads write data to buffers mapped to host memory
- ❑ D4: CPU initiates memory copy to host memory




Design - Relaying

- ❑ 4 relaying path as candidates
- ❑ D3: GPU threads write data to buffers mapped to host memory
- ❑ D4: CPU initiates memory copy to host memory





Design - Relaying





Design - Relaying

- ❑ D3 and D4 achieve the same bandwidth as PCIe, thus limited by PCIe speed and cannot improve relaying efficiency
- ❑ D1 achieves higher data path efficiency
 - ❑ Only one copy step to relay GPUs, leading to low message latency over indirect NIC
 - ❑ Make effect without additional CPU involvement
 - ❑ Hide the extra intra-server relaying latency by overlapping it with pipelined intra- and inter-server communication



Design - Relaying

- ❑ D3 and D4 achieve the same bandwidth as PCIe, thus limited by PCIe speed and cannot improve relaying efficiency
- ❑ D1 achieves higher data path efficiency
 - ❑ Only one copy step to relay GPUs, leading to low message latency over indirect NIC

Use D1 as relaying traffic

~~Reduce the extra intra-server relaying latency by overlapping it with pipelined intra- and inter-server communication~~



Design - Mitigate Contention


❑ Data relaying incurs two types of resources contention

❑ GPU Memory Contention

❑ Allocate extra buffer in relay GPU

❑ Exhaust memory of relay GPU causing OOM

❑ Cannot have precise estimation of memory usage for dynamic workload





Design - Mitigate Contention

- ❑ Relaying without interrupting memory allocation
 - ❑ Set a configurable upper limit for relaying memory usage
 - ❑ Prioritize memory demands of running tasks on relay GPUs




Design - Mitigate Contention

□ Data relaying incurs two types of resources contention

□ Communication Contention

□ Disturb or even interrupt peer GPU communication over direct NICs.





Design - Mitigate Contention

FuseLink provides an interruption-free relaying, including 4 steps

- ❑ Monitor NIC status within the server
- ❑ Aggregate NIC load status on the sender and receiver side
- ❑ Select NICs and initiating data transmission
- ❑ Perform intra-server traffic routing inside receivers via memory remapping



Design - Mitigate Contention


- ❑ Worker-aware NIC monitoring
 - ❑ Mark the status of NICs as idle or busy
 - ❑ ML applications typically have large traffic that can fully occupy NICs when they have communication tasks
 - ❑ Decided by workers that have the highest priorities on the NIC
 - ❑ Monitor NICs workload from the work requests on RDMA connections
 - ❑ FuseLink identifies idle NICs by periodically checking new completions of network operations posted by high-priority workers



Design - Mitigate Contention

❑ Load-aware Scheduling

- ❑ Aggregate NIC load status from receivers
- ❑ Write NIC information into credit
- ❑ Collect credits from receivers






Design - Mitigate Contention

❑ Load-aware Scheduling

- ❑ Aggregate NIC load status from receivers
 - ❑ Write NIC information into credit
 - ❑ Collect credits from receivers
- ❑ Select the NIC for sending
 - ❑ Selects the direct NIC of worker if it is idle






Design - Mitigate Contention

□ Load-aware Scheduling


- Aggregate NIC load status from receivers
 - Write NIC information into credit
 - Collect credits from receivers
- Select the NIC for sending
 - Selects the direct NIC of worker if it is idle
 - Otherwise, select an idle indirect NIC if exists
 - If no idle NICs found, select the direct NIC






Design - Mitigate Contention


credit recv NIC idle: None



credit recv NIC idle: None




credit recv NIC idle: None






Design - Mitigate Contention


credit recv NIC idle: None




credit recv NIC idle: None




credit recv NIC idle: None




Limit the number of operations posted to indirect NICs



identify contention



schedule to direct NIC



adjust sender routing



Design - Scheduling with Efficiency

FuseLink has following tradeoffs to ensure scheduling with efficiency:

- ❑ **Mark NIC load status based on new completions of operations posted by GPU workers, which is later than transmission start time.**
- ❑ **Allows bounded contention over limited number of network operations**
- ❑ **FuseLink allows bounded suboptimal transmission through indirect NICs**



Design - Summary

FuseLink can:

- ❑ Leverage idle NICs through dedicated intra-server fabrics
- ❑ Mitigate contention and interruption among GPUs
- ❑ Be seamlessly integrated into existing communication library
(NCCL, GLOO...)



Agenda

- ❑ Background
- ❑ Motivation
- ❑ Opportunities and Challenges
- ❑ Design
- ❑ Evaluations
- ❑ Conclusion



Evaluation

❑ Setup

- ❑ Intel 8480C CPU, 8*Nvidia Hopper GPUs
- ❑ Inter-server: 8*Connect-X7 50GB/s NICs
- ❑ Intra-server: eight-lane NVLinks & NVSwitches, delivering up to 200GB/s
- ❑ FuseLink scheduler as a daemon process ,FuseLink as a NCCL plugin

❑ Baseline

- ❑ NCCL with PXN



Evaluation

❑ In End-to-End Evaluation

❑ Q1: What are the performance benefits of FuseLink on machine learning tasks with imbalanced traffic?

❑ In Microbenchmark


❑ Q1: What is the maximum inter-server bandwidth for a single GPU with FuseLink compared to the default?

❑ Q2: What is the overhead of FuseLink in traffic scheduling?




Evaluation


- ❑ End-to-End performance (Model serving)
- ❑ FuseLink shows **1.04-2.73 x → speedup over the baseline with NIC binding**
- ❑ OPT models serving(30B) under 1/2/4 tensor parallel degrees



(a) Eight instances with TP=1



(b) Four instances with TP=2



(c) Two Instances with TP=4

Figure 11: FuseLink model serving performance under different TP degrees and number of serving instances

#Instance	Setting	P50	P99
8	NIC Binding	684.54 ms	903.20 ms
	FuseLink	308.48 ms	464.55 ms
4	NIC Binding	174.46 ms	297.49 ms
	FuseLink	122.61 ms	259.73 ms
2	NIC Binding	98.09 ms	175.36 ms
	FuseLink	81.97 ms	160.36 ms

Table 4: Model serving TTFT comparison under different number of serving instances within a server



Evaluation

❑ EP training

❑ FuseLink improves the training throughput by $1.3\times$

❑ Mixtral 8*22B model (TP=4 , EP=8)




Figure 12: EP training times under NIC binding and FuseLink

❑ DLRM training

❑ FuseLink can effectively reduce the training iteration times

❑ Criteo advertisement dataset
(DP=32, batch size=1024)





Figure 13: DLRM training iteration times under different cache sizes (GB)



Evaluation

□ Microbenchmark-bandwidth improvement

- FuseLink achieves up to 212GBps bandwidth with six NICs, compared to the default setting with about 50GBps throughput



- NVLink > Indirect NICs bandwidth
 - Maximum bandwidth: all NICs
- NVLink < Indirect NICs bandwidth
 - Maximum bandwidth:
direct NIC(50GB/s)+Nvlink(200GB/s)



Evaluation

- Microbenchmark-scheduling overhead
 - The scheduling overhead is acceptable and the tradeoff will not harm the overall performance.
 - Relay remapping is called when optimizing the data paths between NICs and GPUs, which brings around 95-193us latency in each remapping.

Controls	Avg. cost(us)
post & pull ops	0.8-1.4
query NIC load	0.9-1.6
process send	2.8-3.5
process recv	4.9-5.6

Table 3: Network operation processing overheads in FuseLink




Figure 10: Average scheduling overhead decomposed into flushing data, changing NIC, and remapping



Agenda

- ❑ Background
- ❑ Motivation
- ❑ Opportunities and Challenges
- ❑ Design
- ❑ Evaluations
- ❑ Conclusion



Conclusion

- Why FuseLink?
- FuseLink enables inter-server GPUs to efficiently communicate over multiple NICs.
- FuseLink exploits dedicated intra-server network for intra-server traffic routing to fully utilize GPU interconnects.
- FuseLink is integrated into NCCL, so that ML applications can benefit from FuseLink seamlessly without code modifications.
- Fuselink implementation has about 3000 lines of code in C++.