

HydraServe: Minimizing Cold Start Latency for Serverless LLM Serving in Public Clouds

Authors: Chiheng Lou, Sheng Qi , Chao Jin, Dapeng Nie,
Haoran Yang,
Yu Ding, Xuanzhe Liu, Xin Jin


Presented by **Jiyang Wang**

2025-11-18



Background: LLM Inference


□ Autoregressive manner





Background: LLM Inference

- ❑ Autoregressive manner
- ❑ Prefill & decode





Background: LLM Inference


- ❑ Autoregressive manner
- ❑ Prefill & decode
- ❑ Service level objectives (SLOs)

- ❖ Time to first token (TTFT)

- Real-time chatbot

- ❖ Time per output token (TPOT)


- Article writing






Background: LLM

- Autoregressive manner
- Prefill & decode
- Service level objectives (SLOs)
 - ❖ Time to first token (TTFT)
 - ❖ Time per output token (TPOT)
- Pipeline parallelism in training






Background: Serverless LLM Serving






Background: Serverless LLM Serving






Background: Serverless LLM Serving






Background: Serverless LLM Serving



For user : Pay-per-use billing



For production : long-tail
models





Background: Serverless LLM Serving


- ❑ Network bandwidth constraints
 - ❖ Sharing among multiple instances



Background: Serverless LLM Serving

❑ Network bandwidth constraints

- ❖ Sharing among multiple instances
- ❖ Cost-efficiency requirement






Background: Serverless LLM Serving

❑ Network bandwidth constraints

- ❖ Sharing among multiple instances
- ❖ Cost-efficiency requirement

User { GPU capabilities is important
Prioritize cost saving



Provider Minimize cost per GPU

Cost of AWS EC2 instance (L40S)

Instance	Mem.(GB)	Band.(Gbps)	#GPU	Cost(\$/h)	Cost/GPU(\$/h)	
g6e.xlarge	32	up to 20	1	1.861	1.861	
g6e.2xlarge	64	up to 20	1	2.24208	2.24208	~120 %
g6e.4xlarge	128	20	1	3.00424	3.00424	
g6e.8xlarge	256	25	1	4.52856	4.52856	
g6e.16xlarge	512	35	1	7.57719	7.57719	~400 %
g6e.12xlarge	384	100	4	10.49264	2.62316	
g6e.24xlarge	768	200	4	15.06559	3.76640	
g6e.48xlarge	1536	400	8	30.13118	3.76640	



Cold start latency is long



Reduce CPU, memory and network resources



Background: Serverless LLM Serving

□ Network bandwidth constraints

- ❖ Sharing among multiple instances
- ❖ Cost-efficiency requirement

User { GPU capabilities is important
Prioritize cost saving

Provider Minimize cost per GPU

Cost of AWS EC2 instance (L40S)

Instance	Mem.(GB)	Band.(Gbps)	#GPU	Cost(\$/h)	Cost/GPU(\$/h)	
g6e.xlarge	32	up to 20	1	1.861	1.861	
g6e.2xlarge	64	up to 20	1	2.24208	2.24208	~120 %
g6e.4xlarge	128	20	1	3.00424	3.00424	
g6e.8xlarge	256	25	1	4.52856	4.52856	
g6e.16xlarge	512	35	1	7.57719	7.57719	~400 %
g6e.12xlarge	384	100	4	10.49264	2.62316	
g6e.24xlarge	768	200	4	15.06559	3.76640	
g6e.48xlarge	1536	400	8	30.13118	3.76640	



Cold start latency is long



Reduce CPU, memory and network resources




Trade-off between cold start latency & cost benefits



Background: Serverless LLM Serving

❑ Network bandwidth constraints

❑ Breakdown of cold starts



❖ **Model fetching** : From remote storage to local memory

❖ **Model loading** : Transfer weights to GPU; CUDA graph & KV cache init

~2s



Observation & Opportunities



Overlapping!

- ❑ Model fetching can be parallelism
 - ❖ Each worker only host a part of model



Observation & Opportunities



Overlapping!

- ❑ Model fetching can be parallelism
 - ❖ Each worker only host a part of model
- ❑ Model loading and library loading use different resource
 - ❖ The two phases can be parallelism



Observation & Opportunities



Overlapping!

- ❑ Model fetching can be parallelism
 - ❖ Each worker only host a part of model
- ❑ Model loading and library loading use different resource
 - ❖ The two phases can be parallelism
- ❑ Model fetching and loading can be pipelined
 - ❖ At tensor granularity to hide overhead





Observation & Opportunities



Overlapping!

- ❑ Model fetching can be parallelism in different worker
- ❑ Parallel Model loading and library loading inside same worker
- ❑ Model fetching and loading pipelined at tensor granularity







HydraServe: Cluster–Level

□ Cluster–level overall design

- ❖ Multiple workers on different servers
- ❖ Each worker fetches a part of model
- ❖ Workers exchange intermediate results during inference



Normal cold start



Cold start with 2 workers




HydraServe: Cluster-Level

□ Trade-off analysis

Setup : 4 servers (A10, 188GB memory, 16 Gbps network bandwidth)

- ❖ Larger parallelism sizes reduce model fetching time
 - Other stage also matters







HydraServe: Cluster-Level

□ Trade-off analysis

Setup : 4 servers (A10, 188GB memory, 16 Gbps network bandwidth)

- ❖ Larger parallelism sizes reduce model fetching time
- ❖ Pipeline parallelism has a modest impact on TPOT
 - 8 KB of inter-layer results per token in Llama2-7B








HydraServe: Cluster-Level

□ Trade-off analysis

Setup : 4 servers (A10, 188GB memory, 16 Gbps network bandwidth)

- ❖ Larger parallelism sizes reduce model fetching time
- ❖ Pipeline parallelism has a modest impact on TPOT
- ❖ Worker colocation leads to longer TPOT





HydraServe: Cluster–Level

□Algorithm design

❖ Target

- Satisfy TTFT & TPOT
- Minimum GPU sharing

❖ Search space

- Pipeline parallelism size (**s**)
 - When $s > 4$, yield little improvement



HydraServe: Cluster–Level

□Algorithm design

❖ Target

- Satisfy TTFT & TPOT
- Minimum GPU sharing

❖ Search space

- Pipeline parallelism size ($1 \leq s \leq 4$)
- GPU memory to each worker
 - The same as non-parallelized (M)
 - The minimal memory required ($\frac{M}{s}$)



HydraServe: Cluster–Level

□Algorithm design

❖ Target

- Satisfy TTFT & TPOT
- Minimum GPU sharing

❖ Search space

- Pipeline parallelism size ($1 \leq s \leq 4$)
- GPU memory to each worker
 - The same as non-parallelized (M)
 - The minimal memory required ($\frac{M}{s}$)
- Number of full-memory workers (w)



HydraServe: Cluster–Level


□Algorithm design

❖ Search space

- Pipeline parallelism size (**s**)
- Number of full–memory workers (**w**)

❖ TTFT estimates

- 1 worker





HydraServe: Cluster–Level


□Algorithm design

❖ Search space

- Pipeline parallelism size (**s**)
- Number of full–memory workers (**w**)

❖ TTFT estimates

- 1 worker
- Inference time for each **s** workers
 - Data transmission : t_n
 - Prefill
 - ✓ Full–memory workers : $t_p \times \frac{1}{s}$
 - ✓ Low–memory workers : t_p





HydraServe: Cluster–Level

□Algorithm design

❖ Search space

- Pipeline parallelism size (**s**)
- Number of full–memory workers (**w**)

❖ TTFT estimates

- 1 worker
- Inference time for each **s** workers

- Data transmission : t_n
- Prefill
 - ✓ Full–memory workers : $t_p \times \frac{1}{s}$
 - ✓ Low–memory workers : t_p



Total inference time is :

$$t_p \times \left(s - w + \frac{w}{s} \right) + t_n \times s$$



HydraServe: Cluster–Level


□Algorithm design

❖ Search space

- Pipeline parallelism size (**s**)
- Number of full–memory workers (**w**)

❖ TTFT estimates

- 1 worker





HydraServe: Cluster–Level


□Algorithm design

❖ Search space

- Pipeline parallelism size (**s**)
- Number of full–memory workers (**w**)

❖ TTFT estimates

- **s** workers





HydraServe: Cluster–Level

□Algorithm design

❖ Search space

- Pipeline parallelism size (**s**)
- Number of full–memory workers (**w**)

❖ TTFT estimates

$$t_c + \frac{M}{s} \times \max_i \left\{ \frac{1}{b_i} + \frac{1}{p_i} \right\} + t_p \times \left(s - w + \frac{w}{s} \right) + t_n \times s$$



HydraServe: Cluster–Level

□Algorithm design

❖ Search space

- Pipeline parallelism size (**s**)
- Number of full–memory workers (**w**)

❖ TTFT estimates

$$t_c + \frac{M}{s} \times \max_i \left\{ \frac{1}{b_i} + \frac{1}{p_i} \right\} + t_p \times \left(s - w + \frac{w}{s} \right) + t_n \times s$$

❖ TPOT estimates

$$t_d \times \left(s - w + \frac{w}{s} \right) + t_n \times s$$



HydraServe: Cluster–Level

Algorithm design

❖ Search space

➤ Pipeline parallelism size (**s**)

➤ Number of full–memory workers (**w**)

❖ TTFT estimates

❖ TPOT estimates

❖ Server selection

➤ Priority smallest $\max_i \left\{ \frac{1}{b_i} + \frac{1}{p_i} \right\}$

Algorithm 1 Resource Allocation Algorithm

Input: time cost of container creation and runtime initialization t_c , data transmission t_n , prefill t_p , and decoding t_d ; model size M ; GPU server network bandwidth b_i and PCIe bandwidth p_i ; user specified requirements SLO_{TTFT} and SLO_{TPOT} .

Output: pipeline parallelism size s , #full-memory workers w , and selected GPU servers g .

```

 $S \leftarrow \emptyset$ 
for  $s \in \{1, 2, \dots, 4\}$  do
    for  $w \in \{0, 1, \dots, s\}$  do
         $i_1, i_2, \dots, i_k \leftarrow$  Servers that fit a model of size  $M$ .
         $j_1, j_2, \dots, j_l \leftarrow$  Servers that fit a model of size  $M/s$ .
         $j'_1, \dots, j'_{l'} \leftarrow$  MergeSort( $(j_1, \dots, j_l), (i_{w+1}, \dots, i_k)$ )
         $g \leftarrow (i_1, i_2, \dots, i_w, j'_1, \dots, j'_{s-w})$ 
        max_ratio  $\leftarrow \max_{x \in g} \left( \frac{1}{b_x} + \frac{1}{p_x} \right)$ 
        TTFT  $\leftarrow t_c + \frac{M}{s} \times \text{max\_ratio} + t_p \times (s - w + \frac{w}{s}) + t_n \times s$ 
        TPOT  $\leftarrow t_d \times (s - w + \frac{w}{s}) + t_n \times s$ 
        if  $TTFT \leq SLO_{TTFT}$  and  $TPOT \leq SLO_{TPOT}$  then
             $S \leftarrow S \cup \{(s, w, g)\}$ 
        if  $S$  is  $\emptyset$  then
            return  $(1, 1, (i_1))$   $\triangleright$  Use single worker if no solution
        else
             $c \leftarrow$  Scheme that incurs minimal GPU sharing from  $S$ 
return  $c$ 

```



HydraServe: Cluster–Level

□ Network–Contention

- ❖ Model fetching & Inference (small intermediate results)



Prioritize inference packets



HydraServe: Cluster–Level

□ Network–Contention

- ❖ Model fetching & Inference (small intermediate results)
- ❖ Workers fetch models on the same GPU



HydraServe: Cluster–Level

□ Network–Contention

- ❖ Model fetching & Inference (small intermediate results)
- ❖ Workers fetch models on the same GPU

A GPU server with N workers comes a new one

- Bandwidth estimates : $\frac{B}{N+1}$
- Worker_i time to fetch the pending model (S_i) : $T_{rest} = S_i \div \frac{B}{N+1}$



HydraServe: Cluster–Level

□ Network–Contention

- ❖ Model fetching & Inference (small intermediate results)
- ❖ Workers fetch models on the same GPU

A GPU server with N workers comes a new one

➤ Bandwidth estimates : $\frac{B}{N+1}$

➤ Worker_i time to fetch the pending model (S_i) : $T_{rest} = S_i \div \frac{B}{N+1}$

➤ S_i estimates : $S'_i = S_i - \frac{B}{N} \times (T_{curr} - T')$

Start or completion of a cold start time

We should meet : $T_{rest} \leq \text{Deadline}_i - T_{curr}$




HydraServe: Worker-Level

□ Model prefetching

- ❖ Shared memory is pre-allocated

□ Parameter Manager

- ❖ An individual thread






HydraServe: Worker-Level

□ Model prefetching

- ❖ Shared memory is pre-allocated
- ❖ Starts before container creation
 - Overlap container and runtime init

□ Parameter Manager

- ❖ An individual thread





HydraServe: Worker-Level

□ Model prefetching


- ❖ Shared memory is pre-allocated
- ❖ Starts before container creation
 - Container and runtime init are overlapped

□ Parameter Manager

- ❖ An individual thread
- ❖ Container first init it
- ❖ Streaming manner with zero-copy






Model prefetch and load can be pipelined





HydraServe: Worker–Level

- ❑ Model prefetching
- ❑ Parameter Manager
- ❑ Prioritize CUDA context init
 - ❖ Overlap library and model load





HydraServe: Inference-Level

□ Worker Scaling

Parameter manager continue to load the remind part

❖ Scaling down

- Only one worker fetched unloaded model
- Default mechanism


❖ Scaling up

- All workers become individual serving endpoint
- Tackle load spikes
 - Use sliding window strategy to predict




HydraServe: Inference-Level


Worker Scaling



Scaling down



Scaling up






HydraServe: Inference-Level

- ❑ Worker Scaling
- ❑ KV cache migration
 - ❖ Stop scheduling & wait for on-the-fly batches
 - ❖ Gather the blocks
 - ❖ Place the block at different layers



HydraServe: Inference-Level


- ❑ Worker Scaling
- ❑ KV cache migration
 - ❖ Stop scheduling & wait for on-the-fly batches
 - ❖ Gather the blocks
 - ❖ Place the block at different layers





HydraServe: Inference-Level

- ❑ Worker Scaling
- ❑ KV cache migration
 - ❖ Stop scheduling & wait for on-the-fly batches
 - ❖ Gather the blocks
 - ❖ Place the block at different layers





Implementation

- ❑ Instance startup optimizations in vLLM
 - ❖ Postpone KV Cache Allocation on CPU
 - ❖ Skip Online Memory Profiling
 - ❖ Direct GPU Tensor Usage (Zero-Copy)



Evaluation : Latency & Pipeline Consolidation

□ Testbed

- ❖ A cluster with 4 A10 servers & 4V100 servers
 - A10 server contains single A10 GPU and 188GiB memory
 - V100 server contains 4 V100 GPUs and 368GiB memory
 - Network bandwidth per server is 16Gps



Evaluation : Latency & Pipeline Consolidation

□ Testbed

- ❖ A cluster with 4 A10 servers & 4V100 servers
 - A10 server contains single A10 GPU and 188GiB memory
 - V100 server contains 4 V100 GPUs and 368GiB memory
 - Network bandwidth per server is 16Gps
 - Remote storage has sufficient network capacity

□ Baselines


- ❖ Servrless vLLM
 - Iterates through all servers to select a available GPU
- ❖ ServerlessLLM^[1]

[1] ServerlessLLM: Low-Latency Serverless Inference for Large Language Models



Evaluation : Cold start Latency

□Latency



HydraServe reduces cold start latency by **2.1x–4.7x** compared to serverless vLLM


1.7x–3.1x compared to serverlessLLM



Evaluation : Cold start Latency

□ Breakdown

- ❖ Prefetch : model prefetcher
- ❖ Stream : streaming loading & implementation optimize
- ❖ Overlap : overlap library & model loading
- ❖ Parallel : parallelize the cold start






Evaluation : Pipeline Consolidation

Llama2-13B on V100, input length is 512 tokens

□ Scaling down




- ❖ Reduce end-to-end generation time by **1.61x–1.70x**
- ❖ Maintain almost same inference speed during early start




Evaluation : Pipeline Consolidation

Llama2-13B on V100, input length is 512 tokens

□ Scaling up



(a) Average TTFT of different loads.



(b) Average TPOT of different loads.

Maximum batch size : 8

- ❖ 128 concurrent requests, reduce average TTFT by **1. 87×**
- ❖ Average TPOT only increase by **1. 08×—1.19×**



Evaluation : End-to-End

□ Testbed

- ❖ A cluster with 2 A10 servers
 - 4 A10 GPUs , 752GiB memory , 64Gbps network bandwidth
- ❖ A cluster with 4 V100 servers
 - 4 V100 GPUs , 368GiB memory , 16Gbps network bandwidth
- ❖ Remote storage has sufficient network capacity

□ Baselines

- ❖ Servrless vLLM
 - Iterates through all servers to select a available GPU
- ❖ ServerlessLLM^[1]

[1] ServerlessLLM: Low-Latency Serverless Inference for Large Language Models



Evaluation : End-to-End

□ Workload

❖ SLO based on warm requests (1024 input tokens, 8 batch sizes)

Model	Model Size	GPU Card	TTFT	TPOT
Llama2-7B	12.5GB	A10	1.5s	42ms
Llama2-13B	24.2GB	V100	2.4s	58ms



TTFT and TPOT of warm request

Application	TTFT	TPOT	Dataset
Chatbot Llama2-7B	7.5s	200ms	ShareGPT
Chatbot Llama2-13	12s	200ms	ShareGPT
Code Completion Llama2-7B	7.5s	84ms	HumanEval
Code Completion Llama2-13B	12s	116ms	HumanEval
Summarization Llama2-7B	15s	84ms	LongBench
Summarization Llama2-13B	24s	116ms	LongBench

SLOs summary



Evaluation : End-to-End

□ Workload

❖ SLO based on warm requests (1024 input tokens, 8 batch sizes)

Model	Model Size	GPU Card	TTFT	TPOT
Llama2-7B	12.5GB	A10	1.5s	42ms
Llama2-13B	24.2GB	V100	2.4s	58ms



TTFT and TPOT of warm request

Application	TTFT	TPOT	Dataset
Chatbot Llama2-7B	7.5s	200ms	ShareGPT
Chatbot Llama2-13	12s	200ms	ShareGPT
Code Completion Llama2-7B	7.5s	84ms	HumanEval
Code Completion Llama2-13B	12s	116ms	HumanEval
Summarization Llama2-7B	15s	84ms	LongBench
Summarization Llama2-13B	24s	116ms	LongBench

SLOs summary

❖ Workloads


➤ Microsoft Azure Function Trace

- Models are mapped, round-robin approach



Evaluation : End-to-End

□ Change the CV




- ❖ Achieve **1.43x–1.74x** higher TTFT SLO during bursty requests
- ❖ Caching further improve **1.11x**



Evaluation : End-to-End

□ Change the SLO scales




- ❖ Tight SLO : Meet ~63% TTFT
- ❖ Loose SLO : Achieve 1. 38x–1.52x improvement
1. 49x–1.58x improvement with caching



Evaluation : End-to-End

□ Application analysis

CV = 8, RPS = 0.6





- ❖ Chatbot & Code completion: 1. **61x–1.70x** improvement in TTFT
- ❖ Code completion has lower TTFT SLO attainment
 - Workers keep alive for shorter time, cause more cold starts



Evaluation : End-to-End

□ TPOT and resource usage penalties

CV = 8, RPS = 0.6




Normalize to serverless vLLM

- ❖ TPOT: A **1.06x** average increase
- ❖ Cost : HydraServe consumes lower in most case
 - Pipeline groups merge is quickly
 - Fast worker startup reduce GPU usage



Evaluation : Brownfield

Llama2-7B on NVIDIA A10 GPUs with 24GB GPU memory



- ❖ **HydraServe achieves an average $2.6\times$ reduction in cold-start TTFT**



Conclusion

□ Pros:

- ❖ Use pipeline parallelism to reduce fetching time
- ❖ Overlap load model and load library
- ❖ Consolidate pipeline to reduce overhead

□ Cons:



- ❖ Remote storage needs sufficient bandwidth
- ❖ Due to lack of NVLink, evaluated models can reside in a GPU



Q&A



Background: Out-of-order execution





HydraServe: Inference-Level

□ Worker Scaling

❖ Scaling down

