

# Pie: A Programmable Serving System for Emerging LLM Applications

In Gim, Zhiyao Ma, Seung-seob Lee, Lin Zhong

Yale University

SOSP' 25

# Outline

---

## □ Background

- ❖ Existing LLM Serving Systems
- ❖ Case Studies
- ❖ Motivation

## □ Design



## □ Evaluation

## □ Discussion

# Existing LLM Serving Systems

## □ LLM Inference

- ❖ Generate tokens **autoregressively**
- ❖ **KV Cache** stores the key/value vectors of previous tokens
  - Modern serving system pre-allocates device memory as KV cache pool and organize KV cache into virtual blocks




# Existing LLM Serving Systems

## □ LLM Serving Systems: vLLM, SGLang, ...

### ❖ Architecture (vLLM V0)


- **Scheduler:** decide which requests go into next iteration
- **Model Executor:** run forward passes
- **Structured Output Manager:** handles constrained decoding



# Existing LLM Serving Systems

## □ LLM Serving Systems: vLLM, SGLang, TensorRT, ...

- ❖ **Architecture**
- ❖ **Generation:** monolithic, schedule – forward – sample loop
- ❖ **KV Cache Management:** system-wide, static policy
- ❖ System behavior is determined by both request parameters and initial system configurations



# **Case Studies on LLM applications**

---

**□ Modern LLM applications and decoding techniques requires flexible generation policy**

**❖ Post-processing**

- Structured Output / Constrained Decoding

**❖ Special Decoding Techniques**

- Speculative Decoding

**❖ Customized Resource Management**

- StreamingLLM / Attention Sink

**❖ Agents**

- Agentic Workflow


# Case Study: Structured Output

## □ Constrained Decoding: Enforce the output format of LLM

```
class CapitalInfo:  
    name: str = Field(pattern=r"^\w+\$")  
    latitude: str = Field(pattern=...)  
    longitude: str = Field(pattern=...)
```


User: Give me the information of the capital of France in the JSON format.

LLM: {  
 "name": "France",  
 "latitude": "48.86N",  
 "longitude": "2.35E",  
}



# Case Study: Post-Processing


- **Constrained Decoding:** Enforce the output format of LLM
- **In Existing Systems:**
  - ❖ ✗ inject logic around sampling step



# Case Study: Special Decoding Technique

## ❑ N-Gram Prompt-Lookup Speculative Decoding

- ❖ Copy from input and generated output as draft token (No draft model)
- ❖ Usage: code edition<sup>1</sup>, summarization, document QA, ...




1: <https://zed.dev/blog/edit-prediction#minimizing-latency-speculative-decoding>.

# Case Study: Special Decoding Technique


## ❑ N-Gram Prompt-Lookup Speculative Decoding

- ❖ Copy from input and generated output as draft token (No draft model)
- ❖ Usage: code edition, summarization, document QA, ...

Prompt: What is the main idea of this paper?  
Current Output: The main idea of the  
Speculation: <None>



Verification



Speculation




Accepted

# Case Study: Special Decoding Technique

## ❑ N-Gram Prompt-Lookup Speculative Decoding

- ❖ Copy from input and generated output as draft token (No draft model)
- ❖ Usage: code edition, summarization, document QA, ...

Prompt: What is the main idea of this paper?  
Current Output: The main idea  
Speculation: of this paper




# Case Study: Special Decoding Technique

## □ Decoding Techniques


- ❖ n-gram prompt-lookup speculative decoding, look-ahead decoding, ...
- ❖ Breaks the schedule-forward-sample loop

## □ In existing systems:

- ❖ **X Inflexible:** implemented as **system-wide toggles**



n-gram prompt lookup spec-decode





look-ahead decoding

# Case Study: Resource Management

## □ StreamingLLM<sup>[1]</sup>: Sliding window w/ Attention Sinks

- ❖ KV cache blocks **masked** (and **evicted**) during the generation process




[1] Efficient Streaming Language Models with Attention Sinks. ICLR' 24

# Case Study: Resource Management

## □ StreamingLLM<sup>[1]</sup>: Sliding window w/ Attention Sinks

- ❖ KV cache blocks **masked** (and **evicted**) during the generation process
- ❖ **X Invasive** modification **deep within** the memory manager




[1] Efficient Streaming Language Models with Attention Sinks. *ICLR' 24*

# Case Study: Agents

## ❑ **Agentic Workflow:** Thought-Action-Observation cycle


- ❖ **Thought:** plan the next step
- ❖ **Action:** take action by calling tool
- ❖ **Observation:** reflect on the response from the tool (*external environment*)



# Case Study: Agents

## ❑ Agentic Workflow: Thought-Action-Observation cycle

- ❖ ✗ Round-trip control between client and server
  - Incurs **network latency**
  - Implicit session persistence → possible **costly re-prefill**



# Motivation

---

## □ Current systems:

### ❖ Inflexible Generation Loop

- Logic injection for controlling output distribution
- Static, system-wide toggles for special decoding technique

### ❖ Implicit Resource Management

- Failed to provide application-specific control

### ❖ Poor Workflow Integration

- Round-trip control switching

## □ The monolithic generation loop couples the application control logic and core execution engine.

# Outline

---

## ❑ Background

## ❑ Design

- ❖ Overview
- ❖ Programming Model
- ❖ System Architecture & Implementation


## ❑ Evaluation

## ❑ Discussion

# Pie – Overview

## ❑ Pie: Serve programs, not prompts


- ❖ breaks monolithic generation loop into fine-grained API calls
- ❖ Inferlet: user-provided program



# Pie – Overview

## ❑ Pie: Serve programs, not prompts

- ❖ Schedule: Application Layer + Control Layer
- ❖ Forward: Inference Layer
- ❖ Sample: Application Layer



# Pie – Overview

## ❑ Pie

### ❖ Schedule

#### ➤ Application Layer:

- Define control logic

#### ➤ Control Layer

- Handle non-GPU API call
- Batch GPU API call

### ❖ Forward: Inference Layer

#### ➤ Handle GPU API call

### ❖ Sample: Application Layer

#### ➤ Sample probs with customized logic

```
model = get_auto_model()
q = model.create_queue()

ids = tokenize(q, "Hello, ")
ilen = len(ids)
limit = ilen + 10

# resource allocation
iemb = alloc_emb(q, ilen)
oemb = alloc_emb(q, 1)
kv = alloc_kvpage(q, limit)

# prefill
pos = len(range(ilen))
embed_txt(q, ids, pos, iemb)
await forward(q, [], iemb, kv[:ilen], oemb)

# decode
for i in range(ilen, limit):
    dist = await get_next_dist(q, oemb)
    ids.extend(dist.max_index())
    embed_txt(q, ids[-1], [i], oemb)
    await forward(q, kv[:i], oemb, kv[i:i+1], oemb)

# output
output = detokenize(q, ids)
print(output)

# resource cleanup
dealloc_emb(q, iemb)
dealloc_emb(q, oemb)
dealloc_kvpage(q, kv)
```

# Pie – Overview

## ❑ Pie

### ❖ Schedule

#### ➤ Application Layer:

- Define control logic

#### ➤ Control Layer

- Handle non-GPU API call
- Batch GPU API call

### ❖ Forward: Inference Layer

#### ➤ Handle GPU API call

### ❖ Sample: Application Layer

#### ➤ Sample probs with customized logic

Application  
Control  
Logic

```
model = get_auto_model()
q = model.create_queue()

ids = tokenize(q, "Hello, ")
ilen = len(ids)
limit = ilen + 10

# resource allocation
iemb = alloc_emb(q, ilen)
oemb = alloc_emb(q, 1)
kv = alloc_kvpage(q, limit)

# prefill
pos = len(range(ilen))
embed_txt(q, ids, pos, iemb)
await forward(q, [], iemb, kv[:ilen], oemb)

# decode
for i in range(ilen, limit):
    dist = await get_next_dist(q, oemb)
    ids.extend(dist.max_index())
    embed_txt(q, ids[-1], [i], oemb)
    await forward(q, kv[:i], oemb, kv[i:i+1], oemb)

# output
output = detokenize(q, ids)
print(output)

# resource cleanup
dealloc_emb(q, iemb)
dealloc_emb(q, oemb)
dealloc_kvpage(q, kv)
```

# Pie – Overview

## ❑ Pie

### ❖ Schedule

#### ➤ Application Layer:

- Define control logic

#### ➤ Control Layer

- Handle non-GPU API call
- Batch GPU API call

### ❖ Forward: Inference Layer

#### ➤ Handle GPU API call

### ❖ Sample: Application Layer

#### ➤ Sample probs with customized logic

### Sampling

```
model = get_auto_model()
q = model.create_queue()

ids = tokenize(q, "Hello, ")
ilen = len(ids)
limit = ilen + 10

# resource allocation
iemb = alloc_emb(q, ilen)
oemb = alloc_emb(q, 1)
kv = alloc_kvpage(q, limit)

# prefill
pos = len(range(ilen))
embed_txt(q, ids, pos, iemb)
await forward(q, [], iemb, kv[:ilen], oemb)

# decode
for i in range(ilen, limit):
    dist = await get_next_dist(q, oemb)
    ids.extend(dist.max_index())
    embed_txt(q, ids[-1], [i], oemb)
    await forward(q, kv[:i], oemb, kv[i:i+1], oemb)

# output
output = detokenize(q, ids)
print(output)

# resource cleanup
dealloc_emb(q, iemb)
dealloc_emb(q, oemb)
dealloc_kvpage(q, kv)
```

# Pie – Overview

---

## □ Pie

### ❖ Customizable generation loop

- Text generation: schedule – forward – sample
- SpecDecode: speculate – schedule – forward – verify – drop KV

### ❖ Explicit, fine-grained resource control

- Attention sink: mask & drop KV cache
- Graph-of-Thought, Tree-of-Thought: KV cache reusing

### ❖ Workflow Integration

- Agents: tool handler within inferlet

# Pie – Programming Model

## □ Pie provides 42 APIs

- ❖ 18 APIs for LLM execution and resource management
- ❖ 24 APIs for core runtime operations and agentic workflow

Trait (Supertraits)	Function	Behavior
	<code>get_arg() -&gt; list[str]</code> <code>send(msg)</code> <code>receive() -&gt; future[str]</code> <code>http_get(url) -&gt; future[str]</code> <code>available_models() -&gt; list[Model]</code> <code>available_traits(model) -&gt; list[str]</code> <code>create_queue(model) -&gt; Queue</code> <code>synchronize(q) -&gt; future</code> <code>set_queue_priority(q, pri)</code>	Gets command-line arguments passed during launch. Sends a message to the client that launched the inferlet. Receives messages from the client. Performs an HTTP GET request to the specified URL. Gets the list of models. Gets the list of given model's traits. Creates a new command queue. Blocks until the queue finishes execution. Hints the controller which queue to process first.
<i>Allocate</i>	<code>export_kvpage(kv, name)</code> <code>import_kvpage(name) -&gt; list[KvPage]</code> <code>alloc_kvpage(q, size) -&gt; list[KvPage]</code> <code>dealloc_kvpage(q, kv)</code> <code>alloc_emb(q, size) -&gt; list[Embed]</code> <code>dealloc_emb(q, emb)</code> <code>copy_kvpage(q, src, dst)</code>	Exports paged KV cache for use in other programs. Imports the paged KV cache. Allocates memory for paged KV cache. Deallocates memory for paged KV cache. Allocates buffer space for embeddings. Deallocates buffer space for embeddings. Copy KV cache contents at token-level.
<i>Forward (Allocate)</i>	<code>forward(q, ikv, iemb, okv, oemb, mask)</code> <code>mask_kvpage(q, tgt, mask)</code>	Transform to KV cache and/or output embeddings. Masks the KV cache in a token-level manner.
<i>InputText (Allocate, Forward)</i>	<code>embed_txt(q, tok, pos, embs)</code>	Embeds text input into embeddings
<i>InputImage (Allocate, Forward)</i>	<code>num_embs_needed(m, size) -&gt; int</code> <code>embed_img(q, blob, embs)</code>	Calculates the number of embeddings needed. Embeds image input into embeddings.
<i>Tokenize (InputText)</i>	<code> tokenize(q, text) -&gt; list[int]</code> <code>detokenize(q, token_ids) -&gt; str</code> <code>get_vocab(q) -&gt; list[list[byte]]</code>	Converts text into a list of token IDs. Converts token IDs back into text. Retrieves the vocabulary list.
<i>OutputText (Allocate)</i>	<code>get_next_dist(q, emb) -&gt; future[Dist]</code>	Gets the next token distribution.

# Pie – Programming Model

## □ Abstractions: Forward pass

❖ A forward pass of LLM: embed – forward – sample

Trait (Supertraits)	Function	Behavior
	<code>get_arg() -&gt; list[str]</code> <code>send(msg)</code> <code>receive() -&gt; future[str]</code> <code>http_get(url) -&gt; future[str]</code> <code>available_models() -&gt; list[Model]</code> <code>available_traits(model) -&gt; list[str]</code> <code>create_queue(model) -&gt; Queue</code> <code>synchronize(q) -&gt; future</code> <code>set_queue_priority(q, pri)</code>	Gets command-line arguments passed during launch. Sends a message to the client that launched the inferlet. Receives messages from the client. Performs an HTTP GET request to the specified URL. Gets the list of models. Gets the list of given model's traits. Creates a new command queue. Blocks until the queue finishes execution. Hints the controller which queue to process first.
<code>Allocate</code>	<code>export_kvpage(kv, name)</code> <code>import_kvpage(name) -&gt; list[KvPage]</code> <code>alloc_kvpage(q, size) -&gt; list[KvPage]</code> <code>dealloc_kvpage(q, kv)</code> <code>alloc_emb(q, size) -&gt; list[Embed]</code> <code>dealloc_emb(q, emb)</code> <code>copy_kvpage(q, src, dst)</code>	Exports paged KV cache for use in other programs. Imports the paged KV cache. Allocates memory for paged KV cache. Deallocates memory for paged KV cache. Allocates buffer space for embeddings. Deallocates buffer space for embeddings. Copy KV cache contents at token-level.
<code>Forward (Allocate)</code>	<code>forward(q, ikv, iemb, okv, oemb, mask)</code> <code>mask_kvpage(q, tgt, mask)</code>	Transform to KV cache and/or output embeddings. Masks the KV cache in a token-level manner.
<code>InputText (Allocate, Forward)</code>	<code>embed_txt(q, tok, pos, embs)</code>	Embeds text input into embeddings
<code>InputImage (Allocate, Forward)</code>	<code>num_embs_needed(m, size) -&gt; int</code> <code>embed_img(q, blob, embs)</code>	Calculates the number of embeddings needed. Embeds image input into embeddings.
<code>Tokenize (InputText)</code>	<code> tokenize(q, text) -&gt; list[int]</code> <code>detokenize(q, token_ids) -&gt; str</code> <code>get_vocab(q) -&gt; list[list[byte]]</code>	Converts text into a list of token IDs. Converts token IDs back into text. Retrieves the vocabulary list.
<code>OutputText (Allocate)</code>	<code>get_next_dist(q, emb) -&gt; future[Dist]</code>	Gets the next token distribution.

# Pie – Programming Model

## ❑ Abstractions: Explicit Resource Managements

- ❖ Embeds (per-token) + KvPage (8 – 32 token)
- ❖ Each inferlet has its own virtual address space (enable KV sharing)

Trait (Supertraits)	Function	Behavior
	<code>get_arg() -&gt; list[str]</code> <code>send(msg)</code> <code>receive() -&gt; future[str]</code> <code>http_get(url) -&gt; future[str]</code> <code>available_models() -&gt; list[Model]</code> <code>available_traits(model) -&gt; list[str]</code> <code>create_queue(model) -&gt; Queue</code> <code>synchronize(q) -&gt; future</code> <code>set_queue_priority(q, pri)</code>	Gets command-line arguments passed during launch. Sends a message to the client that launched the inferlet. Receives messages from the client. Performs an HTTP GET request to the specified URL. Gets the list of models. Gets the list of given model's traits. Creates a new command queue. Blocks until the queue finishes execution. Hints the controller which queue to process first.
Allocate	<code>export_kvpage(kv, name)</code>	Exports paged KV cache for use in other programs.
	<code>import_kvpage(name) -&gt; list[KvPage]</code>	Imports the paged KV cache.
	<code>alloc_kvpage(q, size) -&gt; list[KvPage]</code>	Allocates memory for paged KV cache.
	<code>dealloc_kvpage(q, kv)</code>	Deallocates memory for paged KV cache.
	<code>alloc_emb(q, size) -&gt; list[Embed]</code>	Allocates buffer space for embeddings.
	<code>dealloc_emb(q, emb)</code>	Deallocates buffer space for embeddings.
	<code>copy_kvpage(q, src, dst)</code>	Copy KV cache contents at token-level.
Forward (Allocate)	<code>forward(q, ikv, iemb, okv, oemb, mask)</code> <code>mask_kvpage(q, tgt, mask)</code>	Transform to KV cache and/or output embeddings. Masks the KV cache in a token-level manner.
InputText (Allocate, Forward)	<code>embed_txt(q, tok, pos, embs)</code>	Embeds text input into embeddings
InputImage (Allocate, Forward)	<code>num_embs_needed(m, size) -&gt; int</code> <code>embed_img(q, blob, embs)</code>	Calculates the number of embeddings needed. Embeds image input into embeddings.
Tokenize (InputText)	<code> tokenize(q, text) -&gt; list[int]</code> <code>detokenize(q, token_ids) -&gt; str</code> <code>get_vocabs(q) -&gt; list[list[byte]]</code>	Converts text into a list of token IDs. Converts token IDs back into text. Retrieves the vocabulary list.
OutputText (Allocate)	<code>get_next_dist(q, emb) -&gt; future[Dist]</code>	Gets the next token distribution.

# Pie – Programming Model

---

## ❑ Abstractions: Command Queue

- ❖ A logical sequence of API calls
- ❖ Inform control layer with
  - Data dependency
  - Priority

```
model = get_auto_model()
q1, q2 = model.create_queue(), model.create_queue()

forward(q1, [], iemb1, [], oemb1)
forward(q2, [], iemb2, [], oemb1) batched API calls

await synchronize(q1)
await synchronize(q1)
```

# Pie – Programming Model

## □ Support Library

- ❖ Provides pre-defined control logic

```
model = get_auto_model()  
  
ctx = model.create_context()  
  
ctx.fill("Hello,")  
  
output = ctx.generate_until(max_token=10)  
print(output)  
  
del ctx
```

```
model = get_auto_model()  
q = model.create_queue()  
  
ids = tokenize(q, "Hello,")  
ilen = len(ids)  
limit = ilen + 10  
  
# resource allocation  
iemb = alloc_emb(q, ilen)  
oemb = alloc_emb(q, 1)  
kv = alloc_kvpage(q, limit)  
  
# prefill  
pos = len(range(ilen))  
embed_txt(q, ids, pos, iemb)  
await forward(q, [], iemb, kv[:ilen], oemb)  
  
# decode  
for i in range(ilen, limit):  
    dist = await get_next_dist(q, oemb)  
    ids.extend(dist.max_index())  
    embed_txt(q, ids[-1], [i], oemb)  
    await forward(q, kv[:i], oemb, kv[i:i+1], oemb)  
  
# output  
output = detokenize(q, ids)  
print(output)  
  
# resource cleanup  
dealloc_emb(q, iemb)  
dealloc_emb(q, oemb)  
dealloc_kvpage(q, kv)
```

# Pie – Programming Model

---

## ❑ Integration

- ❖ User-interaction: `send()`, `recv()`
- ❖ Networking: `http_get()`, `http_post()`
- ❖ Inter-inferlet communication: `broadcast()`, `subscribe()`

# Pie – Programming Model

---

## ❑ Integration

## ❑ Extensibility

- ❖ Set of related operations/APIs are defined as Traits
  - Specific model realizes some of (or all of) the traits.
  - Traits has dependencies
    - Forward depends on Allocate
  - Inferlets can query supported traits via `available_traits()`
- ❖ Extension (new modal, advanced optimizations) via defining and implementing new traits
  - Non-invasive: Older traits remains unchanged

# Pie – Programming Model

## ❑ Integration

## ❑ Extensibility

## ❑ Expressiveness

❖ Support a wide range of apps  
and inference techniques

➤ in less than 300 LoC


Technique	R1-3 (§1)	LoC	Wasm	Supported
Support library (§6.3)		728	-	
Text completion		38	129 KB	V, S, L
ToT [75]	R1, R3	198	148 KB	S
RoT [41]	R1, R3	106	152 KB	
GoT [7]	R1, R3	87	171 KB	
SKoT [53]	R1, R3	82	173 KB	S
Prefix caching [38]	R1	45	131 KB	V, S
Modular caching [21]	R1	72	139 KB	
EBNF decoding [13]	R2	225	2 MB	V, S, L
Beam search [18]	R2	98	142 KB	V, L
Watermarking [34]	R2	43	130 KB	
Output validation [35]	R2	52	131 KB	
Speculative decoding [62]	R2	255	152 KB	V
Jacobi decoding [61]	R2	88	96 KB	
Attention sink [74]	R1	60	133 KB	
Windowed attn [5].	R1	60	133 KB	
Hierarchical attn. [66]	R1	42	130 KB	
Agent-ReACT [76]	All	60	309 KB	
Agent-CodeACT [71]	All	62	6.7 MB	
Agent-SWARM [85]	All	95	135 KB	

V: vLLM  
S: SGLang  
L: LMQL

# Pie – System Architecture

## □ Pie: Three-layer Architecture


- ❖ Application Layer
- ❖ Control Layer
- ❖ Inference Layer



# Pie – System Architecture

## Application Layer


- ❖ WebAssembly Runtime for Isolation
- ❖ Inferlet Lifecycle Manager:
  - Manages creation, destruction and communication of inferlets
  - Provides send( ) and recv( ) API for user interaction



# Pie – System Architecture

## □ Control Layer

- ❖ Handle non-GPU API: `get_model()`, `create_queue()`, ...
- ❖ Resource management
- ❖ Batch GPU API calls




# Pie – System Architecture

## □ Control Layer

### ❖ Resource Management

- Pre-allocate device memory at startup
- Manage virtual mapping for Embeds and KvPages
- Resolve resource contention: FCFS




# Pie – System Architecture

## □ Control Layer

### ❖ Batch Scheduling

- Batches compatible GPU API calls
- Command queue provides information about *data dependency* and *priority*
  - Vertical batching: batch non-conflict / data-independent calls
  - Horizontal batching: batch calls from different queue




# Pie – System Architecture

## □ Control Layer

### ❖ Batch Scheduling


- Vertical Batching + Horizontal Batching
- Select the batch containing the longest-waiting call
- Dispatch Timing: as soon as GPU becomes idle



# Pie – System Architecture

## □ Inference Layer

- ❖ Handle GPU API call / Invoke GPU computation
- ❖ Backend: PyTorch + FlashInfer
  - An incomplete but faster C++/CUDA backend also available




# Pie – Summary

## ❑ Programming Model

- ❖ Expressive & Extensible
- ❖ Explicit Resource Management
- ❖ User-friendly Support Library

## ❑ Three-layer Architecture

- ❖ Application Layer
- ❖ Control Layer
- ❖ Inference Layer



# Outline

---

□ Background

□ Design

□ Evaluation

□ Discussion

# Evaluation

---

## ❑ Concerns questions


1. How does Pie **facilitate** the deployment of LLM applications?
2. How does Pie improve end-to-end **latency** and **throughput**?
3. What **overhead** is introduced?

## ❑ Setup

- ❖ Hardware: A single NVIDIA L4 GPU (24 GB)
- ❖ Model: Llama 3 of 1B, 3B, and 8B

# Evaluation – Agentic Workflow

## Workflow comparison




## E2E evaluation on a 1B model

- ❖ up to 15% latency reduction
- ❖ up to 30% throughput speedup

## Speedup related to I/O


- ❖ Smaller models (1,3B): reducing round-trip time
- ❖ Larger models (8B): retaining the KV cache across I/O interactions



# Evaluation – LLM inference techniques

## □ Latency & Throughput

- Improved performance for tasks with various inference techniques
  - except pure text comp. without I/O




## □ Effective expressiveness and implementation on Pie

# Evaluation – System Overhead of Pie

## □ Latency overhead of cold start

- Just like FaaS, cold start exists for user-defined sandboxed code
- Duration measured for a text completion task
  - $T_{\text{inferlet\_start\_forward}} - T_{\text{inferlet\_get\_triggered}}$
  - Negligible compared with per-token generation latency



- However, only the same inferlet is considered here
  - For different inferlets, uploading and JIT overhead would be much higher

# Evaluation – System Overhead of Pie

## □ Latency overhead per decoding step (warm start)

- Monolithic vLLM is faster for pure text completion task without I/O

Component	Latency
Text completion TPOT (vLLM)	64.06 ms
Lack of pipelined sampling on GPU	+1.320 ms
Lack of pipelined input embedding on GPU	+0.070 ms
Overhead of control layer batch scheduling	+0.050 ms
Overhead of returning output distribution	+0.070 ms
Boundary crossing (control-inference layer)	+0.006 ms
Boundary crossing (application-control layer)	+0.001 ms
Wasm processing overhead	+0.001 ms
Text completion TPOT (PIE)	65.59 ms

## □ Reduced overhead ratio

- for larger models

Param. size	vLLM	PIE	Overhead (%)
8B	64.06 ms	65.59 ms	1.53 ms (2.39%)
3B	30.30 ms	32.01 ms	1.71 ms (5.64%)
1B	16.83 ms	18.75 ms	1.92 ms (11.41%)

# Evaluation – Batch Strategy

---

## ❑ Baselines

- **Eager:** no batching
- **K-only:** fixed-size batching based on queue length K
- **T-only:** timeout-based batching based on wait time

## ❑ Adaptive batching in Pie improves throughput

- by 17x vs Eager baseline
- by 8-40% vs K-only and T-only

	<b>Eager</b>	<b>K-only</b>	<b>T-only</b>	<b>Adaptive</b>
Requests/s	5.61	30.09	78.11	84.85

# Outline

---

□ Background

□ Design

□ Evaluation

□ Discussion

# Discussion

---

## □ Contributions

- ❖ Fine-grained, customizable KV cache management
- ❖ Ability to provide client-side information to speed-up inference
- ❖ Flexible generation loop → heterogeneous workload at the same time

## □ Discussions

- ❖ Scalability: only single-instance/GPU inference is supported
- ❖ Security: handling malicious inferlet
- ❖ The inherit complexity of LLM service cannot be removed (Tesler's law)